

Education is Broken

Human Work in Times of AI

09.02.2026

Krishan Mathis

A Normative Manifesto for Education in the Age of AI	11
1. Judgement over recall	11
2. Agency before access	11
3. Education is about Navigating Lives	11
4. Institutions matter because power matters	12
5. Trade-offs must be made explicit	12
The anchor	12
What Exactly Is Broken?	13
Four dimensions of brokenness	13
1. Learning outcomes	13
2. Equity	14
3. Cost	14
4. Labour-market fit	15
Where this book focuses – and where it does not	15
What is our actual claim?	16
Who feels the brokenness: seven groups, seven experiences.....	17
Students	17
Teachers and educators	17
Educational institutions and administrators	18
Employers and corporations.....	19
Policymakers and government bodies.....	19
Families and communities	20
International organisations.....	21
The European dimension	21
The governance trap	22
Evidence vs Conjecture: AI Learning Inequality	22
What comes next	23
Annotated Source List	24
References	26
Agency: The Missing Infrastructure	28
What Agency Is — and What It Is Not	28
Why Agency Becomes the Bottleneck Now	29

How Agency Is Built — and What Prevents It	31
Design Principles for What Follows.....	33
References for This Chapter	35
The Rise of the Intelligence Factory	36
The Intelligence Factory vs. the Traditional University	36
From Batch Processing to Continuous Optimization.....	37
From Credentials to Portfolios.....	37
The Unbundling Threshold: Network Effects and Prestige Inertia	38
Replace, Complement, or Layer?	40
What AI Can Realistically Replace	40
What AI Complements	40
What Remains Institutional and Human.....	41
Comparison: Three Models of Education	41
The Intelligence Factory in Three Futures	42
New Skills for the AI Era	43
The Agency Divide	43
The Realistic Future Is Hybrid.....	44
Annotated Source List	45
References	45
How Systems Actually Change (and Why It's Slow)	46
Who Decides? The Key Actors and What Drives Them.....	46
Governments and Policymakers	46
Accreditation Bodies: Gatekeepers or Blockers?.....	47
University Leaders and Faculty	47
Employers: Loud Demands, Quiet Inertia	48
EdTech Firms and Platforms	49
Students: Consumers, Citizens, and Constrained Agents	49
Why Good Ideas Spread Unevenly.....	49
Micro-Credentials: Same Idea, Different Speeds	50
The WGU Audit: When Regulation Meets Innovation	50
Adoption Reality Checks.....	51
Competency-Based Education (CBE).....	51

Micro-Credentials.....	51
AI-First Educational Models.....	52
European Universities Alliances.....	52
A Scenario Matrix: What Shapes the Next Decade	53
Quadrant 1: Open Rules, Strong Capacity (Accelerated Transformation)	53
Quadrant 2: Open Rules, Weak Capacity (Uneven Disruption)	53
Quadrant 3: Closed Rules, Strong Capacity (Frustrated Innovation)	53
Quadrant 4: Closed Rules, Weak Capacity (Institutional Stagnation)	53
What This Means	54
Five Feedback Loops in Education.....	54
Loop 1: Accreditation → Institutional Structure → Program Design	54
Loop 2: Hiring Practices → Credential Demand → Institutional Investment	54
Loop 3: Funding Formula → Institutional Priority → Learner Access	55
Loop 4: EdTech Market Dynamics → Platform Consolidation → Institutional Lock-in	55
Loop 5: Stratification → Credibility Divergence → Legitimacy Erosion	55
Understanding Loop Dynamics Across Futures	55
In Drift: All Five Loops Remain Unchanged or Tighten.....	55
In Platform World: Loops 1, 3, and 5 Tighten; Loops 2 and 4 Are Hijacked.....	56
In Redesign: Loops Are Intentionally Rebalanced.....	56
The Hard Choices: What We Must Give Up.....	56
1. We must give up the lecture as the default unit of teaching.....	56
2. We must give up the fiction that time spent equals learning.....	57
3. We must give up the idea that fairness means treating everyone the same	57
4. We must give up the degree as a universal proxy for merit.....	57
5. We must give up the illusion of frictionless reform.....	58
6. We must give up the belief that education can remain neutral	58
The uncomfortable truth	58
That is the price of relevance.....	59
New Models of Learning: What's Already Here	60
1. Competency-Based Online Universities.....	60
Westen Governors University (WGU) and Southern New Hampshire University (SNHU).....	60

2. Global Campus-Free Models.....	63
Minerva University	63
3. Peer-Driven, Teacher-Free Models	65
42 School (Ecole 42)	65
4. Employer-Sponsored Learning Ecosystems.....	67
Guild Education	67
5. European University Alliances	69
ECIU, CHARM-EU, and the European Universities Initiative	69
Core Design Principles: What the Models Share.....	72
From knowledge transmission to capability building	75
Continuous curriculum renewal	75
Hybrid human-AI mentoring	76
Agency: How the Models Perform	76
What This Landscape Tells Us	77
Degrees, Skills, and the New Signalling Stack.....	79
Why Did Degrees Become So Powerful?	79
A Brief History of Signalling.....	79
The Degree Premium – and Its Limits.....	80
What Is Changing?	81
The Employer Shift Toward Skills	81
Portfolios and Demonstrated Capability.....	81
Blockchain Credentials and Verifiable Trust	81
AI-Verified Work and Behavioural Signals	82
The Four-Layer Signalling Stack.....	82
Layer 1: Baseline – Degrees	82
Layer 2: Skill – Micro-Credentials, Short Courses, Certifications.....	83
Layer 3: Evidence – Portfolios, Projects, Simulations, Code Repositories	83
Layer 4: Behaviour – AI-Derived Patterns (Speculative).....	84
Degrees Are Not Dead – Yet.....	84
Where Portfolios Already Dominate	84
Where Degrees Remain Hard Gates	84
The Messy Middle	85

A Normative Manifesto for Education in the Age of AI

How Large Is the Degree-Free Labour Market?.....	85
The Credential Inflation Problem	86
Employer Confusion and the Trust Gap.....	86
The Rational Inertia of Degree Requirements.....	87
Measurement, Not Gatekeeping: Why the Signalling Stack Requires Better Data ..	89
Evidence vs Conjecture: The Signalling Stack and the Demotion of Degrees	90
Time Horizons: What Changes When.....	91
Short Term (Next 5 Years): Multiple Possible Paths	91
Medium Term (Next 10 Years).....	92
Long Term (Next 20 Years).....	93
What This Means for You	94
Conclusion	94
Annotated Source List	95
The Stratification Trap: Who Gets Left Behind	96
1. When More Opportunity Means More Inequality	96
2. The Agency Divide.....	97
3. Why “Self-Directed Learners” Are Often the Already Privileged	98
4. When Companies Own Your Learning Path.....	100
5. The Matthew Effect in University Alliances	101
6. Digital Divides and Mental Health	102
Non-Negotiable Design Features.....	105
Operationalising Equity: Three Levels of Intervention	105
AI Tutoring and AI-First Models	109
Peer-Driven and Teacher-Free Models (42 School and similar).....	111
Campus-Free Global Models (Minerva and similar).....	111
Employer-Sponsored Models (Guild and similar)	112
European University Alliances	112
Across All Models	113
The Bottom Line	113
Annotated Source List	114
References	117
Learning in Companies: Second Chance or Parallel System?	119

What Companies Fix – and What They Can’t	119
When Corporate Learning Expands Opportunity: A Typology	120
The New Landscape: From Credentials to Capabilities	123
The Skills-First Revolution	123
The AI Skills Paradox	123
The Continuous Learning Imperative	123
AI as Transformation Engine	124
Personalized Learning at Scale	124
Predictive Analytics: Seeing Around Corners	124
Immersive Technologies: Practice Without Risk	125
Retention and Development: Why Learning Keeps People	125
The Development-Retention Nexus.....	125
Predictive Retention: Intervening Before Resignation	126
Hyperpersonalization.....	126
Career Development: From Ladders to Lattices	126
Dynamic Talent Marketplaces.....	126
Skills-Based Pathing	127
Continuous Assessment	127
Critical Success Factors	127
Human-AI Collaboration	127
Uniquely Human Skills	128
Culture and Change Management	128
Ethics and Privacy.....	128
What This Means for You	129
Annotated Source List	129
References	130
AI in Hiring: How Jobs Really Change Education.....	132
How AI Screens You	132
Skills-Based Hiring: Promise vs Reality	133
New Ways of Being Assessed	134
The Regulation Divide	135
The European Talent Crisis.....	137

Feedback Loops: How Hiring Reshapes Education	138
What This Means for You	139
Annotated Source List	141
References	141
Universities as Public Institutions in an AI Era	143
Why Societies Still Need Universities.....	143
The Crucible of Character – Made Concrete	144
The Selectivity Paradox: How Crucibles Avoid Reproducing Elite Bias	145
The Cost Architecture of Inclusive Crucibles	147
Academic Freedom and Institutional Autonomy in an AI Era	148
Civic Education in an Age of Algorithmic Manipulation	149
Funding Models for the Hybrid University	152
A Concrete Funding Architecture: How Mission Survives Market Pressure.....	153
The Risk of Mission Drift.....	158
Design Anchors for the AI-Era University	159
Conclusion	161
Three Futures for Education: Choice, Not Fate.....	162
How to Read This Chapter.....	162
Future 1: “Drift” – Business as Usual, but More Unequal.....	163
Meet Lena, 18	163
Meet Daniel, 40.....	163
The result	164
Future 2: “Platform World” – Corporations and AI Systems Lead	164
Meet Sofia, 18.....	164
Meet Marco, 40	165
The result	165
Future 3: “Redesign” – A Public-Plus-Platform System	166
Meet Amir, 18.....	166
Meet Priya, 40	166
The result	167
Redesign’s Harder Path	167
What Pushes Us Where: Understanding the Feedback Loops	168

Loop 1: Accreditation → Institutional Structure → Program Design	168
Loop 2: Hiring Practices → Credential Demand → Institutional Investment	168
Loop 3: Funding Formula → Institutional Priority → Learner Access	169
Loop 4: EdTech Market Dynamics → Platform Consolidation → Institutional Lock-in	169
Loop 5: Stratification → Credibility Divergence → Legitimacy Erosion	169
The Three Futures Through the Five Loops.....	169
Drift: All Loops Tighten	169
Platform World: Three Loops Hijacked, Two Tighten	170
Redesign: Loops are Intentionally Rebalanced	170
What Pushes Us Where (Original Framework)	170
What pushes toward Drift.....	170
What pushes toward Platform World.....	170
What pushes toward Redesign	171
Choice, Not Fate	171
Methodological Note.....	172
Evidence Hierarchy.....	172
Tier 1: Meta-analyses and large-scale randomised controlled trials (RCTs).....	172
Tier 2: Quasi-experimental studies and multi-site evaluations.....	172
Tier 3: Single-institution pilots, case studies, and observational data	172
Tier 4: Vendor-reported data, testimonials, and marketing claims.....	173
How to Read the Evidence in This Book.....	173
The Diagnosis–Prescription Evidence Gap	173
Geographic Scope and Its Limits	174
Failure Modes — How the Hybrid Future Goes Wrong	175
Failure Mode 1: “AI on Top of Old Structures”	175
Failure Mode 2: Platform Capture of Learning Pathways	175
Failure Mode 3: The Agency Divide Hardens into a Caste System	175
Failure Mode 4: Credential Chaos and Trust Collapse.....	176
Failure Mode 5: Education Reduced to Continuous Reskilling.....	176
Failure Mode 6: Governance Paralysis by Hybrid Complexity.....	176
Failure Mode 7: Academic Freedom Erodes by Optimization.....	177

A Normative Manifesto for Education in the Age of AI

Failure Mode 8: Human Roles Hollow Out.....	177
Failure Mode 9: Public Trust Collapses Quietly.....	177
The meta-failure: Mistaking inevitability for design	177
References (incomplete)	179
Additional Reading.....	182
Further Reading.....	182
Introductory Material.....	183
Academic Sources	183
Tools and Commercial Offerings	185
Divergent and Opposing Views	186
Funding Models by Countries	187
Main sources of funding.....	187
Germany.....	187
Germany – rough structure of tertiary financing.....	187
France	188
France – rough structure of tertiary financing	188
United Kingdom (especially England).....	188
United Kingdom – rough structure of tertiary funding	188
USA	189
USA – rough structure of tertiary financing	189
Europe as a whole – pattern	189

A Normative Manifesto for Education in the Age of AI

Education is not broken because it has failed to keep up with technology.

It is broken because it has lost clarity about what it is *for*.

This book starts from a simple claim: **education exists to develop human agency and the cohesion of society under conditions of uncertainty**. Everything else—credentials, skills, employability, innovation, even growth—are secondary outcomes, not first principles.

In an age where intelligence can be manufactured at scale, the purpose of education cannot be the efficient transmission of information. That problem has been solved. The purpose must instead be to cultivate the capacities that cannot be automated: **judgment, responsibility, agency, and the ability to act wisely in complex, contested, and changing worlds**.

From this follow five non-negotiable commitments.

1. Judgement over recall

When answers are abundant, the scarce resource is discernment.

Education must shift its center of gravity from memorization to judgment: the ability to frame good questions, evaluate competing claims, integrate evidence, and decide under uncertainty. Systems that optimize for test performance, content coverage, or credential throughput at the expense of judgment are failing their core mission—even if they appear efficient.

AI should be used to *support* judgment, not replace it.

2. Agency before access

Access without agency is not empowerment—it is exposure.

Expanding access to learning tools is necessary but insufficient. Education systems must actively **build the capacity to learn**, not merely provide resources. This includes self-regulation, metacognition, ethical reasoning, and the courage to take responsibility for one's learning trajectory.

Any system that scales content faster than it builds agency will increase inequality, not reduce it.

3. Education is about Navigating Lives

Labour-market relevance matters—but it is not enough.

Education must prepare people not only to *get jobs*, but to **navigate lives**: to adapt when roles change, to resist manipulation, to participate in democratic societies, and to exercise moral judgment when incentives point the wrong way.

Reducing education to employability metrics hollowes out the public justification for universities and invites corporate capture of learning pathways.

4. Institutions matter because power matters

Markets optimise for efficiency, not for fairness. Left unchecked, they concentrate advantage, externalise social costs, and erode the public goods – equal opportunity, social cohesion, democratic participation – that no individual market actor is incentivised to protect. This is not a flaw in markets; it is what markets do.

The corrective has always been institutional: governments, regulators, and public bodies that set boundaries, redistribute access, and hold the line on values that markets cannot price. When those institutions retreat or fail to keep pace, market logic fills the vacuum – and the common interest loses its seat at the table.

In this sense, Universities are not merely service providers. They are public trust institutions responsible for truth-seeking, academic freedom, and the preservation of cultural and scientific memory. These functions require governance, slowness, and resistance to short-term optimization.

AI systems should inform institutional judgment—not replace it. Decisions about what counts as knowledge, competence, and merit must remain human, accountable, and contestable.

5. Trade-offs must be made explicit

There is no neutral redesign.

Every educational reform privileges some values over others: speed over deliberation, flexibility over stability, personalization over shared standards. Pretending otherwise is dishonest.

This book rejects the fantasy of painless transformation. Some traditions must be abandoned. Some efficiencies must be refused. Some innovations must be constrained. The measure of success is not disruption, but **whether the redesigned system produces people capable of using powerful tools without becoming tools themselves.**

The anchor

If education has a single obligation in the age of AI, it is this:

To ensure that as intelligence becomes cheaper, human judgment and agency becomes stronger.

Everything that follows in this book—its analysis, its warnings, its proposed futures—should be read against that standard.

What Exactly Is Broken?

Everyone seems to agree that education is broken. Politicians say it. Employers say it. Students feel it. Parents worry about it. But when you press people to explain what “broken” actually means, the answers scatter in every direction. Some point to cost. Others point to irrelevance. Others still talk about fairness, or mental health, or the fact that a four-year degree no longer guarantees anything it once promised.

If you are picking up this book, you have probably felt it already: something about education no longer fits the world we live in. A degree costs more than a car, new skills go out of date in a few years, and many students and workers feel anxious, burnt out, or simply lost.

At the same time, extraordinary new possibilities have appeared. A teenager in a small town can learn advanced math from an AI tutor at midnight. A nurse can take an online micro-course and qualify for new responsibilities at work. A 40-year-old programmer can retrain for a new role using employer-funded learning, guided by recommendation systems that feel closer to Netflix than to school.

So is education broken – or is it being rebuilt? The honest answer is: both. The old structures are under strain, but the new ones are uneven, confusing, and sometimes unfair.

This chapter tries to pin the problem down. Before we can talk about solutions – and this book is full of them, some promising, some dangerous, some both – we need a shared understanding of what is actually failing, for whom, and why.

Four dimensions of brokenness

When we say education is broken, we are really talking about four things at once. Each one matters on its own, but together they form a system that is failing on multiple fronts simultaneously.

1. Learning outcomes

The most basic question is: are people actually learning? The evidence is not reassuring.

Globally, the World Bank, UNESCO, and UNICEF have documented what they call “learning poverty” – children and young people who attend school but cannot read, write, or do basic math at the level expected for their age. This is not a problem confined to the developing world. In wealthy countries, standardised assessments consistently show stagnant or declining performance in foundational skills, even as spending per student has risen.

In higher education, the picture is more nuanced but still troubling. Students increasingly view themselves as consumers who feel entitled to high grades by virtue of paying tuition – a mindset that research has linked to lower academic performance and weaker critical thinking skills. Grade inflation, where institutions award higher marks without corresponding improvements in learning, has become widespread. Teachers

feel pressured to inflate grades to avoid negative student evaluations, and administrators are sometimes incentivised to manipulate statistics to maintain funding rather than to improve actual education.

Meanwhile, standardised testing – the main tool governments use to measure learning at scale – is itself contested. Critics argue that state-mandated assessments measure narrow abilities rather than genuine cognitive growth. If the tests are flawed, the data they produce is flawed, and the policies built on that data are built on sand.

The result is a system that produces credentials without reliably producing competence. A degree says you finished something. It does not always say you learned something.

2. Equity

Education is supposed to be the great equaliser. In practice, it often reinforces the gaps it claims to close.

Affordability challenges and high debt burdens disproportionately affect lower-income and marginalised students, undermining equity and social mobility goals (OECD 2025a). In Europe, rapid expansion of tertiary attainment has improved overall participation rates, but persistent inequalities by socio-economic status, migrant background, and region continue to challenge the promise of fair access in both public and private institutions (OECD 2025b).

The digital divide adds another layer. During and after the COVID-19 pandemic, poorer households with limited connectivity saw their children fall further behind. The loss was not temporary. Without intervention, those gaps compound over time, narrowing future opportunities for the students who could least afford to lose ground.

Neurodivergent learners remain underserved by mainstream systems that demand conformity rather than providing specialised support. These students often need environments that allow for what researchers call “unmasking” – the ability to learn in ways that match how their minds actually work, rather than forcing adaptation to a system designed for a narrow definition of normal.

And as AI-powered learning tools become more available, a new kind of inequality may be emerging – not based on access to information, but on the agency and discipline to use it. When learning is abundant, the advantage is likely to go to those with the self-direction to build and iterate, rather than to be passively entertained or overwhelmed. If this pattern holds, the divide will fall between those who use these tools as thought partners and those who use them to avoid thinking altogether.

3. Cost

A degree now costs more than a car in many countries. In the United States, student debt has become a defining feature of an entire generation’s economic life. In Europe, while tuition for domestic students is often lower or even free at public universities, the real costs are still rising: living expenses, opportunity costs, and the growing reliance on fees charged to non-EU students to balance institutional budgets (EHEA 2024).

Behind the scenes, institutions face their own financial crisis. Shrinking public funding, rising operating costs, and dependence on tuition make financial sustainability a central risk, especially for small and mid-tier institutions (Goedmo 2025). Growing costs for healthcare, mental health services, compliance, and infrastructure – labs, housing, digital systems – outpace available revenue sources (Deloitte 2025).

Many universities respond with tuition hikes, programme cuts, staff reductions, or even closures, which can further erode access and public confidence (Changing Higher Ed 2025). This creates a vicious cycle: costs rise, public trust falls, enrolment drops, revenues shrink, and institutions cut further. The small and mid-tier schools caught in this spiral are the ones that often serve the students who most need affordable education.

The cost problem is not just about price. It is about whether the price is justified by what you get in return. And for a growing number of students and families, the answer is no longer obvious.

4. Labour-market fit

Students and families increasingly question the return on investment of a degree, given high tuition, student debt, and uncertain job outcomes (KPMG 2020). And their scepticism is grounded in reality.

Employers report persistent skills gaps, especially in digital, STEM, and transversal skills – communication, teamwork, adaptability – that curricula often fail to deliver quickly enough (OECD 2025a). Slow academic governance can delay curriculum updates, making programmes feel outdated compared to fast-moving labour-market demands and alternative credentials (ASU and UDI 2024).

The half-life of a skill is shrinking. A four-year fixed curriculum is structurally outdated before a student even finishes it. What was cutting-edge when a student enrolled may be standard practice – or obsolete – by the time they graduate. This is especially true in technology-adjacent fields, but it increasingly applies across the board. Healthcare protocols change. Regulatory frameworks evolve. Business models shift. The pace of change has outrun the pace of curriculum development.

Meanwhile, companies are shifting their valuation of education. They increasingly look for demonstrated capability – portfolios, projects, and evidence of value creation – rather than relying solely on degrees as proxies for competence. Companies do not hire degrees; they hire the signal that a candidate can think, produce value, and finish hard things. Because AI is making it cheaper to verify skills through projects, simulations, and portfolios, the degree appears to be shifting from a gate to a preference – though the pace and extent of this shift remain uncertain.

The question is moving from “What school did you go to?” to “What can you actually do?” That is a profound shift, and the educational system has not caught up with it.

Where this book focuses – and where it does not

This book is primarily about two domains: higher education and corporate learning. These are the places where the brokenness is most visible, where the most money is at

stake, and where the most dramatic experiments in reform are underway. We also spend significant time on recruiting and hiring, because that is where education meets the labour market and where credentials are tested against real jobs.

We are not writing a book about primary and secondary schooling, though K-12 problems are deeply relevant. Many of the failures in higher education – remedial classes, disengaged students, shallow preparation – trace back to what happened (or did not happen) in earlier years. The global crisis of learning poverty, documented by the World Bank and its partners, is fundamentally a K-12 problem that casts a long shadow over everything that follows. We will reference it where it matters, but a thorough treatment of school systems would require a different book.

Similarly, while we discuss AI and technology throughout, this is not a technology book. It is a book about people and institutions trying to adapt to a world that is changing faster than their structures allow.

What is our actual claim?

Before we move on, it is worth stating plainly what we think is actually broken, and what we think is actually possible.

When we say education is broken, we do not mean that the entire system is beyond repair. Rather, we mean that the foundational bargains that held the old system together – knowledge is scarce and valuable, institutions control access to that knowledge, a degree signals competence and opens doors – have deteriorated faster than institutions have been able to adapt. The system was designed for a world where information was expensive and slow; it is now operating in a world where information is abundant and fast. That mismatch creates failures across every dimension we have described. But mismatch is not permanent. It can be addressed.

Our argument has three parts. **First:** the brokenness is real and structural, not merely cyclical or superficial. The problems in learning outcomes, equity, cost, and labour-market fit are not the result of bad leadership at individual institutions; they flow from the institutions themselves being designed for conditions that no longer exist. **Second:** the brokenness is not uniform. It falls differently on different people, different institutions, and different geographies. A failure that is systemic in one context can be an opportunity in another. **Third, and most importantly:** the brokenness is not destiny. But fixing it requires more than tinkering with curriculum or adopting new technology. It requires institutions and societies to make explicit choices about what education is for, and to align structures, incentives, and resources with those choices. Some institutions will make those choices and adapt. Others will not. Some will adapt to serve their existing students and communities better; others will adapt to serve new markets and demographics. These are not neutral moves. They have winners and losers.

This book is about the choices ahead – what they are, who is making them, what they might lead to, and what is at stake. We focus on three domains: higher education, workplace learning, and hiring. Within these domains, we concentrate on OECD and European contexts, with particular attention to the United States and Europe, where the forces of change are most visible and where the most capital (financial and cultural) is invested in education. Our analysis applies most directly to tertiary education and

professional-level corporate learning, though the underlying patterns often trace back to K-12 failures and ripple outward into lifelong learning. We do not try to fix everything at once. We try to diagnose what is actually broken, show what is actually changing, and help readers – whether they are students, parents, educators, administrators, policymakers, or employers – think more clearly about how to act in a system in transition.

The overall claim is this: education is broken because the structures are out of sync with the world. The system can be repaired, but not through incremental change. It requires explicit choice, disciplined execution, and genuine acceptance of trade-offs. Some of what is happening now is progress. Some is predatory. Most is somewhere in between. Our job is to help readers tell the difference, and to understand what is at stake when they choose which direction to move.

Who feels the brokenness: seven groups, seven experiences

The problems we have described do not land equally on everyone. Education involves at least seven distinct groups of stakeholders, and each one experiences the system's failures differently. Understanding these different perspectives is essential, because a reform that helps one group can easily harm another.

Students

Students are at the centre of the system, but they are often the least powerful actors within it. They are currently processed in batches – taught at a speed that bores top performers and leaves struggling learners behind – and assessed through static grades that provide a snapshot of performance rather than a picture of growth.

In higher education, students face a painful bind. They pay more than ever for degrees whose labour-market value is increasingly uncertain. Many experience financial stress, mental health difficulties, and disengagement, all of which increase dropout risk (SeatsOne 2025).

Specific populations are hit harder. Neurodivergent learners are critical stakeholders often underserved by mandatory integration in mainstream schools, requiring environments that allow for specialised support rather than forced adaptation. First-generation students, students from low-income families, and students with migrant backgrounds face compounding barriers that the system was never designed to address.

The consumer mindset – “I paid for this, so I deserve a good grade” – is itself a symptom of brokenness. When students see education primarily as a transaction, something has already gone wrong with the value proposition. But it is hard to blame them. When a degree costs what a house deposit used to cost, treating it as a product is a rational response to an irrational system.

Teachers and educators

Teachers are asked to do more with less, and the strain is showing. The role is shifting from being a broadcaster of information to being a designer of learning experiences,

acting as mentors, coaches, and guardians of standards. This is a more demanding job, not a less demanding one.

Yet educators face high rates of burnout characterised by emotional exhaustion and depersonalisation, driven by rigorous schedules, lack of resources, and the pressure to manage behavioural problems. In many systems, teacher pay has not kept pace with the demands being placed on the profession, making recruitment and retention increasingly difficult.

Teachers also face the dilemma of grade inflation – often pressured to award higher grades to placate students and avoid retaliation in evaluations. This undermines academic standards and places teachers in an impossible position: uphold rigour and face complaints, or lower the bar and maintain peace. The result is a profession caught between what education should be and what institutional incentives force it to become.

In Europe, universities must invest in staff development for digital and hybrid teaching, but many lack robust, publicly funded bodies to support pedagogy and staff competences at scale (Eurydice 2024). Teachers are expected to master new tools, adapt to hybrid formats, and integrate AI into their practice, while their workload increases and their autonomy shrinks. Universities must upgrade legacy systems and ensure interoperability across borders, all while avoiding the increased workload and burnout that these transitions often cause (EUA 2024a).

Educational institutions and administrators

Traditional universities face an existential question. For generations, the university business model relied on bundling three products: access to knowledge, access to talent networks, and access to a credential trusted by the market. That model thrived when knowledge was scarce, slow, and expensive. Today, all three components are being unbundled by technology, competition, and changing expectations.

Institutions that rely heavily on domestic, full-time, residential undergraduates are particularly exposed to demographic and market changes (Goedmo 2025). In several regions – the United States, parts of Europe, East Asia – fewer traditional college-aged students and more competition from online providers and bootcamps are driving enrolment decline (Hanover Research 2025). Retention is also challenged by students' financial stress, mental health issues, and disengagement, which increases dropout risk (SeatsOne 2025).

In Europe specifically, public universities still rely heavily on state funding, but stagnant or shrinking budgets and more competitive allocation schemes force them to seek private income, third-party projects, and higher fees for non-EU students (EHEA 2024). Private universities are expanding enrolment, especially at master's level, but depend more on tuition, philanthropy, and corporate partnerships, which increases their exposure to market downturns and raises concerns about equity and mission drift (OECD 2025b). Both sectors face rising costs for research, digital infrastructure, and student services while trying to keep tuition politically and socially acceptable (EHEA 2024).

Administrators are caught in the middle. Some are incentivised to manipulate statistics to maintain funding rather than focusing on actual education. Others are trying to

innovate but are blocked by rigid regulatory environments, accreditation rules, and complex internal governance structures that slow experimentation with new models – micro-credentials, stackable degrees, work-integrated learning (ASU and UDI 2024).

Institutions also struggle to modernise legacy systems, integrate learning platforms, and deliver flexible, hybrid, or online learning at scale without fragmenting the student experience (Full Fabric 2024). Rapid adoption of AI and analytics brings opportunities for personalisation and efficiency but raises ethical concerns about data privacy, algorithmic bias, academic integrity, and potential displacement of some academic or administrative roles. Cybersecurity threats and privacy breaches are now routine risks, pushing universities to invest heavily in protection and compliance (KPMG 2020).

We argue that the institutions that survive will need to become something different: not content-delivery machines, but high-trust environments for deep work, identity formation, and character development. They can no longer justify their existence by controlling access to knowledge. Instead, they will need to offer something AI cannot: the crucible of real problems, real stakes, and real human relationships that forge judgment and resilience. If this analysis is correct, the ones that cannot make this transition risk becoming irrelevant.

Employers and corporations

The corporate sector is both a customer of education and an increasingly active participant in it. Employers need people who can think, produce value, and finish hard things. For decades, the degree served as a convenient proxy for these qualities. That proxy is weakening.

Companies report persistent gaps between what graduates know and what the workplace demands. This is especially true for digital and STEM skills, but also for transversal capabilities – communication, teamwork, adaptability – that traditional curricula struggle to teach at all, let alone quickly enough to match the pace of change.

The result is that employers are building their own learning systems: AI-driven training platforms, employer-sponsored tuition programmes, internal skill-tracking tools, and recommendation engines that suggest learning paths tailored to each employee. Education is being pulled into the workplace, sometimes to the benefit of workers, sometimes to the benefit of employers' bottom lines, and often in ways that blur the line between the two.

At the same time, corporations are stakeholders in educational policy. Testing companies like Pearson are criticised for having significant influence over public education policy to protect their bottom line. Education-technology firms shape what gets taught and how it gets measured. This creates tension between educational goals and commercial interests – a tension that deserves far more scrutiny than it currently receives.

Policymakers and government bodies

Governments fund education, regulate it, and set the rules that determine what counts as a legitimate credential. They are also responsible for some of the system's deepest problems.

Policymakers are criticised for politicising education and for failing to give teachers a meaningful seat at the table regarding curriculum and standards. The decisions about what students learn and how they are assessed are often made far from the classroom, by people whose incentives are shaped by electoral cycles rather than educational outcomes.

State-mandated standardised testing is a persistent point of contention, with critics arguing that these mandates measure narrow abilities rather than genuine cognitive growth. Yet without some form of standardised measurement, it is difficult to diagnose problems or hold systems accountable. The tension between accountability and meaningful assessment is one of the oldest in education, and it remains unresolved.

To address crises like global learning poverty, high-level national political commitment is required to prioritise foundational skills and implement recovery plans. International organisations advocate for evidence-based strategies, such as the RAPID framework developed by the World Bank, UNESCO, and UNICEF. But political will is unevenly distributed, and short-term pressures often crowd out long-term investment in education.

In Europe, several systems report growing political interference, including contested programme areas (such as gender studies), constraints on institutional autonomy, and new “research security” regimes that tighten external collaboration and data control (University World News 2025). Governments increasingly tie funding to performance contracts and strategic priorities, which can undermine institutional autonomy and long-term basic research capacity if not carefully designed (EUA 2024b). In some countries, austerity measures and policy volatility undermine universities’ ability to plan strategically at all (EUA 2024b).

Families and communities

Parents and families are essential stakeholders, but their ability to support education varies enormously depending on socio-economic circumstances.

In the context of learning recovery following the COVID-19 pandemic, successful acceleration of learning requires what researchers have called a “national coalition” that includes families engaged in helping their children learn. But the digital divide and lack of connectivity mean that poorer households are least able to provide this support, and their children’s learning losses are the hardest to recover.

Families also face a cultural shift that can be deeply disorienting. For generations, the path was clear: get into the best school you can, get the degree, get the job. That script gave parents a role they understood. Today, they are being asked to evaluate a bewildering landscape of alternatives – online programmes, bootcamps, micro-credentials, apprenticeships, campus-free institutions – without clear information about which ones are worth the investment, which ones are legitimate, and which ones are simply good marketing.

The question is shifting from “What school do you go to?” to “What are you building?” – but most families do not yet have the framework to answer it. This is not a failure of individual parents. It is a failure of the system to provide transparent, trustworthy information about pathways that genuinely lead to good outcomes.

International organisations

On a global scale, organisations such as the World Bank, UNESCO, and UNICEF work to address learning poverty and the loss of human capital. They advocate for evidence-based strategies, such as the RAPID framework, to recover learning losses and ensure that all children acquire foundational literacy. Their role is to set norms, gather data, and push governments to act – but they have limited power to enforce anything.

These organisations matter because education is not just a national concern. The skills gaps, equity failures, and credential crises described in this chapter play out across borders. Students move between countries. Employers hire globally. Credentials are evaluated (or not) across systems with different standards and assumptions.

European university alliances are trying to create new forms of cross-border cooperation and blended mobility, but they must navigate tighter visa rules, geopolitical tensions, and research-security scrutiny of partnerships with certain countries (Verfassungsblog 2024). Competition for international students is intensifying, with fee-paying non-EU students becoming central to many institutions' business models, especially in the Netherlands, the Nordics, France, Italy, and Spain (EUF 2024).

Shifting geopolitics within the European Higher Education Area also challenge the coherence of Bologna norms, particularly around academic values, mobility, and mutual recognition (Taylor and Francis 2024a). The international architecture of education – built over decades through agreements like the Bologna Process – is under strain from the same forces that are disrupting national systems.

The European dimension

Many of the problems described in this chapter are global, but Europe faces its own particular version of them, shaped by its distinctive mix of public funding, political diversity, and cross-border ambition.

European public universities operate within a tension that has no easy resolution. They are committed, in principle, to universal access and low or zero tuition for domestic and EU students. But stagnant public budgets, rising costs for research, digital infrastructure, and student services make this commitment increasingly difficult to sustain (EHEA 2024). Both public and private institutions face questions about the quality and recognition of online, blended, and micro-credential offerings within Bologna structures and national quality-assurance regimes (Eurydice 2024).

Europe's ageing populations and regional disparities create a demographic puzzle. Some systems struggle with declining local cohorts while others manage massification and capacity constraints, often within tight public spending envelopes (OECD 2025c). Fee regimes that charge higher tuition to non-EU students have become an important revenue source but create tensions between income needs, diversification goals, and equity narratives (OECD 2025b).

Digitalisation is a central reform theme across the continent, but many systems lack the infrastructure, staff training, and interoperability frameworks needed to deliver high-quality hybrid teaching at scale. The EHEA reports that digital transformation is a priority, but many systems lack robust, publicly funded bodies to support pedagogy,

staff development, and high-quality hybrid teaching (Eurydice 2024). Universities must upgrade legacy systems, ensure cross-border compatibility for micro-credentials and virtual mobility, and invest in staff competences – all while avoiding the increased workload and burnout that these transitions often cause (EUA 2024a).

Political pressures add a further layer. Academic freedom is being tested in several European systems, with growing political interference in programme areas, constraints on institutional autonomy, and new research-security regimes (University World News 2025). The relationship between universities, governments, and society is being renegotiated in real time, and the outcome is far from certain.

The overall picture is one of a system that has expanded access and raised attainment dramatically over the past generation, but that now faces structural pressures – financial, demographic, political, technological – that its existing governance models struggle to absorb. Global demand for higher education is projected to rise significantly by 2030 (OECD 2025a), even as many European institutions are already stretched thin. Underfunded institutions and entrenched bureaucracies make large-scale redesign hard, even when everyone agrees it is necessary (OECD 2025a).

The governance trap

One theme deserves special emphasis because it runs through every problem described above: institutional inertia.

Many universities face a fundamental misalignment between long-established academic structures and the need for agile, cross-sector collaboration with industry, government, and alternative providers (Goedmo 2025). Rigid regulatory environments, accreditation rules, and complex internal governance structures can slow experimentation with new models to a crawl (ASU and UDI 2024).

This is not because the people inside these institutions are complacent or unimaginative. Many of them see the problems clearly and want to act. But universities are designed for stability, not speed. Their governance structures – shared authority between faculty, administrators, boards, accreditors, and government regulators – make rapid change extraordinarily difficult. A curriculum update that the labour market demands in months can take years to navigate through committee structures.

The governance trap is one reason why so much educational innovation is happening outside traditional institutions – in startups, corporate learning departments, and alternative credentialing systems. Whether that is a good thing or a dangerous one is a question this book takes seriously. Innovation that bypasses established quality controls can expand opportunity, but it can also create a Wild West of unregulated credentials, predatory programmes, and systems that benefit the already-privileged while leaving everyone else more confused than before.

Evidence vs Conjecture: AI Learning Inequality

The claim:

As AI-powered learning tools become more available, a new kind of inequality may be emerging based not on access to tools, but on the agency and discipline to use them effectively as thought partners rather than avenues for passive consumption.

What supports it:

- The disparity in digital skill gaps by socio-economic status is well-documented (Tier 2: research on digital divides and learning outcomes, OECD 2025a).
- The observation that effective learning with AI tools requires self-direction and iterative practice is grounded in adult learning research and cognitive science.
- The “use as thought partner vs. passive entertainment” distinction is a hypothesis about future behaviour and differential adoption, not yet observed at scale across populations.

What would change it:

- If longitudinal studies of AI tool adoption showed that students across socio-economic backgrounds develop similar patterns of use and engagement, this claim would need revision.
- If evidence demonstrates that AI systems effectively scaffold learning for lower-agency users (through better design, guardrails, or adaptive feedback), the inequality risk would diminish.
- If policy or institutional interventions successfully normalize deliberate, high-engagement use across populations, this trajectory could be interrupted.

Confidence level:

The foundational concern about self-direction and equity is well-grounded in existing research on learning disparities. The specific hypothesis about AI-powered inequality as distinct from traditional access inequality is plausible but not yet empirically established at scale. This is a pattern to monitor rather than a settled fact.

What comes next

This book follows the problem through three domains where education meets reality. Part I examines higher education: the traditional university model, the forces unbundling it, and the new institutions and formats – competency-based programmes, campus-free models, peer-driven coding schools, European university alliances, micro-credentials – that are trying to replace or reinvent it. We also take a hard look at the pitfalls: credential inflation, algorithmic bias, stratification between elite and non-elite institutions, and the risk that exciting experiments mainly benefit already-privileged learners. Part II moves into the workplace, where employers are building their own learning systems, sponsoring tuition, and using AI to track and develop skills, raising new questions about who controls what workers learn and how their data is used. Part III looks at hiring and recruiting, where AI tools scan CVs, analyse portfolios, run simulations, and even attempt to read emotions in interviews, and where the shift toward skills-based hiring is uneven, messy, and far from complete. Across all three parts, one question keeps returning: who benefits, who is excluded, and who decides.

Annotated Source List

1. **OECD 2025a** — *Education at a Glance 2025*. Comprehensive annual compendium of cross-national education indicators covering access, completion, spending, and labour-market outcomes across OECD and partner countries. Produced by the OECD Directorate for Education and Skills, an intergovernmental research body. **Tier 2** (multi-country comparative data with standardised methodology).
2. **OECD 2025b** — “How Is Tertiary Education Financed?” Thematic chapter within *Education at a Glance 2025* analysing tuition levels, public subsidies, student debt, and household spending on higher education across OECD systems. Produced by the OECD. **Tier 2** (cross-national policy analysis drawing on standardised fiscal data).
3. **EHEA 2024** — *The European Higher Education Area in 2024*. Stocktaking report assessing progress on Bologna Process commitments including quality assurance, recognition, social dimension, and learning and teaching reforms across 49 signatory countries. Produced by the EHEA/Bologna Follow-Up Group. **Tier 2** (multi-country evaluation based on national self-reports and Eurydice data).
4. **Goedmo 2025** — “Problems in Higher Education.” Blog post summarising common critiques of higher education including cost, relevance, mental health, and pedagogical stagnation. Published by Goedmo, an education technology company. **Tier 4** (vendor blog without primary data or peer review).
5. **Deloitte 2025** — *2025 US Higher Education Trends*. Industry outlook report identifying financial sustainability, enrolment shifts, workforce alignment, and technology adoption as key pressures on US colleges and universities. Produced by Deloitte, a professional services consultancy. **Tier 3** (single-country industry analysis combining survey data and expert opinion).
6. **Changing Higher Ed 2025** — “Higher Ed Challenges 2025: Solution Examples.” Blog-style article listing institutional challenges and vendor-adjacent solutions for US higher education leaders. Published by Changing Higher Ed, an education consulting platform. **Tier 4** (opinion-based content without primary research).
7. **KPMG 2020** — *Future of Higher Education*. White paper forecasting structural shifts in global higher education including digitisation, new credentialing, and business-model disruption. Produced by KPMG, a professional services firm. **Tier 3** (consultancy forecast based on industry interviews and secondary data).
8. **ASU and UDI 2024** — “Not If But How: Redesigning the Future of Higher Education.” Institutional essay arguing for large-scale redesign of university structures, pedagogy, and access models. Published by Arizona State University and its University Design Institute. **Tier 3** (single-institution perspective from a major research university).
9. **SeatsOne 2025** — “The Ultimate Guide for Universities in 2025.” Marketing-oriented article listing challenges and AI-powered solutions for university

operations. Published by SeatsOne, an AI scheduling software vendor. **Tier 4** (vendor marketing content without independent evidence).

10. **Eurydice 2024** — “Chapter 5: Learning and Teaching.” Section of the EHEA stocktaking report analysing student-centred learning, digital pedagogy, and teaching quality across European systems, drawing on institutional surveys and national reports. Produced by the Eurydice network under the European Commission. **Tier 2** (multi-country comparative analysis using systematic data collection).
11. **EUA 2024a** — *Trends 2024*. Flagship survey report on governance, funding, teaching, research, and digitalisation trends across European universities, based on responses from over 700 institutions. Produced by the European University Association. **Tier 2** (large-scale multi-institution survey with established methodology).
12. **Hanover Research 2025** — “5 Ways Higher Education Is Changing in 2025-26.” Brief trends report aimed at institutional decision-makers covering enrolment, workforce development, AI adoption, and student success strategies. Published by Hanover Research, an institutional analytics firm. **Tier 3** (consultancy brief combining survey data and secondary analysis).
13. **Full Fabric 2024** — “Challenges in Higher Education Management.” Article outlining operational and strategic challenges facing university administrators including enrolment management, digital transformation, and student retention. Published by Full Fabric, a CRM software provider for universities. **Tier 4** (vendor content marketing without independent data).
14. **University World News 2025** — “Higher Education Funding and Policy Developments.” News article reporting on recent government policy changes and funding decisions affecting universities globally. Published by University World News, an independent higher-education news outlet. **Tier 3** (journalism based on public records and expert commentary).
15. **EUA 2024b** — *Designing Strategies for Efficient Funding of Universities in Europe (DEFINE)*. Research project report analysing funding models, efficiency, and sustainability across European university systems, based on institutional case studies and financial data. Published by the European University Association. **Tier 2** (multi-site research project with systematic methodology).
16. **Verfassungsblog 2024** — “Academic Freedom and Security.” Legal-academic essay examining tensions between academic freedom, security policy, and institutional governance in European higher education. Published on Verfassungsblog, a peer-reviewed open-access law blog. **Tier 3** (expert commentary on a specialised legal platform).
17. **EUF 2024** — “Internationalisation of Higher Education: Challenges, Trends, Priorities.” Policy briefing on cross-border cooperation, student mobility, and internationalisation strategies in European higher education. Published by the European University Foundation. **Tier 3** (policy organisation briefing based on network experience and secondary data).

18. **Taylor and Francis 2024a** — “Bologna Process and European Higher Education Area Developments.” Peer-reviewed journal article analysing the evolution and impact of the Bologna Process on European higher education structures and policies. Published in a Taylor & Francis academic journal. **Tier 2** (peer-reviewed scholarly analysis).
19. **OECD 2025c** — *Trends Shaping Education 2025*. Forward-looking report examining demographic, economic, technological, and environmental megatrends and their implications for education systems worldwide. Produced by the OECD. **Tier 2** (cross-national trend analysis using OECD datasets and scenario modelling).

References

- Arizona State University and University Design Institute. 2024. “Not If But How: Redesigning the Future of Higher Education.” <https://udi.asu.edu/co-lab/article/Not-If-But-How-Redesigning-the-Future-of-Higher-Education>.
- Changing Higher Ed. 2025. “Higher Ed Challenges 2025: Solution Examples.” <https://changinghighered.com/higher-ed-challenges-2025-solution-examples/>.
- Deloitte. 2025. *2025 US Higher Education Trends*. <https://www.deloitte.com/us/en/insights/industry/public-sector/2025-us-higher-education-trends.html>.
- European Higher Education Area. 2024. *The European Higher Education Area in 2024*. <https://ehea.info/Immagini/the-european-higher-education-area-in-2024-EC0224018ENN.pdf>.
- European University Association. 2015. *Designing Strategies for Efficient Funding of Universities in Europe (DEFINE)*. https://www.eua.eu/images/designing_strategies_for_efficient_funding_of_universities_in_europe_define.pdf.
- European University Association. 2024. *Trends 2024*. <https://www.eua.eu/publications/reports/trends-2024.html>.
- European University Foundation. 2024. “Internationalisation of Higher Education: Challenges, Trends, Priorities.” <https://uni-foundation.eu/2024/02/08/internationalisation-higher-education-challenges-trends-priorities/>.
- Eurydice/EHEA. 2024. “Chapter 5: Learning and Teaching.” https://eurydice.eacea.ec.europa.eu/sites/default/files/2024-05/Chapter_5_Learning_and_teaching.pdf.
- Full Fabric. 2024. “Challenges in Higher Education Management.” <https://www.fullfabric.com/articles/challenges-in-higher-education-management>.
- Goedmo. 2025. “Problems in Higher Education.” <https://goedmo.com/blog/problems-in-higher-education/>.

Hanover Research. 2025. “5 Ways Higher Education Is Changing in 2025-26.” <https://www.hanoverresearch.com/reports-and-briefs/higher-education/5-ways-higher-education-is-changing-in-2025-26/>.

KPMG. 2020. *Future of Higher Education*. <https://assets.kpmg.com/content/dam/kpmg/xx/pdf/2020/10/future-of-higher-education.pdf>.

OECD. 2025a. *Education at a Glance 2025*. Paris: OECD Publishing. https://www.oecd.org/en/publications/education-at-a-glance-2025_1c0d9c79-en.html.

OECD. 2025b. “How Is Tertiary Education Financed?” In *Education at a Glance 2025*. Paris: OECD Publishing. https://www.oecd.org/en/publications/education-at-a-glance-2025_1c0d9c79-en/full-report/how-is-tertiary-education-financed_2845d742.html.

OECD. 2025c. *Trends Shaping Education 2025*. Paris: OECD Publishing. https://www.oecd.org/en/publications/trends-shaping-education-2025_e66587fd-en.html.

SeatsOne. 2025. “The Ultimate Guide for Universities in 2025.” <https://seatsone.ai/the-ultimate-guide-for-universities-in-2025-addressing-higher-educations-challenges/>.

Taylor & Francis. 2024. “Bologna Process and European Higher Education Area Developments.” *European Journal of Higher Education*. <https://www.tandfonline.com/doi/full/10.1080/21568235.2024.2398742>.

University World News. 2025. “Higher Education Funding and Policy Developments.” <https://www.universityworldnews.com/post.php?story=2025010808011331>.

Verfassungsblog. 2024. “Academic Freedom and Security.” <https://verfassungsblog.de/academic-freedom-security/>.

Agency: The Missing Infrastructure

Why Access Without Direction Produces Inequality, Not Opportunity

The previous chapter documented what is broken. This one names what must be built in its place.

If education exists to develop human agency — and that is the central claim of this book — then we need to be precise about what agency actually is, why it matters more now than it did a generation ago, and what it takes to build it in people who do not yet have it. Because the uncomfortable truth running through everything that follows in this book is this: every promising model we will examine works brilliantly for learners who already possess strong agency, and fails — sometimes quietly, sometimes spectacularly — for those who do not.

This is not a chapter about psychology. It is a chapter about the design constraint that every educational institution, every AI learning platform, and every policy reform must confront if it wants to do more than sort people by privilege and call it merit.

What Agency Is — and What It Is Not

Agency is the capacity to direct your own learning, set goals, persist through difficulty, and adapt when things do not work. It means the discipline to sit down and work when no one is watching. It means knowing the difference between using an AI tutor to accelerate your understanding and using it to generate an essay you never read. It means tolerating frustration, asking for help without shame, and keeping going when a problem resists easy solutions.

This sounds like a personality trait. It is not.

Agency is built. It is built through years of experience, modelling, and support. Children who grow up in stable homes with educated parents, consistent routines, access to books and mentors, and low levels of chronic stress develop the self-regulation skills that agency requires earlier and more robustly. Children who grow up in poverty, chaos, or neglect — through no fault of their own — often arrive at education with weaker self-regulation, not because they are less capable, but because their environments gave them fewer opportunities to build those muscles (Zimmerman 2002; Deci and Ryan 2000).

This distinction matters enormously for how we design educational systems.

If agency is an innate trait, then institutions that select for it are acting rationally. If agency is a socially produced capacity — and the evidence overwhelmingly says it is — then institutions that select for it are repackaging privilege as merit.

Agency also needs to be distinguished from several things it is frequently confused with.

It is not motivation. A student can be highly motivated and still lack the ability to structure their learning, manage their time, or recover from setbacks. Motivation is the desire to act. Agency is the capacity to act effectively and sustain that action over time.

It is not autonomy. Autonomy is the freedom to choose. Agency includes the competence to choose well. Giving a struggling learner unlimited freedom to design their own curriculum is not empowerment — it is abandonment.

It is not self-efficacy. Believing you can succeed is important. But belief without the skills to plan, monitor, and adjust your own learning does not produce results. A student can feel confident and still be unable to diagnose why they are stuck.

It is not competence. A person can possess skills and knowledge but lack the capacity to deploy them independently. Agency is what turns dormant competence into directed action.

What makes agency distinctive is that it integrates all of these. It is the operating system that coordinates motivation, skill, belief, and self-regulation into sustained, self-directed behaviour. Researchers in self-regulated learning have documented this for decades: effective learners cycle through planning, monitoring, and reflection repeatedly — setting goals, tracking their progress, identifying what is not working, and adjusting (Zimmerman 2002; Pintrich 2000). Agency is essentially the reliable execution of this cycle across increasingly complex challenges.

The reason this matters for the argument of this book is straightforward. Every chapter that follows describes models and technologies that assume learners can run this cycle on their own. AI tutoring platforms assume it. Competency-based education assumes it. Self-paced micro-credential systems assume it. Peer-driven coding schools assume it. And when that assumption fails — when a learner cannot plan, cannot monitor, cannot reflect, cannot adjust — the model does not adapt to the learner. The learner simply falls out.

Agency is the load-bearing wall. If it is missing, everything built on top of it collapses.

Why Agency Becomes the Bottleneck Now

Agency has always mattered. What has changed is that it has become the primary differentiator between learners who benefit from new educational possibilities and those who are harmed by them.

For most of the twentieth century, educational systems managed the agency problem through structure. Compulsory attendance, fixed schedules, prescribed curricula, teacher-led instruction, and regular deadlines created an external scaffolding that carried learners through the process whether they had strong self-direction or not. You did not need much agency to progress through a system that told you where to be, what to study, and when to submit your work. The system did the directing. The student did the complying.

This was not ideal for learning. It suppressed curiosity, rewarded compliance over initiative, and bored high-performing students while leaving struggling students too little room to develop at their own pace. But it had one underappreciated virtue: it did not require learners to bring their own agency as an entry condition.

Now consider what happens as education shifts toward the models described in this book.

AI tutoring platforms offer infinite patience, immediate feedback, and personalised pathways — but they require the learner to show up, ask the right questions, and persist when the material gets hard. There is no bell, no classroom, no teacher scanning the room for the student who has gone quiet.

Competency-based programmes let you move at your own pace — but “your own pace” for a student without self-regulation skills often means no pace at all. Western Governors University, which has served over 350,000 students through competency-based models, has struggled to produce robust evidence that CBE actually improves learning outcomes. A 2015 meta-review concluded that despite the approach emerging in the 1970s, “there is hardly any evidence for the effectiveness of competence-based education” (Morcke, Dornan, and Eika 2013). The students who succeed in these environments tend to be those who would have succeeded anyway.

Micro-credential systems offer flexibility and modularity — but navigating a landscape of stackable certificates, choosing which ones matter, and assembling them into a coherent career trajectory requires exactly the directional capacity that the most vulnerable learners lack.

The pattern is consistent. The same features that make a model exciting for high-agency learners — flexibility, self-pacing, personalisation, freedom — make it treacherous for everyone else.

We have already seen this play out. Massive Open Online Courses were supposed to democratise education a decade ago. The data tells a different story: completion rates hovered around 5–10 per cent, and the people who completed courses were overwhelmingly those who already held degrees (Reich and Ruipérez-Valiente 2019). The tool was free and open. The benefit flowed upward.

AI tutoring platforms are more sophisticated than MOOCs, but the underlying dynamic is the same. A tool that requires the user to bring direction, persistence, and critical judgment will systematically advantage those who already have these qualities — and those qualities correlate with socioeconomic background, parental education, and the quality of prior schooling.

This creates what we call the **agency divide**: the emerging inequality not between those who have access to learning and those who do not, but between those who can convert access into learning and those who cannot.

The agency divide is more insidious than the traditional digital divide because it is invisible. Two students can sit in front of the same AI tutor. One uses it to master organic chemistry in half the time. The other uses it to get through homework as fast as possible — or not at all. Both had “access.” Only one had the conditions to convert access into learning.

And AI can make this worse in a way that previous technologies could not. When a learner uses AI to avoid thinking rather than to deepen it, the tool does not resist. It complies. It generates the essay, solves the problem, produces the answer. A textbook at least required the student to read. An AI tutor, used without agency, can do the learning for you — and the learner may not even notice what they have lost.

This is the paradox that every remaining chapter in this book must confront: **more opportunity, distributed without the infrastructure to use it, becomes more inequality.**

How Agency Is Built — and What Prevents It

If agency is the bottleneck, then the central question for educational design becomes: how do you build it?

The research base here is substantial, even if educational practice has been slow to absorb it. Three decades of work on self-regulated learning — primarily by Zimmerman (2002), Pintrich (2000), and their successors — converge on a consistent finding: agency develops through repeated cycles of guided challenge, feedback, and reflection. It does not develop through instruction alone, and it does not develop through unstructured freedom.

This matters because the two most common approaches to agency in educational institutions are both wrong.

The first wrong approach is to ignore agency entirely — to treat it as the student’s problem. This is the default in most traditional systems. The curriculum is delivered, assessments are administered, and students are expected to bring whatever self-regulation they need from home. Those who have it progress. Those who do not are labelled unmotivated, underprepared, or unsuitable. The institution has done its part; the failure belongs to the learner.

The second wrong approach is to leap to full autonomy — to give students “ownership of their learning” without building the capacity that ownership requires. This is the trap that many progressive and technology-driven models fall into. They mistake the endpoint for the starting point. Self-direction is the goal of education, not its prerequisite.

What the research actually shows is that agency develops through **scaffolded progression**: a deliberate, staged movement from high structure to increasing independence.

In practice, this means something specific. A first-year student who has never managed their own learning needs structured deadlines, explicit instruction on how to break problems into manageable pieces, regular check-ins with a mentor, and clear feedback on what is and is not working. The structure is not a constraint on their development — it is the condition for it. As the student demonstrates they can plan, monitor, and adjust on their own, the scaffolding is gradually removed. By the third year, the same student might be running independent projects with minimal guidance — not because they were born self-directed, but because the institution invested in building that capacity.

This is not a new idea. Vygotsky described the zone of proximal development a century ago. Deci and Ryan’s self-determination theory (2000) established that autonomy supports learning only when accompanied by competence and relatedness — connection to others who care about your progress. What is new is the urgency. As education moves toward models that assume self-direction, the cost of failing to build it rises dramatically.

Four conditions, specifically, must be present for agency to develop.

Structured challenge with immediate feedback. Agency grows when learners face problems that are hard enough to require effort but not so hard that they produce helplessness. The feedback must be prompt enough to be useful — not a grade returned three weeks later, but a signal that lets the learner adjust in real time. AI can be genuinely powerful here, providing the kind of rapid, personalised feedback that a single teacher with thirty students cannot. But only if the learner is in a position to use that feedback, which requires the other three conditions.

Explicit metacognitive instruction. Most educational programmes treat the skills of learning itself — how to set realistic goals, how to diagnose why you are stuck, how to manage your time, how to ask for help — as if learners arrive with them fully formed. They do not. These are teachable skills, as teachable as mathematics or writing. An institution serious about building agency would dedicate real course time to metacognition — treating it with the same rigour it gives to its subject matter, not as a workshop sandwiched into orientation week (Flavell 1979; Schraw and Dennison 1994).

Human mentorship and accountability. Self-paced, algorithm-driven systems cannot notice when a learner is silently struggling. A mentor can. This does not mean surveillance or returning to traditional lectures. It means building relationships — regular check-ins where mentors know learners well enough to spot when someone is disengaging, when the self-pacing is moving at the wrong speed, when external stress is making self-direction impossible. For learners from backgrounds where self-direction is less developed, these relationships are not luxuries. They are prerequisites.

Emotional and psychological stability. A learner battling depression does not lack agency. They lack the internal bandwidth that agency requires. When nearly half of incoming students report persistent sadness or hopelessness — as a 2023 CDC survey found among U.S. high school students — any model that does not invest in mental health support is not building agency. It is filtering for learners resilient enough to overcome isolation on their own.

These four conditions are expensive. They require smaller class sizes, more instructor contact time, dedicated mentorship infrastructure, and comprehensive wellbeing services. They cost more than a lecture hall and an AI platform. And they fall disproportionately on the institutions that serve the students most likely to benefit from them — lower-income students, first-generation learners, students in under-resourced regions.

This creates the perverse dynamic at the heart of the equity problem: the students who most need agency-building support cannot attend the institutions that have the budgets to provide it.

Now consider what actively prevents agency from developing.

Most traditional educational systems do not merely fail to build agency — they suppress it. Rigid curricular pathways that leave no room for choice. Time-based progression models that move everyone at the same speed regardless of readiness. High-stakes summative assessments that punish error rather than treating it as information. Reward

systems that value compliance over initiative, correct answers over good questions, and obedience over independent thinking.

In the German context, this is particularly visible. The early tracking of students into Gymnasium, Realschule, and Hauptschule at age ten sorts children by perceived ability at an age when their self-regulation capacities are still developing — and those capacities are distributed along socioeconomic lines. A child from an educated household with books, routines, and parental engagement is more likely to demonstrate the behaviours that earn a Gymnasium recommendation. A child from a less stable background, equally intelligent but with fewer opportunities to develop self-regulation, is more likely to be tracked into a pathway that offers less autonomy, less challenge, and less opportunity to develop the agency that the system claims to reward. The tracking does not discover ability. It reflects background — and then locks the result in place.

The irony is severe. The system that most needs to build agency in students who lack it is structured to deny those students the environments where agency develops.

Design Principles for What Follows

This book examines many models: AI tutoring platforms, competency-based education, peer-driven schools, employer-sponsored learning, micro-credential systems, and various institutional hybrids. Each of them will be assessed on its own terms — what it promises, what it delivers, who it serves, and who it excludes.

But running beneath all of those assessments are three questions that this chapter has established, and that every remaining chapter must answer.

First: Does the model treat agency as an entry requirement or a learning outcome?

This is the sharpest diagnostic. A model that selects for self-directed learners and celebrates their success is not expanding opportunity — it is filtering for privilege and calling it merit. A model that deliberately builds self-direction in learners who arrive without it is doing the harder, more expensive, and more genuinely transformative work.

When we examine 42 School, Minerva, Western Governors University, and others in later chapters, this will be the first question we ask. Not “does it work?” but “for whom does it work, and what happens to everyone else?”

Second: Who bears the cost when agency is absent — the institution or the learner?

In most current systems, the answer is the learner. If you cannot self-direct, you drop out. If you drop out, you carry the debt without the credential. The institution reports a retention rate; you report a setback that may take years to recover from.

An equitable system would distribute that cost differently. It would invest upfront in the scaffolding, mentorship, and mental health support that prevent the dropout in the first place — rather than optimising for the students who need the least help and accepting the loss of everyone else as an acceptable rate of attrition.

This is expensive. The chapters on financing and institutional design will confront that cost directly. But the cost of not investing is also real — it is just borne by individuals and societies rather than appearing on institutional budgets.

Third: How does the model move learners from dependence to independence?

The endpoint of education is a person who can direct their own learning, their own career, and their own contribution to society. But the path to that endpoint cannot begin at that endpoint. It must begin with structure, guidance, and support — and it must gradually, deliberately, and measurably reduce that support as the learner develops the capacity to proceed without it.

Any model that offers only structure (traditional lecture-based education) fails because it never develops independence. Any model that offers only freedom (pure self-paced, algorithm-driven learning) fails because it assumes an independence that does not yet exist.

The models that work — and we will find some in the chapters ahead — are the ones that manage this transition deliberately: high scaffolding early, progressive autonomy over time, with clear markers for when a learner is ready for more independence and support for when they are not.

These three questions are not the only lens through which we will evaluate what follows. Cost, scalability, labour-market fit, and institutional governance all matter. But agency is the load-bearing wall. If a model does not address it, everything else is decoration.

One final point. There is a risk in making agency central that we must name honestly.

If agency becomes the new metric by which learners are judged, it can become the new mechanism of exclusion — the latest way to blame individuals for systemic failures. A student who drops out of a self-paced programme is not demonstrating a “lack of agency.” They may be demonstrating a lack of support, a lack of stability, or a system that was designed for someone else.

Similarly, highly personalised algorithmic systems can create the appearance of agency — giving learners choices, tracking their preferences, adjusting difficulty — while quietly directing behaviour in ways that serve the platform’s goals rather than the learner’s development. This is pseudo-agency: the illusion of self-direction within a system that has already decided where you are going.

And continuous self-direction without community or structure can produce its own pathology. Decision fatigue, burnout, and the anxiety of perpetual self-optimisation are not signs that a learner lacks agency. They are signs that the system has placed too much weight on individual shoulders without providing the collective infrastructure that sustained effort requires.

Agency is the goal. But it is a goal that can only be reached through institutions willing to invest in the conditions that make it possible — not through systems that demand it as the price of admission.

The next chapter examines the force that makes all of this urgent: the rise of manufacturable intelligence, and what it means for the institutions that have traditionally held the monopoly on learning.

References for This Chapter

- Deci, E.L. and Ryan, R.M. 2000. “The ‘What’ and ‘Why’ of Goal Pursuits: Human Needs and the Self-Determination of Behavior.” *Psychological Inquiry* 11(4): 227–268.
- Flavell, J.H. 1979. “Metacognition and Cognitive Monitoring: A New Area of Cognitive-Developmental Inquiry.” *American Psychologist* 34(10): 906–911.
- Morcke, A.M., Dornan, T. and Eika, B. 2013. “Outcome (Competency) Based Education: An Exploration of Its Origins, Theoretical Basis, and Empirical Evidence.” *Advances in Health Sciences Education* 18(4): 851–863.
- Pintrich, P.R. 2000. “The Role of Goal Orientation in Self-Regulated Learning.” In M. Boekaerts, P.R. Pintrich and M. Zeidner (eds.), *Handbook of Self-Regulation*. San Diego: Academic Press, 451–502.
- Reich, J. and Ruipérez-Valiente, J.A. 2019. “The MOOC Pivot.” *Science* 363(6423): 130–131.
- Schraw, G. and Dennison, R.S. 1994. “Assessing Metacognitive Awareness.” *Contemporary Educational Psychology* 19(4): 460–475.
- Zimmerman, B.J. 2002. “Becoming a Self-Regulated Learner: An Overview.” *Theory Into Practice* 41(2): 64–70.

The Rise of the Intelligence Factory

How Manufacturable Intelligence Reshapes What Universities Do

If the previous chapter mapped what is broken, this one maps what is arriving to fill the cracks. The central image is simple: for the first time in history, intelligence is becoming manufacturable. Data and electricity can be converted into tokens of useful thinking on demand. The result is an “intelligence factory” that operates 24 hours a day, speaks every language, and adapts instantly to every learner. Traditional universities are no longer competing solely with one another; they are competing with this new infrastructure. Understanding the nature of that competition – what it threatens, what it cannot touch, and what it might improve – is the purpose of this chapter.

The Intelligence Factory vs. the Traditional University

For generations, the university business model relied on bundling three distinct products into a single package:

20. **Access to knowledge** – lectures, libraries, laboratories.
21. **Access to talent networks** – peers, alumni, faculty connections.
22. **Access to a credential** – the degree, a market signal that a graduate can think, produce value, and finish hard things.

This bundle thrived in a world where knowledge was scarce, slow, and expensive. An expert standing before a lecture hall was the most efficient bandwidth solution available. Students had to move to the knowledge, because the knowledge could not move to them.

AI changes each term in that equation. Knowledge is now abundant and retrievable on demand. A personal AI tutor can explain quantum mechanics or contract law with infinite patience, instant adaptation to confusion, and – crucially – privacy from the fear of looking stupid. The lecture, once the core delivery mechanism, is exposed as an artifact of scarcity: an “ancient bandwidth solution” designed for an era when experts were the rare resource. When every student carries a tutor in a pocket, the lecture becomes among the least efficient ways to transmit understanding.

But the intelligence factory does not stop at content delivery. It reaches into assessment, feedback, scheduling, and even career guidance. Its logic is relentless: anything that can be encoded as a pattern – grading rubrics, prerequisite maps, skill-gap diagnostics – can be automated, personalized, and scaled at near-zero marginal cost.

This is the unbundling of the university. AI is separating the three products that the traditional model held together, making each one independently available, and in many cases, independently superior to what a single institution can offer.

The “intelligence factory” is a deliberately provocative metaphor. It names a real force: the relentless logic of AI-driven automation, personalisation, and scale applied to functions that universities have bundled for centuries. But the metaphor risks implying total replacement, which is not what the evidence supports. The analysis that follows

shows a fundamentally hybrid reality: AI will displace some university functions, complement others, and leave some stubbornly institutional. The sections below map where the pressure points lie – not as a prediction of replacement, but as a guide to which parts of the university bundle remain essential, which should be rethought, and where the interesting design challenge is the interface between intelligent machines and human institutions.

From Batch Processing to Continuous Optimization

The structural mechanics of education trace back to the Prussian system of the eighteenth century: compulsory schooling, standardized curricula, state certification, students processed in age-based cohorts. This is batch processing. Students move through the system in herds – taught at a speed that bores the top performers and leaves the bottom behind – and are assessed via static grades that provide a snapshot of performance at a single moment.

AI disrupts this by introducing adaptive learning that tracks trajectory rather than just status. Instead of asking “What grade did you get?” the system asks “How fast are you improving, and on what?” Education shifts from a scheduled event to a continuous operating system.

If this trajectory continues, graduation as a fixed endpoint may become an outdated concept. One plausible direction is “leveling up in real time,” where credentials become modular and skills are stacked rather than earned all at once over four years. The half-life of a skill is shrinking – estimates suggest many technical competencies become obsolete within two to five years – rendering a four-year fixed curriculum structurally outdated before a student even finishes it.

Self-paced mastery tracks replace the cohort lockstep. A student who grasps statistics in three weeks is not held back by a semester calendar; a student who needs eight weeks is not failed for being slower. The intelligence factory does not care about the clock. It cares about demonstrated understanding.

In principle, dynamic curricula could update weekly based on new tools, research, and industry demands, rather than waiting for the glacial cycle of faculty committee approvals that can take years. If such systems become widespread, the gap between what the economy needs and what the curriculum teaches – always present – would narrow dramatically. This process is constrained by feedback loops that operate at different speeds: regulatory accreditation (Loop 1), funding formulas (Loop 3), and institutional priorities all operate at different time scales, and loosening any one without loosening the others creates friction.

From Credentials to Portfolios

Universities have long relied on the degree as a defensive fortress. Employers require degrees, so students need universities, so the model sustains itself. But this logic is weakening.

Companies do not hire degrees. They hire the signal that a candidate can think, produce value, and finish hard things. Because AI makes it cheap to verify skills through projects, simulations, and portfolio analysis, the degree is shifting from a gate to a preference.

Grade inflation and a broader decline in academic rigor accelerate the erosion: when students are treated as customers entitled to high grades, the degree's informational value hollows out from within.

We are moving toward a culture of demonstrated capability. In this world, the question is not "What school did you attend?" but "What have you built?" Projects become the new transcripts. Creation is the only honest proof of competence left in a world where AI can generate standard answers on demand.

This does not mean degrees disappear overnight. Regulated professions – medicine, law, engineering – still require them, and many high-status employers still filter on institutional prestige. But the trajectory suggests that a layered "signalling stack" is emerging, with degrees as one layer among several, alongside micro-credentials, project portfolios, and AI-verified evidence of learning velocity and collaboration patterns.

The realistic time horizon matters. If current trends continue, it is plausible that within five years, portfolios and skills-based hiring will dominate in technology, design, and parts of business. Over the next twenty years, if regulatory reform keeps pace, degrees may recede across many more sectors. But for that to happen, accreditation bodies, professional associations, and insurance and liability models all have to evolve – and institutions change slowly.

The Unbundling Threshold: Network Effects and Prestige Inertia

Unbundling does not happen just because better tools exist. It happens when the incentives that hold the bundle together weaken enough to create a cascade.

Understanding when – and whether – that happens requires looking at what actually keeps universities together in the first place.

The traditional university bundle is held together by two interlocking incentives, each stronger than it appears:

Network effects. Students do not attend Harvard for the lectures alone. They attend for access to peers, faculty, and alumni who themselves attended Harvard. This creates a self-reinforcing cycle: employers hire from Harvard because Harvard graduates are disproportionately successful; talented students attend Harvard because employers hire from it; the success of graduates makes the network more valuable. The degree is the formal credential, but the network is the real product. AI tutors, no matter how good, cannot replicate this. A student learning quantum mechanics from the best AI tutor in the world is not building a relationship with someone who will sit beside her in a startup ten years later.

Prestige inertia. Even when alternatives are superior in measurable ways, institutional prestige is sticky. A degree from an Ivy League university signals not just capability but access – it certifies that you survived a gauntlet that most people did not attempt. Employers use this signal partly because it works, but also because it is legally defensible and culturally understood. When everything else is uncertain, saying "we hire from top schools" is a safe decision. Changing that decision requires not just demonstrating that portfolios work better, but accepting the social and legal risk that comes with being wrong about a person who looks good on paper but fails on the job.

These are not technical problems. They are equilibrium problems. The bundle persists not because it is optimal, but because everyone coordinated on it. For unbundling to happen, three conditions must be met:

23. **Employer confidence in alternatives must exceed employer confidence in degrees.** This is not a small gap to close. Degrees have been filtered through decades of hiring experience. Portfolios, micro-credentials, and skills assessments are newer, messier, and require more judgment to interpret. The risk of making a bad hire using a portfolio feels higher, even if the statistical likelihood is lower. This perception gap persists until the evidence becomes overwhelming and the social cost of ignoring it becomes public and visible.
24. **The value of the network must erode.** If the best people still cluster in traditional universities, then the degree retains value as a network gateway regardless of what it certifies academically. This erosion happens if: (a) the quality of peers becomes available through other mechanisms – bootcamps, online cohorts, professional associations – (b) remote work and distributed teams reduce the value of physical proximity, and (c) employers start matching candidates based on what they can do rather than where they went to school, breaking the feedback loop that sends the best students to the best universities.
25. **A critical mass must defect together.** This is the coordination problem. A single employer dropping degree requirements gains nothing if all other employers still use them as a filter. A single student skipping university gains nothing if employers still expect degrees. But if enough employers simultaneously shift, the signal travels backward: students stop incurring the cost of degrees, universities scramble to adapt, and the threshold flips. The tech industry approached this threshold in certain roles; most of the economy has not.

The realistic implication is this: unbundling will not be universal or inevitable. It will be sectoral. In technology, design, and parts of business where demonstrated work is easily verifiable and prestige-neutral (you can assess a programmer by looking at her code regardless of where she learned), unbundling can happen rapidly. The existence of alternatives – bootcamps, self-taught learning, open-source contribution – plus the rise of remote work, plus the genuine gap between what universities teach and what tech companies need, have already eroded the bundle significantly for software engineers.

In sectors where prestige inertia is stronger – management consulting, law, medicine, investment banking – the bundle will persist even as the intelligence factory grows. These sectors have built institutional moats: they hire from specific universities, those universities produce disproportionately successful recruits, and the cycle perpetuates. Breaking it would require not just that an alternative emerge, but that a critical mass of employers in the field simultaneously believe the alternative is less risky than the status quo. That is a harder coordination problem.

The bandwidth that AI creates may accelerate unbundling, but it does not guarantee it. The intelligence factory is not coming to obliterate the university. It is coming to pressure each function to justify itself on its own merits. Which functions can justify themselves will depend not on what is technically possible, but on what employers believe they can trust, and what students believe will give them access to opportunity.

Replace, Complement, or Layer?

The strongest critique of the intelligence factory metaphor is that it implies total replacement. It does not have to. The realistic future is hybrid: AI handles what it does well, institutions handle what they do well, and the interesting design work lies in the interface between them. Clarity about which functions fall where is essential.

What AI Can Realistically Replace

- **Content delivery.** The traditional lecture – one expert broadcasting to hundreds of passive listeners – is the function most vulnerable to replacement. An AI tutor can deliver the same material, personalized to the learner’s level, pace, and preferred modality, available at any hour, in any language. This is not a speculative claim; it is already happening at scale through platforms like Khan Academy’s Khanmigo, Duolingo’s AI features, and dozens of university pilots.
- **Routine assessment.** Multiple-choice tests, standardized problem sets, formulaic essay grading – any assessment that follows predictable rubrics can be automated. AI can grade faster, more consistently, and with immediate feedback. The hours faculty spend on routine grading can be redirected to higher-value activities.
- **Information retrieval and navigation.** Finding the right reading, locating a relevant paper, mapping prerequisite knowledge, answering procedural questions about course logistics – these are search-and-retrieval tasks that AI handles with ease. The reference librarian’s factual function (though not their curatorial judgment) is already largely automated.

What AI Complements

- **Tutoring and personalized feedback.** AI does not replace the tutor so much as multiply the tutor’s reach. A human tutor working with AI-generated diagnostics can focus on the moments that matter – conceptual breakthroughs, motivational crises, creative leaps – while the AI handles the routine scaffolding. The combination is more powerful than either alone.
- **Learning analytics.** AI can surface patterns invisible to the human eye: which students are falling behind, which concepts cause systematic confusion, which course sequences produce the best long-term outcomes. Faculty and administrators can use these insights to redesign curricula, allocate support resources, and intervene early. But interpreting the analytics and deciding what to do about them remains a human judgment call.
- **Personalized feedback on complex work.** AI can provide a first pass on essays, code, or design projects – flagging structural issues, suggesting references, identifying logical gaps. But the evaluative judgment about whether an argument is genuinely original, whether a design solves the right problem, or whether a student’s thinking shows real growth requires human expertise and relationship.

What Remains Institutional and Human

- **Research and knowledge creation.** Universities are not just teaching institutions; they are research institutions. The creation of new knowledge – formulating questions no one has asked, designing experiments, interpreting ambiguous results, making conceptual leaps – remains fundamentally human. AI accelerates research (better literature reviews, faster data analysis, hypothesis generation), but it does not replace the researcher's judgment, curiosity, or accountability for truth claims.
- **Socialization and community.** Learning is not purely cognitive. Students learn how to collaborate, argue, lead, fail publicly, and recover. They form friendships, professional networks, and identities. These are embodied, relational processes that require physical or at least synchronous human presence. No AI tutor replicates the experience of debating a peer at midnight, organizing a student newspaper, or navigating a team project with difficult personalities.
- **Character formation and purpose.** AI can generate options, but it cannot generate purpose or taste. The universities that survive will be those that function as "crucibles" – high-trust environments that push students into real problems with real stakes to forge resilience, ethical judgment, and a sense of vocation. In an age where tools are commoditized, having the character to wield them responsibly becomes paramount.
- **Public mission and democratic citizenship.** Universities serve a civic function that no intelligence factory can replicate: preserving cultural memory, training citizens for democratic participation, providing spaces for open debate and dissent, and anchoring regional communities. Fully outsourcing learning to corporations or platforms risks reducing education to job training, stripping away the broader purposes that justify public investment.
- **Governance and accountability.** Who decides what gets taught? Who ensures quality? Who protects academic freedom? These are institutional questions that require human governance structures – faculty senates, accreditation bodies, public oversight. An intelligence factory optimizes for engagement and completion; a university must also optimize for truth, fairness, and the long-term public interest, even when those goals conflict with efficiency.

Comparison: Three Models of Education

The following table contrasts the traditional university, the pure intelligence factory, and the AI-era hybrid that is most likely to emerge in practice.

Dimension	Traditional University	Intelligence Factory	AI-Era Hybrid
Content Delivery	Expert lectures to large cohorts; fixed schedule and pace	AI tutors deliver personalized content 24/7 in any language	AI handles routine delivery; faculty design experiences, lead seminars, and mentor
Assessment	Standardized exams, essays graded by	Continuous, adaptive assessment with	AI handles routine grading and diagnostics; faculty

Dimension	Traditional University	Intelligence Factory	AI-Era Hybrid
	faculty; periodic snapshots	instant feedback; AI-graded at scale	assess complex, creative, and ethical work
Socialization	Campus life, peer networks, clubs, residential experience	Minimal; asynchronous and individual by default	Intentional community design: project teams, cohort rituals, debate, and civic engagement
Signalling	The degree as primary credential; institutional prestige	Portfolios, micro-credentials, AI-verified skill profiles	Layered signalling stack: degree + micro-credentials + portfolio + behavioral evidence
Research	Faculty-led labs, peer review, long-term inquiry	AI-assisted literature review and data analysis; no original research agenda	AI accelerates research workflows; human researchers set agendas, interpret results, ensure integrity
Governance	Faculty senates, accreditation bodies, public boards	Platform algorithms, engagement metrics, market demand	Shared governance with AI-informed analytics; public oversight and academic freedom protections

The Intelligence Factory in Three Futures

The forces described above – unbundling, continuous optimization, platform economics, algorithmic assessment – are real and unfolding now. But they do not lead to a single outcome. How these forces mature and whom they serve depends on the institutional choices made in the next few years. We introduce three scenarios, the “Three Futures”: Drift, Platform World and Redesign.

In Drift: Universities adopt AI tools incrementally without rethinking underlying structures. Lectures get smarter screens. Assignments get plagiarism detection. But the degree remains the primary credential and the four-year lockstep persists. The intelligence factory exists, but it serves as a supplement to the traditional model rather than a replacement. Employers talk about skills but still sort resumes by degree. AI assessment tools proliferate, but without standardization, creating a confusing landscape of proprietary systems. The agency divide widens because students must navigate fragmented tools and credential stacks without institutional scaffolding. Those with the resources to hire tutors and assemble coherent learning pathways thrive; those without, drift.

In Platform World: The intelligence factory becomes the dominant infrastructure. A handful of large platform companies control the assessment, credentialing, and talent-matching systems. Universities shrink to serving elite students and conducting research. Workers experience continuous algorithmic evaluation tied to employment, with little portability or control over their own data. The optimization is ruthless and efficient for organizational productivity, but optimized for corporate needs, not learner growth or public mission. The agency divide is acute: those who game the algorithm and self-optimize thrive; those who need slower, more human-centered support are left

behind. The intelligence factory is powerful, but it serves as an instrument of corporate control rather than democratization.

In Redesign: Universities proactively use AI as a catalyst for intentional redesign. They shed functions that platforms do better (content delivery, routine grading, scheduling). They double down on functions only they can provide (research, character formation, civic mission). Micro-credentials and portfolios become standard alongside degrees, but within a quality framework that employers can trust. Lifelong learning accounts and portable credentials are backed by public investment and shared standards. AI assessment tools exist, but their use is governed by institutional frameworks accountable to learners and the public. The agency divide narrows because students have clear pathways, institutional support, and genuine choices. The intelligence factory serves as a tool of individual and public purpose, not corporate concentration.

New Skills for the AI Era

As AI proves capable of reproducing answers better than humans, we argue that the educational value stack must shift from memorization to judgment, and from compliance to curiosity. The true elite skill of the AI era, we believe, is first-principles thinking – the ability to strip away assumptions and build solutions from fundamentals rather than templates.

Two archetypes of success appear to be emerging:

- **The Questioner.** When machines can generate ten thousand answers, the human advantage lies in asking the right question – the one that reframes the problem, unlocks a breakthrough, or reveals a hidden assumption. Education must cultivate the art of questioning, not just the art of answering.
- **The Domain Elite.** As natural language becomes the primary coding interface, power shifts to those who deeply understand reality – biology, supply chains, law, materials science, public health. The goal is to produce AI-enabled experts who know the world and use the tool to reshape it, rather than AI specialists who only know the tool. A biologist who can direct an AI to design a novel protein is more valuable than a prompt engineer who knows nothing about biology.

The curriculum implications are significant. General education cannot remain a set of loosely connected distribution requirements. It must become a coherent core that builds integrated thinking: the ability to move between disciplines, connect patterns, and exercise judgment in ambiguous situations. The fragmented, cafeteria-style curriculum of many universities is poorly suited to this goal.

The Agency Divide

Perhaps the most uncomfortable implication of the intelligence factory is that while AI will make high-quality education cheaper and more available, it may increase inequality based on agency rather than access. When learning tools are abundant and free, the scarce resource is no longer information – it is the discipline and direction to use those tools productively.

The divide will be between those who use AI to avoid thinking and those who use it as a thought partner to accelerate their learning velocity. A student who asks an AI to write their essay learns nothing; a student who uses AI to critique their drafts, challenge their assumptions, and explore alternative arguments learns faster than any previous generation could.

This is not a hypothetical concern. Early evidence from AI tutoring deployments suggests that students with higher prior motivation and self-regulation benefit disproportionately, while students who most need support sometimes use AI as a crutch rather than a scaffold. If this pattern persists as AI tools scale, the agency divide risks becoming a new axis of stratification, layered on top of existing inequalities of wealth, geography, and social capital.

The design implication is clear: the hybrid model must include deliberate agency-building. Chapter 3 identifies the four conditions this requires – scaffolded challenge with feedback, explicit metacognitive instruction, human mentorship, and psychological stability – and the cost of providing them. The question is not whether AI tutoring is good or bad for equity. It is whether the institutions deploying it invest in these conditions or skip them. An intelligence factory that simply provides access without building capacity to use it will widen the gaps it promises to close.

The Realistic Future Is Hybrid

The intelligence factory is not coming to abolish the university. It is coming to unbundle it, pressuring each function to justify its existence on its own merits rather than hiding inside the bundle. Some functions – content delivery, routine assessment, information retrieval – will migrate almost entirely to AI-powered platforms, because the platforms do them better, faster, and cheaper. Other functions – research, socialization, character formation, civic mission – will remain stubbornly institutional, because they depend on human relationships, physical presence, and governance structures that algorithms cannot provide.

If this analysis is correct, the institutions that thrive will be those that accept this division of labor honestly: shedding the functions AI handles better, doubling down on the functions only they can provide, and designing the interface between human and machine with care. They will stop bragging about their lecture halls and start bragging about their crucibles – the experiences that push students into real problems with real stakes, forge resilience and ethical judgment, and produce graduates who can wield powerful tools responsibly.

The institutions that fail will be those that try to bolt AI onto old methods without rethinking the underlying model, or those that deny the shift entirely and hope prestige alone will protect them. History suggests that bundled models survive unbundling only when the bundle still creates value greater than the sum of its parts. For many universities, that case has yet to be convincingly made.

The intelligence factory is not a utopia. It carries real risks: the agency divide, corporate capture of learning pathways, erosion of the public mission, and a reduction of education to narrow job training. But it also carries a genuine promise: that the ancient dream of a personal tutor for every learner, once available only to aristocrats, might

finally become universal. Whether that promise is fulfilled or betrayed depends not on the technology, but on the choices institutions, governments, and learners make in the next decade.

We argue that the most critical skill for navigating this transition – for individuals and institutions alike – is learning velocity: how fast you can detect errors, pivot, and build again. The traditional map is gone. The future belongs to those who stop preparing for a stable world and start adapting to a reality of constant motion.

Annotated Source List

26. **Huang 2024** — “The Future of AI and Education.” An AI Video citing Jensen Huang, CEO of NVIDIA, discussing how artificial intelligence will reshape education, workforce preparation, and the value of traditional degrees. Presented as a public talk. Cited for its framing of the AI-education narrative rather than as empirical evidence. **Tier 4** (individual executive opinion without peer review or empirical data).

References

Huang, Jensen. 2024. “The Future of AI and Education.” Video.
<https://www.youtube.com/watch?v=sjGFJNY2v1k>.

How Systems Actually Change (and Why It's Slow)

It is tempting, after surveying the brokenness of higher education and the power of AI, to leap straight to solutions. New models exist. Some of them work. So why isn't everything already different?

Because education systems are not software. You cannot push an update and watch millions of users migrate overnight. Higher education is shaped by overlapping actors with conflicting incentives, protected by regulations written for a different era, and held in place by cultures that reward continuity over experimentation. Understanding who these actors are, what they want, and what constrains them is essential before any discussion of innovation can be taken seriously.

This chapter maps the playing field. It identifies the main actors who shape whether and how higher education changes, explains why good ideas spread unevenly, and offers a realistic check on where the most-discussed innovations actually stand today. It closes with a simple framework for thinking about what comes next.

Who Decides? The Key Actors and What Drives Them

Governments and Policymakers

Governments hold the most powerful levers: public funding, regulation, accreditation standards, and student financial aid rules. In most European countries, public universities still rely heavily on state budgets, and even in the United States, federal financial aid eligibility determines which institutions can survive. When governments tie funding to performance contracts and strategic priorities, they can accelerate change, but they can also undermine institutional autonomy and long-term research capacity if the formulas are poorly designed.

The catch is political risk. Education reform affects millions of voters, parents, and workers. Politicians who push too fast risk backlash from faculty unions, alumni networks, and communities that depend on local campuses for economic stability. Politicians who do nothing risk being blamed when graduates cannot find jobs or when tuition keeps rising. The result is a pattern familiar across democracies: bold language in policy documents, cautious incrementalism in practice.

Funding formulas deserve special attention. When governments allocate money based on enrollment headcounts and degree completions, they implicitly reward the traditional model: students in, degrees out. Shifting to performance-based or outcomes-based formulas can encourage innovation, but it can also punish institutions that serve harder-to-reach populations with lower completion rates. Several European systems have introduced performance contracts linking a portion of funding to strategic goals, but the share of funding at stake is often modest enough that institutional behavior changes only at the margins.

Accreditation Bodies: Gatekeepers or Blockers?

Accreditation agencies exist to protect students and the public from low-quality education. They set standards for faculty qualifications, curriculum design, assessment rigor, and institutional governance. Without accreditation, an institution's degrees carry no recognized value, and its students cannot access government financial aid.

The problem is that many accreditation standards were built around a traditional model: credit hours, seat-time, full-time faculty with terminal degrees, physical campuses. When an institution like Western Governors University designs a competency-based program where students advance by demonstrating mastery rather than completing fixed terms, it collides with these legacy frameworks. WGU's 2017 federal audit is the most vivid example. The Office of the Inspector General argued that WGU's self-paced, mentor-supported model constituted "correspondence education" rather than distance education and recommended the university repay \$713 million in federal financial aid. The Department of Education ultimately rejected the most extreme findings, but the audit created years of perception risk: prospective students questioned degree value, and other institutions considering CBE models took note of the regulatory hazard.

Accreditation reform is happening, but slowly. Some agencies are developing competency-based standards, and the European approach to quality assurance for micro-credentials emphasizes that providers should have robust internal mechanisms rather than requiring each individual credential to be evaluated. But reform timelines are measured in years, not months, and the agencies themselves face capacity constraints: small staffs, limited technical expertise in emerging models, and boards composed largely of leaders from traditional institutions.

University Leaders and Faculty

University presidents and provosts often understand the case for change. Many have read the same reports about demographic decline, employer frustration, and technological disruption. But they operate within governance structures that distribute power widely. In most universities, faculty senates control curriculum, hiring, and academic standards. Departments guard their budgets and course enrollments. Tenure protects individual faculty from institutional pressure to change but also reduces the institution's ability to redeploy resources quickly.

Faculty culture matters as much as formal governance. Academics are trained to be skeptical, to evaluate evidence carefully, and to resist trends that lack rigorous support. These are virtues in research. In institutional decision-making, they can produce a default toward the status quo. When a new teaching model requires faculty to shift from lecturing to coaching, to learn new technologies, to redesign assessments, and to accept that their role may evolve, the response is often cautious. Many instructors lack training in AI tools and digital pedagogy, and professional development is chronically underfunded.

Unions add another layer. In systems with strong faculty unions, changes to workload, class size, online delivery, and assessment methods are subject to collective bargaining. This is not inherently bad; unions protect working conditions and academic

freedom. But it does mean that reforms requiring changes to faculty roles face negotiation timelines that can extend for years.

Framing faculty response as “inertia” or “resistance to change” can obscure what is often a rational political calculation. Faculty opposition to reform frequently reflects not opposition to innovation itself, but opposition to the terms on which change is imposed – particularly in systems where faculty have been precarious through casualisation, where reforms are experienced as deskilling through algorithmic mediation, or where institutional restructuring threatens disciplinary autonomy and professional identity. In European systems especially, faculty unions, tenure structures, and disciplinary identities function as primary political obstacles to reform not because faculty are conservative by temperament, but because they are organised and have legitimate grounds for concern. A German university considering employer-aligned micro-credentials must contend not with individual reluctance but with powerful union organisations representing precarious Wissenschaftliche Mitarbeiter who correctly perceive that algorithmic credentialing threatens their gatekeeping authority. A French university proposing AI-integrated assessment faces faculty councils where disciplinary associations have formal power and real incentives to defend pedagogical autonomy against platform standardisation.

Reform narratives that treat faculty resistance as an obstacle to be overcome through better change management miss a critical political reality: successful reform in systems with organised faculty requires either genuine faculty buy-in through co-design (which substantially slows and constrains what is possible), or sufficient political capital – government mandate, funding leverage, or generational turnover – to override resistance. The first path is slower but produces durable change rooted in practitioner expertise. The second is faster but creates lasting resentment and often collapses when political will shifts. Acknowledging faculty resistance as partly rational and politically grounded is not a counsel of despair. It is a prerequisite for reform strategy that accounts for actual power relations rather than pretending they do not exist.

Employers: Loud Demands, Quiet Inertia

Employers consistently report skills gaps. Surveys find persistent mismatches in digital competencies, communication, teamwork, and adaptability. Many companies have publicly announced moves toward skills-based hiring, dropping degree requirements for some positions, and investing in workforce upskilling through platforms like Guild Education.

Yet employer behavior often lags employer rhetoric. Hiring managers still rely heavily on degrees as screening filters, particularly for roles that carry regulatory requirements (healthcare, law, engineering, finance) or reputational stakes. A 2024 survey found that 60% of hiring managers consider alternatives to traditional degrees, but “consider” and “consistently hire based on” are different things. When a company receives hundreds of applications, the degree checkbox remains a fast, cheap filter, even if it is an imperfect one.

Employers also send mixed signals about micro-credentials. While 92% report they would hire candidates with GenAI-specific micro-credentials, close to half (46%) say they cannot assess program quality when a candidate presents a non-degree

credential. The demand signal is noisy: employers want specific skills but lack the frameworks to evaluate alternative ways of certifying them.

EdTech Firms and Platforms

The educational technology market has grown rapidly, fueled by venture capital, pandemic-era acceleration, and genuine unmet needs. Platforms offering adaptive learning, AI tutoring, micro-credential delivery, and learning management systems have become essential infrastructure for many institutions.

But the incentive structure of venture-backed EdTech deserves scrutiny. Investors seek rapid growth, market capture, and returns on timelines (5-10 years) that do not match the pace of educational change. This creates pressure to scale fast, lock institutions into proprietary ecosystems, and prioritize engagement metrics over learning outcomes. Vendor lock-in is a real risk: once a university builds its curriculum around a particular platform's architecture, switching costs become prohibitive.

Platform economics also tend toward consolidation. The largest providers can invest more in AI, data infrastructure, and content, making it harder for smaller or public alternatives to compete. This raises questions about who controls the learning infrastructure of the future: public institutions accountable to learners and society, or private platforms accountable to shareholders?

Students: Consumers, Citizens, and Constrained Agents

Students are often described as the primary beneficiaries of educational innovation, and they should be. But their agency is constrained in ways that reformers sometimes overlook. Most students choose from a limited set of options shaped by geography, family resources, information access, and credential requirements set by employers and regulators. A first-generation student in a mid-sized European city does not face the same menu of choices as an upper-middle-class student in a global capital.

Students are also subject to what the existing literature calls the “agency divide.” Self-directed, digitally fluent, motivated learners thrive in flexible, technology-rich environments. Students who need more structure, mentorship, and scaffolding can struggle in the same settings. Innovations designed for the first group risk excluding the second, and the second group disproportionately comes from lower-income and marginalized backgrounds.

The consumer framing itself is incomplete. Students are not just purchasing a credential; they are developing as thinkers, citizens, and people. Reducing education to a transaction between a buyer and a provider misses the developmental, social, and civic dimensions that make higher education a public good, not just a private investment.

Why Good Ideas Spread Unevenly

If a model works, why doesn't everyone adopt it? Because “works” is context-dependent, and the conditions for adoption vary dramatically across countries, institutions, and regulatory regimes.

Micro-Credentials: Same Idea, Different Speeds

Micro-credentials illustrate the pattern clearly. The European Commission has championed a common European approach to micro-credentials, aligned with Bologna structures (ECTS credits, EQF levels, quality assurance). The ECIU alliance became the first European University Alliance to issue jointly recognized, blockchain-secured digital micro-credentials using the Europass wallet in 2023. The European Commission is advancing a joint European degree label to be rolled out from mid-2026.

Yet adoption varies enormously by country. Some national qualifications frameworks have embedded micro-credentials and linked them to established quality assurance processes. Others treat them as unregulated continuing-education offerings with no formal recognition. Government financial aid rarely covers micro-credentials, creating a 23% barrier to adoption according to the research literature. The result: the same innovation moves at different speeds depending on whether national regulators have cleared the path or left it blocked.

In the United States, the landscape is even more fragmented. Over one million unique educational credentials now exist, a threefold increase since 2018, offered by universities, industry providers, bootcamps, and platforms. Without standardized definitions, learning outcomes, or assessment rigor, micro-credentials vary wildly in value. Employers report confusion: 46% cannot assess program quality when a candidate presents a non-degree credential. The innovation exists, but the ecosystem needed to make it trustworthy and portable does not yet exist at scale.

The WGU Audit: When Regulation Meets Innovation

Western Governors University serves over 200,000 students with fully online, self-paced, competency-based degree programs. Its model lets learners complete pre-assessments, work through only the material they need, and take final assessments when ready. Some students have completed both a bachelor's and a master's degree in nine months by leveraging prior knowledge.

But the 2017 federal audit exposed a structural tension that affects every CBE program. The OIG argued that WGU's model, which contracts course content from various sources and employs mentors rather than traditional instructors, did not meet the regulatory definition of "distance education" because it lacked sufficient "regular and substantive interaction" between faculty and students. The recommended \$713 million repayment would have been existential. The Department of Education rejected the finding, but the message to other institutions was clear: innovate at your own regulatory risk.

This is not an isolated case. Institutions without a substantial core of qualified faculty with full-time responsibility face steep eligibility challenges for accreditation. The debates about whether mentor-based, content-contracted models meet standards for faculty governance of curriculum and academic quality absorb resources and create uncertainty that discourages smaller institutions from attempting similar reforms.

Adoption Reality Checks

For each of the most-discussed innovations, it is worth asking three simple questions: How big is this actually? Where is it concentrated? What is blocking wider adoption?

Competency-Based Education (CBE)

- **Current scale:** CBE remains concentrated in a small number of large institutions, primarily in the United States. WGU (200,000+ students) and SNHU are the dominant players. Most traditional universities have not adopted CBE at scale, and many CBE pilots at smaller institutions have quietly ended.
- **Where concentrated:** Overwhelmingly in the US, serving adult and non-traditional learners. European adoption is minimal, partly because Bologna structures already define learning outcomes but retain credit-hour and seat-time norms in practice.
- **Main blockers:**
- Accreditation frameworks built around seat-time and faculty structures that CBE disrupts.
- The measurement trap: defining and assessing competencies at scale is harder than it sounds. A meta-review found that despite CBE approaches emerging in the 1970s, there is very limited evidence for their effectiveness.
- CBE privileges self-directed, digitally literate learners and struggles to serve students who need social scaffolding and structured progression.

Micro-Credentials

- **Current scale:** Growing rapidly in volume but unevenly in quality and recognition. Over one million unique credentials exist in the US alone. In Europe, the ECIU alliance and several national systems have begun formal integration, but most micro-credentials remain outside national qualifications frameworks.
- **Where concentrated:** Tech-adjacent fields, professional development, and corporate upskilling. University-issued micro-credentials are a small fraction of the total market. Employer recognition remains inconsistent.
- **Main blockers:**
- Employer confusion: nearly half cannot assess the quality of non-degree credentials.
- Fragmentation: stackability promises often fail in execution because institutions lack credit equivalency frameworks and transparent metadata for true portability.
- Financial aid exclusion: government student aid rarely covers micro-credentials, limiting access for lower-income learners.

AI-First Educational Models

- **Current scale:** No comprehensive AI-first university operates at scale in higher education as of early 2026. The Alpha School example frequently cited in discussions is a K-12 prototype with a small cohort and no longitudinal outcomes data. Individual AI tools (adaptive platforms, predictive analytics, AI tutors) are deployed across many institutions, but none has built an entire institutional model around AI from the ground up.
- **Where concentrated:** Pilot programs at large, well-resourced institutions like Arizona State University and SNHU. Conceptual proposals in academic literature and policy documents.
- **Main blockers:**
- No proven model at scale with demonstrated learning outcomes.
- Faculty unpreparedness: professors often lack training in AI tools, and implementation requires substantial investment in platforms, cybersecurity, and continuous monitoring.
- Regulatory uncertainty: the EU AI Act, which entered into force in stages starting in 2024, classifies AI systems used in education as “high-risk,” requiring conformity assessments, transparency obligations, and human oversight. Compliance costs and legal uncertainty slow institutional adoption.
- Algorithmic bias: 80% of AI systems in education showed some form of bias when not properly audited, raising equity concerns that regulators and institutions have not fully resolved.

European Universities Alliances

Current scale

- Over 60 alliances now exist under the European Universities Initiative, launched in 2017. These involve hundreds of institutions across Europe and represent significant public investment through Erasmus+ and Horizon Europe funding.

Where concentrated

Predominantly Western European, research-intensive universities. Over 75% of EUI member institutions rank among the top 500 globally. Institutions in Central and Eastern Europe, smaller universities, and applied-sciences institutions are disproportionately underrepresented.

Main blockers

- The Matthew effect: alliances built on pre-existing partnerships and elite status consolidate advantages for already-privileged institutions, deepening vertical stratification rather than spreading innovation broadly.
- Resource asymmetries: smaller institutions face higher relative costs to participate (staff time, travel, co-funding) while receiving proportionally smaller benefits.

- Regulatory heterogeneity: cross-border recognition of joint degrees and micro-credentials requires aligning diverse national frameworks, quality assurance regimes, and legal structures, a process that moves at the speed of the slowest participant.

A Scenario Matrix: What Shapes the Next Decade

Rather than predicting a single future, it is more honest and more useful to map the space of plausible outcomes. Two axes capture the forces that matter most:

- **Regulatory openness:** How willing are governments and accreditation bodies to adapt rules for new models? At one extreme, regulators actively create space for competency-based credentials, AI-integrated pedagogy, and cross-border recognition. At the other, existing rules remain unchanged, and innovators must fit new models into old frameworks or operate outside the recognized system.
- **Institutional capacity:** How able are universities to invest in transformation? This includes financial resources, technical infrastructure, faculty development, and organizational agility. At one extreme, institutions have the funding, talent, and governance flexibility to redesign programs. At the other, they are constrained by budget cuts, legacy systems, and internal resistance.

These two axes produce four quadrants:

Quadrant 1: Open Rules, Strong Capacity (Accelerated Transformation)

Regulators clear the path. Well-resourced institutions invest in CBE, AI integration, micro-credential ecosystems, and cross-border alliances. Innovation spreads relatively quickly, though concentration at elite institutions risks deepening stratification. This is the scenario most innovation advocates imagine.

Quadrant 2: Open Rules, Weak Capacity (Uneven Disruption)

Rules allow experimentation, but most institutions lack the money, talent, or organizational flexibility to take advantage. A few large, well-funded players (and EdTech platforms) capture the opportunity. Smaller and public institutions fall further behind. The gap between innovative and traditional institutions widens.

Quadrant 3: Closed Rules, Strong Capacity (Frustrated Innovation)

Institutions have the resources and will to change but are blocked by rigid accreditation standards, funding formulas tied to seat-time, and regulatory uncertainty (e.g., EU AI Act compliance costs). Innovation happens at the margins: internal pilots, continuing education, corporate partnerships that sit outside the regulated degree system. Change is slow and incremental.

Quadrant 4: Closed Rules, Weak Capacity (Institutional Stagnation)

Neither the regulatory environment nor institutional resources support transformation. Universities continue operating with legacy models, facing declining enrollment, employer frustration, and fiscal stress. Change comes, if at all, through crisis: closures,

mergers, or external disruption by platforms and employers who build parallel credentialing systems outside traditional higher education.

What This Means

No country or institution sits neatly in one quadrant. Most occupy messy positions between them, with some programs innovating while others remain unchanged, some regulators experimenting while others hold firm. The point of the matrix is not to predict but to clarify the variables that matter.

Two conclusions follow. First, innovation alone is not enough. Without regulatory adaptation, even the best new models remain marginal or legally precarious. Without institutional capacity, even the most permissive rules produce nothing. Both levers must move.

Second, the pace and direction of change are choices, not inevitabilities. Governments can redesign funding formulas and accreditation standards. Universities can invest in faculty development and organizational agility. Employers can build credible frameworks for evaluating alternative credentials. Students can demand transparency and quality. None of these things will happen automatically. All of them require actors to accept short-term costs and political risks for long-term gains.

Five Feedback Loops in Education

![[feedback-loop-in-education-2.png]]

Higher education is shaped by five reinforcing loops, each connecting regulatory choices, institutional strategy, hiring practices, and learner pathways. Understanding which loops are active and which are blocked tells you why change stalls or accelerates.

Loop 1: Accreditation → Institutional Structure → Program Design

When accreditation bodies require seat-time, faculty with terminal degrees, and physical classrooms, universities design programs that fit these rules. When rules allow competency-based models, micro-credentials, and distributed faculty, universities can – but do not automatically – redesign. The tighter the accreditation constraint, the more institutions converge on similar models. This loop is now *tight* in most places, creating inertia. Loosening it is necessary but not sufficient; institutions with tight budgets or risk-averse leadership will converge on the same model anyway.

Loop 2: Hiring Practices → Credential Demand → Institutional Investment

When employers filter resumes by degree, students demand degrees, and universities invest in bachelor's programs. When employers hire based on demonstrated skills and portfolios, students seek micro-credentials and apprenticeships, and universities invest in stackable credentials and work-integrated programs. This loop is *misaligned* today: employers say they want skills-based hiring but still filter by degree, sending contradictory signals that leave institutions uncertain where to invest. Resolving this misalignment requires explicit employer commitment to alternative hiring practices.

Loop 3: Funding Formula → Institutional Priority → Learner Access

When governments fund universities based on enrollment headcounts, institutions prioritize scale and completion, which can mean easier courses and larger cohorts. When governments fund based on learning outcomes or performance contracts, institutions prioritize assessment rigor, but this can inadvertently exclude low-achieving or under-prepared students. This loop is *mediated by redistribution*: governments can design formulas that reward both completion and serving difficult populations, but most do not. Shifting this loop requires deliberate policy design to avoid perverse incentives.

Loop 4: EdTech Market Dynamics → Platform Consolidation → Institutional Lock-in

When venture-backed platforms compete on feature richness and user engagement, they grow quickly, gain market share, and create switching costs for universities that integrate their tools into curriculum. Early adopters buy into a platform's infrastructure (content, assessments, data architecture), and later replacements become prohibitively expensive. This loop is *rapid and concentrated*: a handful of platforms now control large shares of learning management, adaptive content, and credential verification. Breaking this loop requires public alternatives and transparent standards for data portability.

Loop 5: Stratification → Credibility Divergence → Legitimacy Erosion

When universities divide into tiers (elite research universities, mid-tier teaching institutions, struggling regional campuses), the same credential means different things depending on its source. Employers learn that a degree from a top university predicts something different than a degree from a struggling institution. This credibility divergence increases as institutions in different tiers can no longer afford equivalent faculty, resources, or opportunities. The loop is *self-reinforcing*: tier-based legitimacy makes it harder for lower-tier institutions to attract talented students and faculty, which erodes quality further. This loop is particularly acute in systems with weak regulatory quality assurance and wide resource variation.

Understanding Loop Dynamics Across Futures

The scenario matrix describes the conditions that favor change or stagnation. But which loops activate, and which tighten, depends on the future that unfolds. Understanding how loop dynamics work across different scenarios clarifies what choices matter most.

In Drift: All Five Loops Remain Unchanged or Tighten

Accreditation stays rigid (Loop 1 tightens). Hiring practices send mixed signals, with employers continuing to filter by degree despite rhetoric around skills (Loop 2 remains misaligned). Funding formulas reward enrollment, not outcomes (Loop 3 continues to reward completion over access). EdTech consolidates as a handful of platforms dominate (Loop 4 accelerates). Stratification deepens as institutions in different tiers can no longer afford equivalent resources (Loop 5 tightens). In this future, the dominant pattern is inertia. Change happens only at the margins.

In Platform World: Loops 1, 3, and 5 Tighten; Loops 2 and 4 Are Hijacked

Regulation and public funding weaken (Loops 1 and 3 lose traction as accreditation and government funding become less relevant). Employers and platforms become the de facto accreditors and funders, replacing government and universities. Loop 2 is hijacked: employers use algorithmic hiring to sort candidates, but only within platform ecosystems. Loop 4 is not broken; it is weaponized. A single platform (or tight oligopoly) controls the infrastructure. Stratification accelerates as platforms concentrate resources among large institutions they partner with (Loop 5 tightens dramatically). The system becomes more efficient for organizational productivity, but less accountable to learners or the public.

In Redesign: Loops Are Intentionally Rebalanced

Accreditation loosens, creating space for competency-based models and alternative credentials (Loop 1 loosens). Hiring practices shift as employers commit to skills-based hiring at scale (Loop 2 aligns). Funding formulas reward learning outcomes and equity, not just completion (Loop 3 is restructured). EdTech is governed as public infrastructure, with open standards and data portability (Loop 4 is broken). Investment is redistributed to prevent stratification; quality assurance is strengthened in lower-tier institutions (Loop 5 is actively managed to prevent deepening). In this future, the dominant pattern is intentional redesign.

This translation shows that the futures are not random: they follow from which loops policymakers choose to tighten or loosen. The scenario matrix identifies the conditions. The five-loop framework identifies the mechanisms.

The Hard Choices: What We Must Give Up

Every serious reform story hides a refusal problem. We like to talk about innovation, modernization, and redesign—but we avoid naming what must end. Education will not change because better tools arrive. It will change only when institutions, governments, and learners accept losses that feel uncomfortable, politically risky, and culturally destabilizing.

Let's talk about those losses.

Not everything worth preserving can be preserved. Not every tradition deserves protection. And not every inefficiency is a virtue. If education is to remain legitimate in an age of AI, it must give up several assumptions that once made sense—but no longer do.

What follows are the hard choices we cannot evade.

1. We must give up the lecture as the default unit of teaching

The lecture was a rational solution to a world of scarce expertise and expensive transmission. That world is gone.

Continuing to center education around one-to-many content delivery is no longer tradition—it is denial. When a personalized AI tutor can explain concepts more clearly,

patiently, and adaptively than a human lecturer, insisting on lectures as the core experience is not pedagogy; it is nostalgia.

This does *not* mean giving up explanation, storytelling, or inspiration. It means giving up the idea that explaining content *is* teaching.

What replaces it is harder: designing experiences that require judgment, collaboration, confrontation with uncertainty, and responsibility for outcomes. Many institutions will resist this shift because it is more demanding, less scalable, and harder to standardize. That resistance must be overcome.

2. We must give up the fiction that time spent equals learning

Seat time, credit hours, semesters, and degrees are administrative conveniences masquerading as epistemology.

They persist because they are legible to funding formulas, accreditation systems, and bureaucracies—not because they reliably indicate competence. In a world of continuous learning and rapid skill decay, fixed-duration credentials increasingly misrepresent what learners can actually do.

Letting go of time-based proxies means embracing uncomfortable questions: - How do we recognize uneven learning trajectories? - How do we certify mastery without standard pacing? - How do we fund institutions when progress is nonlinear?

There are no easy answers. But clinging to time as a proxy for learning is easier only in the short term. In the long term, it corrodes trust.

3. We must give up the idea that fairness means treating everyone the same

Uniformity is not equity.

Mass education systems were built on standardization: same curriculum, same pace, same assessment. This created scalability—but it also embedded structural advantage for those already aligned with the system’s norms.

AI exposes this trade-off brutally. Personalization can unlock extraordinary gains—but only if systems deliberately build agency and support. Without that, personalization becomes a sorting machine that amplifies prior advantage.

True fairness in the AI era requires differentiated pathways, targeted scaffolding, and explicit investment in agency-building for those who need it most. That means abandoning comforting narratives of “equal treatment” and replacing them with harder commitments to unequal support.

This will feel unjust to those who benefited from uniformity. It is still necessary.

4. We must give up the degree as a universal proxy for merit

Degrees will not disappear—but they must lose their monopoly as the sole proxy for competence.

For too long, degrees have served as a blunt sorting device: expensive, slow, and increasingly disconnected from actual capability. Their dominance has distorted incentives throughout the system—fueling credential inflation, grade inflation, and performative learning.

In the AI era, demonstrated capability—projects, portfolios, simulations, and evidence of learning velocity—must carry greater weight alongside degrees. This requires employers, regulators, and institutions to accept messier, less standardized signals. It does not mean that degrees become irrelevant—in regulated professions and for entry-level hiring, they remain the most legible and trusted filter available. What must change is the assumption that the degree alone is sufficient.

The alternative is worse: a world where degrees persist as the sole gatekeeper long after they have ceased to be informative. That path preserves hierarchy, not quality.

5. We must give up the illusion of frictionless reform

There is no pain-free transition.

Some institutions will shrink. Some roles will disappear. Some academic identities will be destabilized. Some communities will lose familiar anchors. Pretending otherwise erodes trust and fuels backlash.

The goal is not to avoid loss, but to *choose it consciously*. To decide which frictions protect core values—and which merely protect legacy structures.

Universities that try to add AI on top of existing models without shedding obsolete functions will exhaust their staff, confuse their students, and dilute their mission. Survival will require subtraction, not just addition.

6. We must give up the belief that education can remain neutral

Education always encodes values. The question is not whether it does so, but whether those values are examined or smuggled in under the guise of neutrality.

In the AI era, refusing to take a stance defaults power to platforms, markets, and algorithms whose incentives are not aligned with public purpose. That is not neutrality; it is abdication.

Education systems must be explicit about what they stand for: - judgment over compliance - agency over passivity - public value over private optimization

This will invite conflict. Avoiding that conflict is no longer an option.

The uncomfortable truth

The hardest choice is this: **we must stop designing education for a stable world that no longer exists.**

That means letting go of predictability, linear pathways, and the promise that following the rules guarantees a good outcome. It means preparing learners not for certainty, but for responsibility in uncertainty.

AI makes intelligence cheap. That forces a reckoning.

If education refuses to give up what no longer serves its purpose, it will preserve its forms while losing its meaning. If it accepts these losses deliberately, it may finally recover what mattered all along: forming people capable of thinking, choosing, and acting wisely when no script is available.

That is the price of relevance.

The chapters that follow examine the most promising innovations in detail, their real evidence base, their pitfalls, and the design choices that determine whether they expand opportunity or deepen inequality. But every one of those innovations operates within the system described here. Understanding the system is the first step toward changing it.

New Models of Learning: What's Already Here

The traditional university is not the only game in town anymore. Over the past two decades, a handful of radically different institutional models have moved from experiment to scale, each testing a different hypothesis about how higher education could work. Some bet on technology. Others bet on peers, employers, or cross-border networks. A few bet on eliminating the campus, the professor, or the credit hour altogether.

None of these models has “solved” higher education. Each carries its own blind spots, trade-offs, and failure modes. But together they offer a practical catalogue of what is already possible – and a clear-eyed look at what goes wrong when good ideas meet messy reality.

This chapter profiles five archetypes. For each, we ask the same four questions: - What is it? - Why does it look promising? - What does the evidence actually say? - And what can go wrong?

The goal is not a shopping list of innovations - but an honest inventory – one that helps readers judge - which ideas deserve scaling, which need redesign, - and which are more marketing than substance.

1. Competency-Based Online Universities

Western Governors University (WGU) and Southern New Hampshire University (SNHU)

What it is

Competency-based education (CBE) replaces the credit hour with demonstrated mastery. Students do not sit through a semester of lectures and take an exam at the end; instead, they work through material at their own pace and prove competence whenever they are ready.

- WGU, founded in 1997 by 19 U.S. state governors, now serves over 200,000 students with fully online, self-paced degree programs.
- Learners complete pre-assessments to identify what they already know, study only the material they need, and take final assessments when ready – no waiting for the semester to end.
- Students must complete a minimum of 12 competency units per term but can accelerate beyond that as much as effort and prior knowledge allow.
- One graduate completed both a bachelor’s and a master’s degree in leadership and management within nine months.

SNHU extends the CBE model to underserved populations – refugees, frontline workers, low-income adults – through partnerships with community organizations. Its Duet model combines competency-based curricula with high-touch wraparound support: coaching, advising, and mental-health resources.

SNHU is also piloting generative-AI tutors built on its Learner Information Framework (LIF) to identify at-risk students early and deliver personalized interventions before challenges escalate.

Why it looks promising

- **Time becomes flexible.** The principle is simple: “learning is the constant and time is the variable.” Students who already know the material move fast; those who need more time get it. This is the opposite of the industrial-age model where everyone moves at the same speed regardless of preparation.
- **Cost drops when pace increases.** WGU’s flat-rate tuition per six-month term means that students who accelerate pay less overall. For working adults with prior professional experience, this can cut the cost of a degree by half or more.
- **Reaches populations traditional universities miss.** SNHU’s partnerships with community organizations bring degree programs to refugees and frontline workers who would never set foot on a residential campus. The Duet model’s wraparound support – coaching, advising, mental-health resources – yields higher persistence and completion rates among non-traditional learners.
- **Employer alignment is built in.** CBE curricula are mapped directly to employer-valued competencies. SNHU structures its direct-assessment programs around real-world projects rather than abstract coursework, eliminating the “course” as a time-based container entirely.

What we know works

The evidence base for CBE is thinner than advocates suggest – and the honest reader should pay close attention to the quality labels.

Meta-review (independent research, 2015)

Despite competency-based approaches emerging in the 1970s, a systematic meta-review concluded that “there is hardly any evidence for the effectiveness of competence-based education.” Four decades of implementation produced remarkably little rigorous, independent evaluation of learning outcomes.

Graduate outcomes (vendor/institutional data)

WGU and SNHU report strong employment and satisfaction numbers for their graduates, but these are institutional self-reports rather than independent studies. WGU graduates’ employer satisfaction rates are frequently cited in university marketing materials; independent verification is limited.

Pilot data (SNHU institutional report)

SNHU’s Duet model reports higher persistence and completion rates among non-traditional learners compared to standard online programs, though published results come from the institution itself rather than external evaluators. SNHU’s AI-driven personalization pilot, powered by the Learner Information Framework, combines

behavioral analytics with generative AI to deliver personalized feedback via email and the learning management system.

[Employer survey data \(independent, 2024\)](#)

A 2024 survey found 60% of hiring managers consider alternatives to traditional degrees, and 92% report they would hire candidates with GenAI micro-credentials. This suggests growing employer receptivity to CBE outputs, though receptivity is not the same as demonstrated effectiveness.

[The CBE evidence paradox.](#)

There is a genuine tension in the evidence base that deserves explicit acknowledgment. On one hand, the most rigorous independent meta-reviews conclude that four decades of CBE implementation have produced “hardly any evidence” for its effectiveness as a pedagogical approach.

On the other hand, the two largest CBE providers – WGU and SNHU – report strong completion rates, graduate employment outcomes, and employer satisfaction figures that exceed many traditional programs.

How can both be true?

The most likely explanation is selection effects and institutional design: WGU and SNHU serve motivated adult learners with prior professional experience, wrap their CBE models in substantial support infrastructure (mentoring, advising, AI-driven early intervention), and measure outcomes that may reflect who enrolls as much as what the model teaches. Their success is real but may not generalize to CBE as a universal approach.

The honest conclusion is that CBE works well for specific populations under specific conditions, but the broader claim that competency-based education is inherently superior to time-based models remains unsubstantiated by independent research. For the purposes of this book, CBE should be treated as a plausible and partially validated hypothesis about how to structure learning – not as an evidence-based recommendation. Where we discuss CBE models approvingly in subsequent chapters, we do so on the basis of design logic and early institutional outcomes, not on the basis of the kind of independent, controlled, multi-site evidence that would justify confident prescription.

[What can go wrong](#)

[The measurement trap.](#)

Defining and assessing competencies at scale proves far more complex than replacing credit hours with learning outcomes. When competencies are tied narrowly to current employer needs – as CBE advocates often recommend – they fail to provide the transferable, adaptable foundations learners need for uncertain futures.

When institutions try to assess underpinning knowledge separately from performance, assessment becomes overloaded and cumbersome – undermining the very efficiency CBE promises. Some CBE programs embedded in digital instructional software deem

students “proficient” at 80% pass rates, effectively reinforcing lower standards and allowing knowledge gaps to accumulate.

Without coherent design, CBE can feel like a collection of disconnected modules rather than an integrated educational experience, sacrificing the interdisciplinary synthesis and critical reasoning that traditional curricula – at their best – cultivate.

The accreditation battle.

WGU’s 2017 federal audit exposed how CBE models collide with legacy regulatory frameworks. The Office of the Inspector General argued that WGU’s self-paced, mentor-supported courses constituted “correspondence education” rather than distance education, recommending the university repay \$713 million in federal financial aid.

The Department of Education ultimately rejected the most extreme findings, but the audit created lasting perception risk: prospective students questioned degree value and financial aid eligibility.

The underlying tension persists: many accrediting bodies base standards on measurable metrics tied to traditional time-based structures, making competency-based programs difficult to quantify and slow to gain approval.

WGU’s model – contracting course content from various sources, employing mentors rather than traditional instructors – triggered debates about whether such structures meet accreditation standards for faculty governance of curriculum and academic quality.

For smaller or newer CBE programs, these battles absorb resources and create existential uncertainty.

Not for all learners.

CBE’s self-paced, individually focused model privileges learners who are already self-directed, digitally literate, and capable of sustained motivation without social scaffolding.

It sidelines the importance of social learning, peer interaction, and the developmental role of structured deadlines and cohort progression – supports that benefit many students, particularly those from under-resourced backgrounds.

Some CBE providers exploit the “competency-based” branding to promote programs of dubious educational value, further eroding public trust in the model as a whole.

2. Global Campus-Free Models

Minerva University

What it is

Minerva University operates without a traditional campus. Instead, students rotate through global cities – San Francisco, Berlin, Buenos Aires, Seoul, Taipei – over four years, immersing themselves in diverse cultural, economic, and policy contexts.

All classes are conducted as small, fully active-learning seminars via a proprietary digital platform; there are no lectures. Minerva's curriculum is skills-focused and competency-based, mapping every course to granular, measurable learning outcomes – data analysis, argumentation quality, ethical reasoning – which are reinforced iteratively across disciplines.

Students apply classroom knowledge to real-world challenges through local internships, community projects, and field research in each city.

Why it looks promising

Active learning replaces passive listening.

Every class session requires student participation; the platform tracks engagement and performance in real time. There is no place to hide in a Minerva seminar.

Global immersion builds adaptability.

Living and working in multiple countries over four years exposes students to different institutional, cultural, and economic contexts – a form of preparation that residential campuses in a single location cannot match.

Lower sticker price than elite peers.

Tuition is set at approximately \$10,000 per year, less than one-third of Harvard's. The absence of a physical campus, tenure-track faculty, and research infrastructure enables this cost structure.

Measurable learning gains (internal data).

Internal studies report at least a 50% improvement in student learning over three years using active pedagogy, compared to a 92% knowledge-loss rate within six months for traditional lecture-based methods.

What we know works

Internal longitudinal study (vendor data)

Minerva's own research shows at least 50% improvement in student learning outcomes over three years using its active-learning pedagogy. The comparison baseline – a 92% knowledge-loss rate for traditional lecture methods – comes from independent cognitive science research on retention curves.

Pedagogical design (independent research)

Active learning's superiority over lectures is well-established in the broader education literature, with multiple meta-analyses confirming the effect. Minerva's contribution is implementing this at institutional scale, not discovering the underlying principle.

Graduate trajectories (vendor/institutional data)

Minerva reports strong graduate placement outcomes, but the institution is young and cohorts are small. Long-term, independent outcome data does not yet exist.

What can go wrong

The cost barrier is real despite the marketing.

While tuition is \$10,000 annually, total costs – tuition, room, board, and living expenses across seven global cities – reach approximately \$29,000 per year. This places Minerva out of reach for many qualified students, particularly those from the Global South. The model advertises affordability but delivers a price point that, while lower than Harvard, still exceeds what most families worldwide can manage.

Elite selectivity is maintained, not eliminated.

Minerva targets the “best students” and positions itself as a meritocratic alternative to artificially scarce elite institutions. But “qualified” is defined by academic preparation, English fluency, and the capacity to navigate intensive, small-seminar active-learning pedagogy.

Admissions are need-blind, but the model does not address the systemic barriers – under-resourced secondary schools, lack of college counseling, digital divides – that prevent talented learners from lower-income or marginalized backgrounds from becoming “qualified” in the first place. Critics question whether Minerva genuinely expands access or simply creates a new tier of selectivity at a slightly lower price point.

Institutional sustainability questions.

The absence of tenure and research programs, while enabling cost savings, raises questions about long-term intellectual depth, faculty governance, and whether the model can sustain itself without the revenue diversification – endowments, research grants, alumni giving tied to campus identity – that traditional universities rely on.

Evidence is largely internal.

The learning-gains data comes from Minerva’s own studies. Independent, peer-reviewed validation of the model’s long-term outcomes – career trajectories, civic engagement, intellectual development over decades – does not yet exist at scale.

3. Peer-Driven, Teacher-Free Models

42 School (Ecole 42)

What it is

42 School, originally founded in Paris, now operates a global network of 54 campuses across 31 countries. It eliminates teachers, classes, and tuition fees entirely.

Students learn by completing real-world coding and software-engineering projects, accessing freely available internet resources and peer support in open-plan computing labs. Peer-to-peer learning is formalized: a randomly assigned fellow student marks each project, and learners advance through 21 levels by demonstrating competencies in a gamified progression structure.

Graduation typically takes three to five years and awards a certificate rather than a formal degree. The model is supported by corporate partnerships and governments seeking to address digital-talent shortages.

Why it looks promising

Zero tuition removes financial barriers to entry.

Unlike virtually every other higher education model, 42 charges nothing. Corporate sponsors and government funding cover costs, removing the debt burden that shapes (and distorts) educational choices for millions of students.

Self-direction and peer accountability mirror real workplaces.

The absence of teachers forces students to develop the exact skills employers say they want: independent problem-solving, resourcefulness, collaboration under ambiguity.

Employers value the graduates.

Companies report that 42 graduates are more adept at independent problem-solving and less reliant on supervisor direction than graduates of conventional programs. The model produces workers who can function without being told what to do.

Rapid global scaling.

From one campus in Paris, 42 has expanded to 54 campuses in 31 countries in roughly a decade – a pace of institutional replication that traditional universities cannot match.

What we know works

Employer feedback (informal industry data):

Companies hiring 42 graduates consistently report stronger independent problem-solving skills and lower reliance on supervisor direction. This is informal industry feedback, not a controlled study, but it is consistent across multiple employers and countries.

Scaling record (observational)

The expansion from 1 campus to 54 across 31 countries demonstrates replicability of the model. However, outcome data across campuses is not publicly reported in standardized form.

Selection process outcomes (institutional data)

The piscine process successfully identifies learners who will persist through the program. Dropout data during and after the piscine is not publicly available in disaggregated form.

What can go wrong

Selection bias is the model, not a bug.

42's intensive month-long “piscine” (French for “swimming pool”) selection process requires applicants to navigate close collaboration, peer feedback, and high-pressure

problem-solving for weeks before admission. Those who struggle with interpersonal dynamics or need more explicit instruction drop out during selection.

The model is designed for “resourceful, self-directed learners” – it does not develop these traits in learners who lack them. This means 42’s impressive outcomes are partly a function of who gets in, not just what happens inside.

No mentorship means no safety net.

With no teachers and no structured guidance, learning quality depends entirely on the knowledge and generosity of peers. When peer expertise is insufficient – a common problem in advanced or niche topics – or when interpersonal conflicts arise, students lack recourse to authoritative support.

The absence of mentorship is particularly risky for students from backgrounds where academic and professional networks are thin: the very populations that most need guidance receive the least.

Who thrives versus who struggles.

The model works for learners who already possess high intrinsic motivation, digital fluency, and collaborative capacity. It does not serve populations requiring developmental support, remedial education, or socio-emotional services.

42 scales horizontally – more campuses – but it does not scale inclusively to serve broader populations. The model is a powerful solution for a specific demographic; it is not a replacement for the diverse functions of a university.

No degree, limited portability.

Graduation awards a certificate, not a degree recognized within national qualifications frameworks. In labor markets and countries where degree requirements remain standard for hiring, immigration, or professional licensing, 42 graduates face structural disadvantages regardless of their actual competence.

4. Employer-Sponsored Learning Ecosystems

Guild Education

What it is

Guild Education disrupts traditional tuition-reimbursement models by facilitating upfront, employer-paid tuition for frontline workers. Rather than the old model – where employees pay tuition, submit receipts, and hope for partial reimbursement – Guild curates a marketplace of vetted degree programs and certificates from university and industry partners, connects employees to offerings aligned with their employer’s strategic needs, and manages payment flows so workers incur no debt.

Over six million employees accessed Guild in the past year. The model operates as a business-to-business (B2B) platform: employers are the clients, not individual learners.

Why it looks promising

Eliminates debt for frontline workers.

Workers in retail, food service, and healthcare – populations that traditional universities often fail to serve – can earn credentials without borrowing money. The employer pays upfront, removing the single largest barrier to adult education.

Measurable career impact.

Guild participants are 2.6 times more likely to experience internal mobility and earn wage increases compared to non-participants. Chipotle reported that crew members enrolled in Guild programs are six times more likely to advance into management roles.

High retention signals mutual value.

Chipotle reports that Guild enrollees have 89% retention rates after nine months and are nearly twice as likely to be promoted. Internal job mobility increases by 3.5 times. Both employers and employees appear to benefit.

Addresses both sides of the skills gap.

Employers secure a more capable, loyal workforce; employees gain credentials and career progression without upfront financial barriers.

What we know works

Employer outcome data (vendor/client reports)

Guild's published numbers – 2.6x internal mobility, Chipotle's 6x management advancement rate, 89% nine-month retention – come from Guild and its corporate clients. These are vendor and client reports, not independent studies. The numbers are impressive but produced by parties with clear incentives to report positive results.

Scale of reach (vendor data)

Over six million employees accessed Guild in the past year. This demonstrates market demand and employer willingness to invest, though it does not by itself demonstrate learning effectiveness.

No independent analysis

No independent meta-analysis or controlled study of Guild's model has been published. The model is relatively new and proprietary, making independent evaluation difficult.

What can go wrong

Corporate capture of curriculum.

Guild curates which programs employers fund, and employers decide which offerings to support and which employees to enroll. Over 80% of degree and certificate learners enroll in business-aligned programs selected to meet employers' strategic workforce needs.

This ensures relevance to the employer but narrows educational offerings toward instrumental, skills-focused training at the expense of liberal arts, critical inquiry, and public-good missions. Universities dependent on Guild partnerships may face pressure to design programs that prioritize employer satisfaction over intellectual rigor, disciplinary integrity, or student-centered pedagogy – a dynamic that risks transforming universities from sites of broad human development into corporate training arms.

Worker lock-in.

Guild's business model succeeds partly because it increases retention: participants are 2.6 times more likely to remain with their employer. While employees gain credentials and wage increases, they do so within a closed ecosystem that incentivizes staying rather than exploring external opportunities.

Workers who leave their employer before completing programs may forfeit progress or face financial penalties, creating economic lock-in. From an individual standpoint, Guild provides opportunity; from a labor-market standpoint, it may reduce worker bargaining power by tying educational benefits to continued employment and limiting the portability of skills and credentials outside the employer's network.

Who decides what you learn?

Unlike traditional financial aid or tuition reimbursement, Guild's B2B model places employers – not learners – at the center of decision-making. Employees choose from a curated set of offerings rather than open access to all higher education options, restricting autonomy and potentially directing learners toward credentials that serve employer needs more than individual aspirations or long-term career flexibility.

The question is whether employer-sponsored education is a second chance or a gilded cage.

5. European University Alliances

ECIU, CHARM-EU, and the European Universities Initiative

What it is

The European Universities Initiative (EUI), launched in 2017 and now comprising over 60 alliances, represents a bold experiment in transnational, integrated higher education. Alliances such as ECIU (European Consortium of Innovative Universities) and CHARM-EU (Challenge-Driven, Accessible, Research-Based, Mobile European University) jointly develop curricula, award micro-credentials or full joint degrees, and offer students seamless mobility and recognition across partner institutions.

ECIU became the first European University Alliance to issue jointly recognized, blockchain-secured digital micro-credentials using the Europass wallet in 2023, establishing a template for distributed credentialing across its 14 partner institutions. ECIU structures its entire educational offering around learner-selected “challenges” co-designed with societal partners – companies, municipalities, NGOs. Students form interdisciplinary teams, define problems within challenge areas, investigate solutions using domain-appropriate knowledge, and deliver actionable outputs: prototypes, advisory reports, proof of concept. Faculty shift from lecturers to coaches.

CHARM-EU offers a fully joint master's degree, the first of its kind in Europe, with curriculum co-created and delivered across nine partner universities. Every module is jointly designed to ensure coherent integration of theory, disciplines, and dissertation projects.

The European Commission is advancing a joint European degree label to be rolled out from mid-2026, enabling automatic recognition of jointly awarded programs everywhere in the EU.

Why it looks promising

Micro-credentials with cross-border recognition.

Micro-credentials – typically 4-6 ECTS in size – certify specific competencies and can stack toward larger qualifications or stand alone. ECIU's blockchain-secured credentials, issued through the Europass wallet, offer a model for portable, verifiable, learner-owned recognition that works across 14 institutions and multiple countries.

Challenge-based learning at institutional scale.

ECIU structures learning around authentic societal problems co-designed with external partners, shifting pedagogy from lectures to interdisciplinary team projects with real deliverables. Tilburg University's Digital Interventions Incubator exemplifies this: students develop virtual-reality tools for children with neurodevelopmental disorders or augmented-reality apps to encourage physical activity in patients, working directly with clinical partners.

Joint degrees eliminate institutional silos.

CHARM-EU's jointly designed master's degree integrates theory, disciplines, and dissertation projects across nine universities, offering students a genuinely transnational educational experience rather than a patchwork of exchange semesters.

Innovation labs embed experimentation.

ECIU established Innovation of Education Labs at each campus – physical and virtual spaces where faculty experiment with AI-supported competence recommendations, challenge-based pedagogy, and digital tools.

What we know works

Micro-credential infrastructure (pilot/institutional data)

ECIU's blockchain-secured credentials using the Europass wallet represent the first working implementation of cross-border digital micro-credentialing in European higher education. This is a successful infrastructure pilot, not yet a proven system at scale.

Challenge-based learning outcomes (pilot studies)

Utrecht University's Da Vinci honors program, integrating challenge-based learning with design thinking, reports cultivating competencies including stakeholder engagement, empathy, navigating uncertainty, and integrative problem-solving. These are pilot-level observations from specific programs.

Joint degree implementation (institutional data)

CHARM-EU's joint master's degree is operational across nine partner universities. This demonstrates institutional and regulatory feasibility but does not yet have outcome data comparing graduates to those of conventional programs.

What can go wrong

The Matthew effect: the rich get richer.

Over 75% of EUI member institutions rank among the top 500 globally, and most alliances build on pre-existing partnerships. Institutions with international networks, co-funding capacity, and strong reputations secured alliance membership; those without – particularly demand-absorbing lower-tier universities in Central and Eastern Europe – were disproportionately excluded.

The partial reallocation of Erasmus+ and Horizon Europe funds to predominantly Western European universities consolidates their competitive advantages. Scholars describe this as the Initiative's "Matthew effect": the system designed to foster cooperation instead reproduces existing hierarchies.

West-east stratification deepens.

Geographic coverage, resource asymmetries, and opportunity costs create structural barriers. Smaller institutions in countries with limited national support for internationalization face higher relative costs to participate – staff time, travel, co-funding requirements – while receiving proportionally smaller benefits.

The Initiative services the agendas of excellence and competitiveness more than inclusion and cooperation, contrasting with other European Commission strategies (such as Teaming and Twinning) that explicitly target peripheral regions and less-resourced universities. Institutions positioned higher in the European landscape leverage alliances to enhance their own competitiveness, further stratifying the sector between participating and non-participating universities.

ECTS challenges and credential fragmentation.

The promise of "stackability" – earning micro-credentials that accumulate toward degrees – often fails in execution. Without thoughtful curricular design, stackable credentials become disjointed modules lacking coherent progression or interdisciplinary integration.

Credit transfer between institutions remains complex despite blockchain and open-badge infrastructure: most systems lack the credit equivalency frameworks and transparent metadata required for true portability. National frameworks diverge: some countries embed micro-credentials within national qualifications frameworks and link them to established quality assurance, while others treat them as unregulated continuing-education offerings. The result is a patchwork where recognition depends more on geography than on learning.

Employer confusion about micro-credentials.

When job applicants present non-degree credentials, close to half of employers (46%) cannot assess program quality, and 42% cannot judge the skills and competencies acquired. A systematic review found that 80% of employers cite consistency as their primary concern. Authenticity worries follow: 38% of studies report employer concern about fraudulent credentials due to the sheer variety and lack of verification infrastructure.

The United States alone is now home to over one million unique educational credentials, a threefold increase since 2018. When credentials proliferate faster than quality assurance mechanisms can keep pace, the result is noise, not signal.

Learners report that 40% do not know where to start in the digital credential landscape, and 60% worry that costs will be out of reach. Government financial aid rarely covers micro-credentials, creating a 23% barrier to adoption according to the literature. Narrow knowledge attainment is another documented risk: 43% of the literature reviewed notes that micro-credentials deliver competencies limited to one specific niche of a job scope, restricting career mobility and adaptability.

Core Design Principles: What the Models Share

Beneath the surface differences, the most effective innovations converge on three principles that invert traditional academic assumptions. These principles are not tied to any single model; they appear, in various forms, across the archetypes described above.

Understanding why these models work – or fail – requires looking beneath the design innovations to the learning science embedded in each approach. The five archetypes differ not just in delivery modality or funding, but in fundamental assumptions about motivation, cognitive load, transfer, and the conditions under which learners develop metacognition and agency. Traditional universities bundle these elements invisibly; innovative models must make them explicit or risk collapse.

Motivation varies across the models in ways that matter more than marketing suggests. CBE's self-paced structure works because WGU and SNHU attract motivated adult learners with professional stakes in credential completion. But the model also reveals an uncomfortable truth: removing structural deadlines and cohort progression scaffolds self-directed motivation into a selection criterion rather than a developmental outcome. For learners without intrinsic urgency – many teenagers from under-resourced backgrounds, for instance – the flexibility becomes paralysis. The research on motivation distinguishes between extrinsic drivers (external rewards, deadlines, social pressure) and intrinsic ones (autonomy, competence, purpose). CBE optimizes for intrinsic motivation but assumes it exists. The question Krishan's analysis leaves open: Which pedagogical structures cultivate intrinsic motivation in learners who do not arrive with it already formed? 42's peer-driven model partially answers this through social accountability (peer review, public standing in the cohort), but at the cost of excluding learners uncomfortable with that form of pressure. This is not a design flaw in isolation – it is a necessary trade-off, but one that requires honest acknowledgment.

Cognitive load operates at multiple levels across these models. CBE proponents argue that self-paced learning prevents overload by letting students control pace. But the measurement trap reveals a different problem: when institutions pile assessment atop competency demonstration atop feedback loops, the total cognitive demand can exceed what it was in a traditional semester course, particularly for students managing other work and family obligations. Minerva's active-learning seminars deliberately compress content delivery into small, peer-tested exchanges to reduce passive information load, but they simultaneously increase social and performance pressure on each student in the room. This is not inefficient; it is a conscious shift of cognitive burden from retention to judgment. But it assumes students have sufficient executive function, emotional regulation, and social confidence to operate under constant observation. Students without these capacities do not simply perform worse; they experience the environment as hostile.

Transfer – the ability to apply learning in contexts beyond the training setting – is where the models show the starker divergence. Guild, 42, and employer-sponsored programs promise employer-aligned competencies that transfer directly to the job. But decades of learning science research shows that skills trained in context-specific ways often do not transfer well to novel situations. A student who learns Excel through job-task examples may not recognize how to apply spreadsheet thinking to a different problem. SNHU and ECIU attempt to build transfer capacity by embedding skills across multiple contexts and disciplines, but this is effortful and invisible in curriculum documents. Minerva's approach – teaching argumentation, ethical reasoning, and data analysis across different content domains and geographies – is explicitly designed for transfer, but this kind of cross-disciplinary reinforcement requires substantial instructional overhead. The book identifies that many CBE curricula “fail to provide the transferable, adaptable foundations learners need for uncertain futures,” but it does not fully explore *how* to design for transfer without reverting to the broad, slow curricula that CBE was meant to replace. This is a design problem that remains partially open – but learning science offers more concrete guidance than is typically applied in modular programme design.

Three principles, grounded in decades of cognitive research, can be systematically embedded in competency-based curricula without reverting to the slow, undifferentiated breadth that CBE was designed to replace. The first is *contextual variation*: transfer fails when learners encounter a skill in only one context, because they encode it as bound to that context. The solution is not to teach broadly, but to deliberately vary the contexts in which core competencies are practised. A data-analysis competency developed through spreadsheet exercises, then statistical software applied to health datasets, then algorithmic problem-solving in code, then visualisation for policy decision-makers builds what learning scientist David Perkins calls “flexible knowledge.” This requires identifying which three to four contexts will develop a particular competency most effectively, sequencing them strategically, and linking them explicitly. ECIU’s challenge-based model approximates this through problem variation, but most CBE systems do not systematise it.

The second principle is *explicit abstraction and metacognitive scaffolding*. Transfer depends not just on varied contexts but on learners becoming aware of the underlying principle connecting them. Instructional design that works for transfer includes

structured “bridging” moments where students identify the core principle operating across contexts and anticipate where it might apply next. This can be embedded in competency assessments: rather than asking “Can you do X?” ask “Where else might you apply this? What would change? What would stay the same?” Metacognitive scaffolding is particularly important in self-paced, low-contact models like WGU, where there is less dialogue to naturally surface these connections.

The third is what Perkins terms “hugging” – anchoring abstract principles to concrete examples and allowing learners to move between concrete and abstract levels. A competency in “critical evaluation of evidence” is too abstract to transfer unless practised concretely through evaluating evidence in healthcare claims, then marketing assertions, then research papers, with periodic reflection on what “evaluating evidence” means across these domains. This requires design choices that many modular systems resist: accepting that transfer capacity takes instructional time, that the module map must include deliberate redundancy across domains, and that assessment must test transfer explicitly – whether the learner can apply a skill in a new domain, not just demonstrate mastery in the training context.

Metacognition – the capacity to monitor one’s own thinking, recognize gaps in understanding, and adjust strategy – is fostered by different mechanisms across the models. Traditional seminars build metacognition through dialogue: when a peer or professor questions your reasoning, you are forced to examine your own thinking explicitly. Adaptive AI systems promise to do this at scale, but the evidence is mixed. AI tutors can identify when a student’s answer is wrong, but diagnostic precision about *why* a student is wrong – whether it is a conceptual misunderstanding, a computational error, a misreading of the problem, or a gap in background knowledge – requires human judgment or extraordinarily sophisticated AI. SNHU’s Learner Information Framework combines behavioral analytics with AI feedback, which is promising, but the manuscript does not explore whether algorithmic feedback alone can replicate the developmental effect of human dialogue. The question underlying 6.1 is: In what conditions does metacognition develop, and which model innovations support it versus undermine it?

Agency – the capacity to set goals, evaluate options, and act on deliberate choices rather than follow prescribed paths – is the deepest pedagogical challenge across all models. Crucibles supposedly forge agency through authentic challenge and high stakes. But authentic challenge can also be traumatizing if the learner lacks psychological resources or perceives the stakes as arbitrary rather than meaningful. The European alliances’ challenge-based learning approach explicitly builds agency by allowing students to define problems within broad domains rather than follow preset pathways, but this assumes a baseline of confidence and self-advocacy that not all students bring. The question Krishan hints at but does not fully answer: How do crucibles foster judgment and agency without reproducing social hierarchies – that is, without becoming spaces where students from privileged backgrounds naturally feel at home and others feel excluded? And how to scaffold agency for students who arrive already self-directed? These are not rhetorical questions. They are design constraints that determine whether innovations expand or narrow educational possibility.

The through-line connecting these elements is that pedagogical effectiveness depends on alignment between learner readiness, instructional design, and assessment. The five models do not fail because their underlying science is wrong. They fail at scale because

they assume a particular kind of learner or require substantial human support infrastructure that disappears under cost pressure. The honest reading is that no single pedagogical model works for all learners. The task of the AI-era university is not to choose between active learning and scaffolded instruction, or between self-paced and cohort-based, but to deploy multiple pedagogical approaches and explicitly match them to learner profiles, developmental readiness, and goals. This requires moving away from one-size-fits-all curricula toward genuinely adaptive systems that preserve human judgment at critical decision points.

From knowledge transmission to capability building

AI's ability to generate, summarize, and deliver content renders traditional lecture-based dissemination less defensible. If a student can get a clear explanation of any concept from an AI tutor at 2 a.m., the value of a professor reading slides aloud in a 300-seat auditorium drops to near zero.

The universities that are adapting – Minerva, ECIU, SNHU – are pivoting toward cultivating judgment, critical thinking, relational understanding, and ethical reasoning: capacities that machines cannot replicate. This shift demands pedagogies that emphasize active learning, authentic problem-solving, and reflective practice.

GenAI also exposes how curricula and assessments have drifted from deeper purposes of higher learning toward performance and reproduction. Institutions responding strategically are redesigning assignments to be AI-resistant – emphasizing process, collaboration, and context-specific reasoning – and integrating critical AI literacy into curricula so students understand algorithmic bias, data privacy, and the epistemological limits of datafication.

But there is a risk embedded in this shift: over-reliance on AI threatens to degrade the very capacities universities claim to build. Generative AI presents information with confidence, including hallucinations, leading users to accept outputs uncritically. Management educators warn that if students do not learn to critically engage with AI – questioning its sources, biases, and limits – they will carry uncritical acceptance into professional contexts. A 2021 study found that 80% of AI systems in education showed some form of bias when not properly audited. The tools are powerful, but they are not neutral.

Continuous curriculum renewal

Traditional curriculum review cycles – three to five years or longer – cannot keep pace with labor-market and technological change. A living curriculum treats program design as continuously reviewed, informed by disciplinary developments, employer feedback, student experience, and societal challenges.

The Quintuple Helix model provides a framework for this ongoing dialogue, bringing together academia, industry, government, civil society, and environmental considerations to co-design relevant, adaptive learning pathways.

In practice, this means moving from periodic committee reviews to real-time feedback loops. ECIU's Innovation of Education Labs – physical and virtual spaces at each

partner campus – let faculty experiment with AI-supported competence recommendations, challenge-based pedagogy, and digital tools.

The challenge is institutional: universities are built for stability, not speed. Governance structures, accreditation cycles, and faculty incentives all favor the status quo. Continuous renewal requires not just new tools but new organizational cultures.

Hybrid human-AI mentoring

AI-powered career guidance and mentoring systems analyze academic records, interests, skills, and labor-market trends to recommend personalized pathways. Platforms provide real-time job-market alignment, virtual internships, and AI tutors available 24/7.

Research on individual development plans in STEM graduate education has identified two hybrid models: **Sequential Integration**, where students consult AI mentors first and human mentors then refine guidance; and **Concurrent Collaboration**, where AI and human mentors operate in parallel and students synthesize inputs.

These models combine AI's speed and data breadth with human mentors' empathy, contextual understanding, and ethical judgment. They hold particular promise for underrepresented minorities in STEM, where mentorship gaps are widest.

But hybrid mentoring only works if the human side is adequately resourced. When universities replace human mentors with AI agents to cut costs rather than to augment support, they risk compounding the isolation and loneliness that already plague higher education. Teens and young adults are forming emotional attachments to AI chatbots, sometimes substituting them for human relationships. Mental health professionals report rising social anxiety and loneliness as drivers.

Universities that automate mentoring without investing in human connection may save money in the short term while deepening the mental-health crisis they claim to address.

Agency: How the Models Perform

Chapter 3 established that agency – the capacity to direct your own learning, set goals, persist through difficulty, and adapt to change – is socially produced, not innate, and identified four conditions required to build it: scaffolded challenge with feedback, explicit metacognitive instruction, human mentorship, and psychological stability. It also posed three diagnostic questions for every model: Does it treat agency as an entry requirement or a learning outcome? Who bears the cost when agency is absent? How does it move learners from dependence to independence?

Applied to the five archetypes in this chapter, the results are sobering.

WGU and SNHU split along this axis. WGU's self-paced model assumes agency as a given: students must bring their own discipline, time management, and persistence. There is no cohort, no structured progression, no mentor who notices you have gone quiet. For motivated adults with professional experience, this works. For learners without those capacities, “your own pace” often means no pace at all. SNHU's Duet model is the more honest design: it wraps CBE in coaching, advising, and mental-health support – investing in the scaffolding that agency requires rather than assuming it.

Minerva demands exceptionally high agency from day one. Its small-seminar, fully active-learning format leaves no room to disengage – every session requires participation, and the platform tracks performance in real time. This builds agency in students who already have enough of it to survive the environment. But the model does not address the upstream barriers – under-resourced secondary schools, lack of college counselling, digital divides – that prevent talented learners from lower-income backgrounds from arriving prepared for that intensity. Minerva develops agency in its students; it does not develop the conditions that produce agency-ready applicants.

42 School is the starker case. Its month-long piscine selection process is explicitly designed to identify learners who already possess high intrinsic motivation, digital fluency, and collaborative capacity. It does not develop these traits in learners who lack them. The absence of teachers and structured mentorship means there is no institutional mechanism for building agency – the model selects for it and calls the result merit. 42 scales horizontally (more campuses) but not inclusively (broader populations).

Guild Education restricts agency in a different way. Workers choose from a curated set of offerings pre-selected by their employer, not from the full landscape of higher education. The scaffolding exists – employer funding removes the financial barrier – but the learner’s direction is shaped by corporate strategy rather than personal aspiration. Agency in the sense of self-determined learning goals is structurally constrained.

ECIU and CHARM-EU come closest to building agency by design. Challenge-based learning explicitly requires students to define problems within broad domains rather than follow preset pathways, developing the goal-setting and self-direction that agency requires. But this assumes a baseline of confidence and self-advocacy that not all students bring – and the Matthew effect in alliance membership means the students most likely to benefit are those at already-privileged institutions.

The pattern across all five models is consistent: the more a model relies on learner agency, the better it works for those who already have it – and the more it excludes those who do not. The models that invest in building agency (SNHU’s Duet, ECIU’s challenge-based approach) are more expensive and harder to scale. The models that select for agency (42, WGU, Minerva) are leaner and more impressive in their headline outcomes – but those outcomes partly reflect who gets in, not just what happens inside.

What This Landscape Tells Us

Five models, five different bets on what matters most in education. None has cracked the code.

WGU and SNHU prove that competency-based learning can reach hundreds of thousands of students – but four decades in, the evidence for CBE’s effectiveness remains strikingly thin.

Minerva shows that active learning and global immersion work – for students who can afford \$29,000 a year and navigate intensive seminars in English.

42 demonstrates that peer-driven, tuition-free coding education produces strong graduates – by selecting for self-directed learners and filtering out everyone else.

Guild proves that employers will fund education when it serves retention – but the question of who controls the curriculum, and for whose benefit, remains open.

European alliances show that transnational cooperation and micro-credentials are technically feasible – while inadvertently concentrating resources among already-privileged institutions.

The pattern is consistent: innovation solves some problems while creating others. The models that expand access often narrow curriculum. The models that deepen learning often restrict who can participate. The models that reduce cost often shift control from learners to employers or algorithms.

But there is a harder conclusion that the pattern demands. This is not a limitation of individual models that better design could fix. It is a structural feature of the alternative education landscape as a whole. Nearly every celebrated model in this chapter selects for the very quality – agency – that Chapter 3 identified as the scarcest and most unequally distributed resource in education. The reform agenda's central promise is that innovation can expand opportunity for those whom traditional systems have failed. The evidence from these five models suggests something narrower: the most impressive alternatives succeed partly because they filter for learners who already possess strong self-direction. Their outcomes are real. But those outcomes partly reflect who gets in, not just what happens inside.

This forces an honest choice among three responses. First: accept these models as serving a motivated tier, and stop claiming they address systemic inequality wholesale. Second: invest seriously in upstream agency development – preparatory programmes, bridge curricula, mentorship, psychological support – that builds self-direction before entry. SNHU's Duet model points in this direction, but at costs far exceeding headline tuition. Third: design explicitly different models for different populations – self-paced and peer-driven models for high-agency learners, structured progression with human mentorship for learners developing agency, and stabilisation services for learners in crisis before any learning model becomes meaningful.

All three are preferable to the current default: expanding access to models that implicitly select for privilege while using the language of equity. This is not a reason for cynicism. It is a reason for design honesty. The next chapter of higher education will not be written by any single model but by the deliberate choices institutions, governments, and learners make about which trade-offs they are willing to accept – and which ones they refuse to tolerate.

Degrees, Skills, and the New Signalling Stack

Earlier in this book we made a bold claim: the degree is shifting from a gate to a preference. Later, when we looked at how employers actually hire, we found a messier picture – companies that drop degree requirements on paper still struggle to interpret what candidates bring instead. One chapter announces the death of credentials; another shows the funeral has been poorly attended.

Both observations are true. They just operate on different timescales and in different sectors. This chapter reconciles them by asking three questions: Why did degrees become so powerful in the first place? What is genuinely changing? And what should we realistically expect over the next five, ten, and twenty years?

The answer, we argue, is not that degrees disappear. It is that they become one layer in a broader signalling stack – a multilevel system where degrees, micro-credentials, portfolios, and AI-derived behavioural data each carry different weight depending on the profession, the employer, and the moment in a person’s career.

Why Did Degrees Become So Powerful?

A Brief History of Signalling

For most of human history, you proved competence by doing the work. A blacksmith apprenticed for years, and the village could see whether the horseshoes held. A merchant’s reputation was local and personal. Signalling – the economist’s term for communicating your qualities to strangers who cannot directly observe them – was not a problem because communities were small and information travelled slowly.

Industrialisation changed that. When employers began hiring strangers at scale, they needed a fast, cheap filter. The modern university degree emerged as exactly that filter. It bundled three signals into one piece of paper: cognitive ability (you could handle four years of intellectual work), conscientiousness (you showed up and finished), and socialisation (you absorbed the norms of a professional class). Employers did not need to know what you studied in your third-year seminar. They needed to know you survived the process.

The Prussian education system of the 18th century established compulsory schooling, standardised curricula, and state certification – a template that spread globally and entrenched the idea that credentials should be issued by institutions authorised by the state. Over two centuries, this model became so deeply embedded in regulation, professional licensing, and cultural expectation that questioning it felt almost absurd.

By the late 20th century, the degree had become what sociologists call an “institutional fact” – valuable not because of what it contained but because everyone agreed to treat it as valuable. Employers required degrees. Governments subsidised them. Professional bodies mandated them. Parents expected them. The cycle reinforced itself.

The strongest version of this argument comes from economist Bryan Caplan, whose *The Case Against Education* (2018) contends that the degree premium is largely a signalling

premium: employers reward credentials not because education makes workers more productive, but because completing a degree signals pre-existing traits – intelligence, conscientiousness, conformity – that employers value. If Caplan is substantially right, then most of what universities teach is economically irrelevant; the degree functions as an expensive, slow sorting mechanism that could in principle be replaced by cheaper signals.

This book does not adopt Caplan's conclusion, but it takes the signalling critique seriously – more seriously, in fact, than the human-capital orthodoxy does. The Signalling Stack framework developed later in this chapter implicitly accepts a core Caplan premise: that degrees are valued partly for their signal, not solely for their substance. The question is what follows. Caplan argues for reducing public investment in education. We argue for building better signals. If the degree's value is partly signalling, then the solution is not to abolish the signal but to create signals that are more informative, more granular, and more closely tied to demonstrated capability. A portfolio of shipped projects is a better signal of programming ability than a computer science degree – not because the degree taught nothing, but because the portfolio shows what the candidate can do now, not what they survived four years ago. The Signalling Stack is, in this sense, a partial answer to Caplan: it accepts the signalling diagnosis while rejecting the nihilistic prescription.

Where the Caplan critique does bite is on the question of public subsidy. If a significant share of the degree's labour-market value comes from signalling rather than human capital formation, then subsidising degrees is partly subsidising a sorting mechanism that could be replaced by cheaper alternatives. This does not mean universities have no value beyond signalling – research, civic formation, intellectual development, and social mobility are real functions that signalling theory does not capture. But it does mean that the economic case for universal degree subsidies is weaker than policymakers typically assume, and that investment in alternative signalling infrastructure (quality-assured micro-credentials, verifiable portfolios, standardised skills assessments) may deliver more efficient returns.

The Degree Premium – and Its Limits

The economic case for degrees remains real. Across OECD countries, graduates earn significantly more over a lifetime than non-graduates. But this premium is an average that conceals enormous variation by field, institution, and individual. A computer science degree from a well-regarded university pays for itself many times over. A degree in a field with weak labour-market demand from an institution with high dropout rates may leave a graduate worse off than someone who entered the workforce directly.

Grade inflation has further eroded the degree's informational value. When students are treated as customers entitled to high grades, the signal weakens. An A no longer reliably distinguishes deep understanding from surface compliance. Employers notice. As one hiring manager put it, degrees tell you someone finished something – but not what they can actually do.

What Is Changing?

The Employer Shift Toward Skills

Skills-based hiring has moved from conference rhetoric to measurable practice. In 2025, 81% of employers globally report using skills-based approaches, up from 73% in 2023. In Europe, 77% of companies plan to prioritise competency-based recruiting in 2026. The talent pool expansion is striking: skills-based hiring increases the pool of eligible candidates by 15.9 times in the United States and 10.4 times in the United Kingdom. For Gen Z workers, the multiplier reaches 17.6 times in the U.S. and 12.7 times in the UK.

These numbers reflect a genuine structural shift, not a temporary fashion. The persistent talent shortage – cited by 24% of global CEOs as a primary business threat – forces employers to look beyond traditional credential filters. When you cannot find enough people with the right degree, you start asking what people can actually do.

AI-powered skills assessments, used by 41% of talent acquisition teams, make this feasible at scale. Online simulations, gamified tests, and project-based evaluations provide data that a resume cannot. Some 83% of companies now use some form of gamified assessment. The technology that makes degrees less necessary for filtering also makes alternatives more practical.

Portfolios and Demonstrated Capability

The shift from credentials to portfolios is already visible in specific industries. In software development, a GitHub profile with real contributions often matters more than a degree listing. In design, a portfolio of shipped products is the entry requirement. In data science, a well-documented Kaggle competition or a published analysis demonstrates capability in ways a transcript cannot.

Projects are becoming the new transcripts. In a world where AI can generate standard answers, creation is the only honest proof of competence. The question is no longer “What school did you attend?” but “What have you built?”

This shift is accelerated by the economics of verification. AI makes it cheap to assess skills through projects, simulations, and portfolios. When verification is expensive and slow, employers default to institutional proxies like degrees. When verification becomes fast and cheap, the proxy loses its monopoly.

Blockchain Credentials and Verifiable Trust

Resume fraud and credential verification bottlenecks cost employers billions annually. Blockchain-based credentials offer a technical solution: cryptographically verifiable proof of education, certifications, and employment history. Platforms like TrueProfile.io, Blockcerts, and Credly by Pearson create immutable digital records that employers can verify instantly.

The European approach is particularly ambitious. The European Consortium of Innovative Universities (ECIU) became the first European University Alliance to issue jointly recognised, blockchain-secured digital micro-credentials using the Europass

wallet in 2023. Open Badges 3.0, introduced in 2022, integrates blockchain technology with verifiable credentials, creating a decentralised system where each badge contains detailed metadata – learning path, objectives achieved, criteria met – embedded in a digital wallet that learners own and control.

The European Blockchain Partnership is developing unified standards for interoperable credential verification across the single market. North America leads in private-sector implementation, particularly in tech and healthcare where certification is critical.

Beyond fraud prevention, blockchain credentials enable lifelong learning records where workers accumulate micro-credentials from multiple sources, automated skill matching where AI systems read verified credentials without manual data entry, and simplified cross-border recognition of foreign qualifications.

AI-Verified Work and Behavioural Signals

Perhaps the most disruptive development is still emerging: AI systems that assess not just what you know or what you have built, but how you work. Adaptive learning platforms already track trajectory rather than just status – measuring learning velocity, the speed at which a student can unlearn, pivot, and build again. Predictive analytics evaluate behavioural patterns: consistency, collaboration habits, response to feedback, ability to sustain focus over time.

These behavioural signals are qualitatively different from traditional credentials. A degree tells an employer what you survived. A portfolio shows what you built. A behavioural profile shows how you operate – and whether you are likely to keep growing. In recruitment, AI algorithms already evaluate factors like job change frequency, behavioural traits, and assessment responses to flag candidates likely to leave early or to forecast on-the-job performance.

The ethical implications are significant – algorithmic bias, privacy concerns, and the risk of reducing people to data patterns are real dangers. But the direction of travel is clear: the information available to employers about candidates is becoming richer, more granular, and more continuous than a degree could ever be.

The Four-Layer Signalling Stack

Rather than replacing degrees with a single alternative, we are moving toward a layered system. Think of it as a signalling stack, where different layers carry different weight depending on the context.

Layer 1: Baseline – Degrees

Degrees are not going away. In regulated professions – law, medicine, engineering, nursing, teaching, architecture – they remain hard gates enforced by professional bodies, government regulation, and liability frameworks. You cannot practise medicine without a medical degree, regardless of how impressive your portfolio is. You cannot sign off on a bridge design without an accredited engineering qualification. These gates exist for good reasons: they protect the public from unqualified practitioners in fields where mistakes can kill.

Even outside regulated professions, degrees serve as a baseline signal that is difficult to replicate. They certify a sustained period of intellectual engagement, exposure to a discipline's methods and debates, and the ability to complete a multi-year commitment. For entry-level hiring, where candidates have little work history, this baseline remains useful.

The degree's role, however, is narrowing. It is becoming a necessary-but-not-sufficient condition in many fields rather than the dominant signal it once was.

Layer 2: Skill – Micro-Credentials, Short Courses, Certifications

Micro-credentials address labour-market velocity by unbundling degrees into smaller, skill-focused units. Typically 4-6 ECTS in size, they certify specific competencies and can stack toward larger qualifications or stand alone. A 2024 survey found that 60% of hiring managers consider alternatives to traditional degrees, and 92% report they would hire candidates with GenAI micro-credentials.

Harvard Extension School permits two-course micro-certificates and three-course certificates to accumulate toward bachelor's or master's degrees. Miami Dade College maps stackable pathways from short college-credit certificates through career-technical certificates to full degrees. The European approach aligns micro-credentials with Bologna structures – EQF levels, ECTS, quality assurance – to ensure cross-border recognition.

Industry certifications from companies like Google, AWS, Microsoft, and Salesforce operate in the same layer. They signal current, specific, verifiable skills in a way that a four-year degree issued a decade ago cannot.

Layer 3: Evidence – Portfolios, Projects, Simulations, Code Repositories

This layer provides direct evidence of capability. It includes GitHub repositories, design portfolios, published analyses, open-source contributions, completed simulations, and documented project outcomes. Unlike credentials, which certify that someone once demonstrated competence, evidence shows what someone can do right now.

VR and AR assessment technologies are expanding this layer. Walmart uses VR to place candidates in high-pressure retail scenarios, observing decision-making under stress. Manufacturing firms create virtual factory floors where candidates troubleshoot equipment failures. Healthcare organisations simulate patient interactions to evaluate clinical judgment. These immersive assessments generate objective, standardised data that AI can analyse for pattern matching against top performers.

Project-based hiring platforms enable candidates to complete real-world tasks, giving companies direct insight into work approach and quality. Software companies present coding challenges mirroring actual product problems. Marketing agencies ask candidates to develop campaign concepts. These exercises reveal not just technical skill but creativity, problem-solving approach, and professional judgment.

Layer 4: Behaviour – AI-Derived Patterns (Speculative)

A note on status: Layers 1 through 3 describe observable, currently functioning elements of the hiring landscape. Layer 4 is qualitatively different. It describes a plausible future development, not a current practice. No hiring system currently uses behavioural signals as a primary input at scale. The inclusion of this layer in the stack reflects a directional trend – not demonstrated adoption. Readers should weight it accordingly.

The newest and most speculative layer – one that does not yet exist at scale – would capture behavioural signals derived from how people learn and work over time. Learning velocity – how fast someone adapts to new material, detects errors, and iterates – is measurable through adaptive platforms. Collaboration patterns, consistency of output, and response to feedback can be tracked across digital work environments.

This layer is embryonic and raises serious questions about privacy, bias, and consent. But it represents the logical endpoint of a trend: employers want to predict future performance, and AI systems can now infer patterns from behavioural data that no credential or portfolio captures.

The signalling stack is not a hierarchy where higher layers are better. Different employers, industries, and career stages will weight the layers differently. A hospital hiring a surgeon will care overwhelmingly about Layer 1. A startup hiring a developer will focus on Layer 3. A large corporation filling a mid-career role might weight Layer 2 and Layer 4 most heavily. Most people will navigate a stack, not rely on a single credential.

Degrees Are Not Dead – Yet

The narrative that degrees are dying makes for compelling conference talks. The reality is more textured.

Where Portfolios Already Dominate

In software engineering, design, data science, and parts of the creative industries, portfolios and demonstrated work have become primary hiring signals. Many leading tech companies – including Google, Apple, and IBM – have publicly dropped degree requirements for a significant share of roles. The 42 School model, now spanning 54 campuses across 31 countries, awards certificates rather than formal degrees; employers report that its graduates excel at independent problem-solving and are less reliant on supervisor direction.

In these fields, the evidence layer of the signalling stack already outweighs the baseline layer. What matters is what you have shipped, not where you studied.

Where Degrees Remain Hard Gates

In medicine, law, engineering, nursing, accounting, teaching, pharmacy, and architecture, degrees are not preferences – they are legal requirements. Professional licensing bodies, government regulations, and liability insurance frameworks mandate specific educational credentials. These requirements exist because the consequences of incompetence in these fields are severe and often irreversible.

Changing these gates requires regulatory reform, not just employer preference shifts. It requires changes to professional body governance, liability and insurance models, and public trust frameworks. These changes happen slowly – measured in decades, not product cycles.

The Messy Middle

Most professions fall between these extremes. Business, communications, marketing, human resources, project management – these fields nominally require degrees, but the requirement is softening. Skills-based hiring expands the talent pool, but many employers still use degree requirements as a first filter, even when they say they do not.

The gap between stated policy and actual practice is significant. Companies announce skills-based hiring initiatives; hiring managers still sort resumes by university name. The cultural inertia of credentialism runs deeper than corporate policy changes can quickly reach.

How Large Is the Degree-Free Labour Market?

The honest answer is: smaller than the skills-based-hiring rhetoric suggests, and larger than sceptics assume.

Across OECD economies, roughly 20–25% of employment sits in regulated or licensed professions where specific degrees remain legal requirements – medicine, law, engineering, nursing, teaching, pharmacy, accounting, architecture, and their adjacent fields. For these roles, the Signalling Stack does not displace the degree. It supplements it. A surgeon with an excellent portfolio and strong behavioural signals still needs a medical degree. This share of the labour market is unlikely to change substantially within a decade, though competency-based pathways may begin to open at the margins.

At the other end, perhaps 5–8% of employment sits in sectors where portfolio-based hiring is already the dominant practice – primarily software development, UX/UI design, data science, parts of the creative industries, and some areas of digital marketing. Here, the evidence layer of the Signalling Stack already outweighs the baseline layer. This share is growing, but it remains a small fraction of total employment.

The remaining 65–75% is the messy middle: business administration, sales, human resources, project management, communications, finance (non-regulated roles), operations, and the broad landscape of knowledge work where degrees are expected but not legally required. This is where the real contest plays out – and where the gap between employer rhetoric and employer practice is widest.

Seniority compounds the picture. At entry level, where candidates have little track record, degrees function as the cheapest available filter. The Signalling Stack is weakest here precisely because the other layers – skills, evidence, behaviour – have had the least time to accumulate. Mid-career professionals, by contrast, are rarely hired on their degree; track record, network, and demonstrated capability dominate. For senior roles, degrees function as pedigree signals (which school, not which subject) rather than competence signals, and their weight in hiring decisions is marginal compared to reputation and relationships.

The implication is uncomfortable but necessary to state: the Signalling Stack primarily displaces the degree for mid-career professionals in unregulated knowledge work, and for entry-level candidates in a narrow band of portfolio-driven industries. For everyone else – entry-level candidates in the messy middle, workers in regulated professions, and anyone in a labour market where employer coordination around skills-based hiring has not yet reached critical mass – the degree remains either a hard gate or a rational default. The Signalling Stack supplements the degree far more often than it replaces it.

This does not mean the shift is trivial. Even partial displacement matters: when skills-based hiring expands candidate pools by 10–17 times in affected roles, the individual and economic impact is significant. But the claim that degrees are “dying” misstates the scale and the timeline. What is dying is the degree’s *sufficiency* – the idea that a degree alone signals everything an employer needs to know. Its *necessity*, in most of the labour market, remains intact.

The Credential Inflation Problem

The United States is now home to over one million unique educational credentials – a threefold increase since 2018. This explosion includes offerings from universities, industry providers, bootcamps, and platforms, creating what observers call “micro-credential inflation.”

When credentials proliferate faster than quality assurance can keep pace, the result is noise, not signal. Close to half of employers (46%) cannot assess the quality of a non-degree credential when a job applicant lists one. Forty-two percent cannot judge the skills and competencies acquired. A systematic review found that 80% of employers cite consistency as their primary concern: without standardised definitions, learning outcomes, or assessment rigour, micro-credentials vary wildly in value. Authenticity concerns follow: 38% of studies report employer worry about fraudulent credentials.

For learners, the landscape is equally confusing. Forty percent report not knowing where to start in the digital credential landscape. Sixty percent worry that costs will be out of reach despite marketing promises of affordability. The promise of stackability – earning smaller credentials that accumulate toward degrees – often fails in execution. Without thoughtful curricular design, stackable credentials become disjointed modules lacking coherent progression.

The irony is sharp: micro-credentials were supposed to solve the problem of degree inflation. Instead, we risk replacing one form of credential inflation with another, more fragmented one. Government financial aid rarely covers micro-credentials, creating a 23% barrier to adoption. Forty-three percent of the literature notes that micro-credentials deliver competencies limited to one specific niche of a job scope, restricting career mobility and adaptability.

Employer Confusion and the Trust Gap

The signalling stack only works if employers can read it. Right now, many cannot.

When a candidate presents a traditional degree, the employer has a rough mental model of what it means – imperfect, but shared. When a candidate presents a collection of micro-credentials, a GitHub profile, three online certificates, and a portfolio link, the

employer faces a parsing problem. Which of these are rigorous? Which are vanity certificates from a weekend course? How do they compare to one another or to a degree?

This trust gap is the central obstacle to the new signalling system. Blockchain verification solves the authenticity problem – you can confirm that a credential is real. But it does not solve the quality problem – you still cannot tell whether the credential is worth anything. The default employer response remains what researchers call “trust but verify”: employers ask their own questions to confirm knowledge rather than relying on the credential itself.

Without robust quality assurance frameworks, transparent metadata standards, and shared recognition systems, the signalling stack risks fragmenting into a confusing patchwork that serves no one well. The European approach – embedding micro-credentials within national qualifications frameworks, linking them to established quality assurance processes, and using blockchain for verification – offers one path forward. But implementation remains uneven across countries.

The Rational Inertia of Degree Requirements

The gap between “81% of employers use skills-based hiring” and “most still require degrees” is not a contradiction. It is rational. Understanding why reveals why the shift to portfolios is slower than the technology and the data would suggest.

From a hiring manager’s perspective, the decision tree looks like this:

Cost of screening by degree: Low. It is simple, fast, legally defensible, and based on decades of calibration. If you filter for candidates from accredited universities with a minimum GPA, you are following a standard that any candidate can understand and challenge. If you make a bad hire, you can point to the credential and claim you followed reasonable procedure. The credential does not require interpretation. Everyone knows what a degree from Stanford means.

Cost of screening by skills: High in three ways. First, it demands expertise. Interpreting a GitHub profile, a portfolio of projects, and performance on skills assessments requires judgment. This judgment is harder to systematize and easier to challenge as biased. Second, it demands infrastructure. You need assessment tools, people trained to interpret them, and time in the interview process to evaluate them. Third, and most important, it concentrates legal risk on the hiring manager. If you hire someone based on a portfolio and they fail spectacularly, you cannot point to an industry standard that you followed. You made a judgment, and if the judgment was wrong, that is on you in a way that degree requirements are not.

Cost of false positives: When screening by degree, the main risk is hiring someone with a credential who cannot perform. This risk is managed through other filters – interview quality, reference checks, trial projects. When screening by skills, the main risk is hired someone without a credential who cannot perform. This feels different. In a lawsuit or a regulator question, the question “Why did you hire someone without a degree for a role that traditionally requires one?” is harder to answer than “Why did you hire someone

with a degree who turned out to be incapable?” The former puts you on the defensive about your process; the latter puts you on the defensive about your judgment. Corporations have legal departments and insurance models built around the latter, not the former.

This is not theoretical. In regulated professions – law, medicine, engineering – degrees are hard gates not because they perfectly predict performance, but because they are the legally defensible basis for gatekeeping. You cannot open yourself to the liability of hiring an unaccredited lawyer, even if that person might perform well.

In less regulated fields, the pattern is subtler but real. A hiring manager at a bank or a Fortune 500 company who eliminates degree requirements is making an individual decision to take on individual risk. A hiring manager who maintains degree requirements is following industry practice and can share the risk across an entire sector. Until the industry standard shifts, individual incentives point toward maintaining the status quo.

This means that skills-based hiring will not replace degree requirements until one of three things happens:

27. **Risk is pooled at the industry level.** If a critical mass of major employers in a sector simultaneously commit to hiring without degree requirements and report positive outcomes, the decision shifts from individual to collective. At that point, an individual hiring manager can say “our industry has moved to skills-based hiring, and we follow that practice.” This creates legitimacy and shared liability. Tech is approaching this threshold in some roles; most sectors are far from it.
28. **Skills assessment becomes as legible as degrees.** Right now, interpreting a portfolio is harder than reading a resume. If a standardized, transparent, industry-recognized system emerged for certifying skills – something more credible than the current fragmented landscape of micro-credentials and bootcamp certificates – then the cost of screening by skills would drop significantly. Blockchain-based credentials with immutable metadata, combined with widespread employer training on interpretation, could move the needle here, but it requires sustained coordination that has not yet emerged.
29. **The cost of maintaining degree requirements becomes visible and high.** This happens when employers explicitly demonstrate that degree filters are excluding high-performing candidates or when the talent shortage becomes acute enough that degree requirements become economically painful. Europe is approaching this in tech – the unfilled IT positions are not hypothetical, and the pain is driving regulatory and employer-level coordination toward skills-based alternatives. But in most of the labor market, the pain is not yet visible enough to justify changing a process that works well enough.

The implication is that the timeline for shifting from degree-based to portfolio-based hiring is not driven by the capability of assessment technology. It is driven by the coordination problem of shifting industry standards and the regulatory timing of reducing legal risk. In technology, where this coordination is most advanced and the pain is most acute, the shift is visible and ongoing. In most other sectors, it is slower and more patchy. This is not because hiring managers are irrational or stubborn. It is

because the rational incentives that have made degree requirements standard are still very much in place.

The shift from “we plan to adopt skills-based hiring” to “we actually do it at scale” requires not just better tools. It requires a change in the risk calculus, a change in industry norms, and often a change in regulation. All three are possible, but none is automatic.

Measurement, Not Gatekeeping: Why the Signalling Stack Requires Better Data

Earlier in this book, we criticised standardised testing for measuring narrow abilities rather than genuine cognitive growth – for creating systems that optimise learning for what is easily measurable rather than what is genuinely important. That critique remains valid. But it requires precision.

The problem with state-mandated standardised testing is not that it measures, but that it measures *once, at a gate*. A single snapshot of performance, taken at a fixed moment and used to sort students into trajectories they may follow for years, inevitably narrows what schools teach and what students learn. Teachers feel pressure to teach to the test. Students experience assessment as judgment rather than feedback. The measurement apparatus becomes the curriculum.

The opportunity with AI-driven continuous assessment is fundamentally different. It measures *iteratively, in real time, as part of the learning process itself*. An adaptive learning platform that adjusts difficulty based on performance is not taking a standardised test; it is using moment-to-moment data to personalise instruction. A micro-credential that certifies mastery of a specific skill is not sorting students into fixed tracks; it is creating a granular record of what someone can actually do. These are not the same thing.

But the risk is real: AI measurement can replicate the narrow-ness of standardised testing at a much larger scale and with far less transparency. If we use AI analytics to decide who gets promoted, who gets flagged as “at risk,” or what learning path someone should follow, without auditing the systems for bias and without maintaining human oversight, we have not fixed the standardisation problem – we have just made it invisible.

The solution is not to reject measurement. The solution is to insist on three conditions: First, measurement must be *continuous*, not a single gate – giving learners and educators feedback about growth rather than imposing judgment at a fixed moment. Second, measurement must be *transparent* – learners and educators must understand what is being measured, how, and why, with particular vigilance against algorithmic systems trained on historical data that may encode past discrimination. Third, measurement must be *human-readable* – the data must inform human judgment, not replace it.

When these conditions are met, the signalling stack becomes legible: employers, educators, and learners can all understand what a micro-credential, a portfolio, a

learning velocity score, or a behavioural signal actually means. When they are not met, we have simply replaced crude mass standardisation with sophisticated, opaque algorithmic sorting.

The systems we describe in this book – AI-personalised learning, micro-credentials, continuous assessment – can genuinely improve on the narrow testing regime we critique. But they can only do so if they are built with deliberate constraints against the very narrowing dynamic we criticised in the first place.

Evidence vs Conjecture: The Signalling Stack and the Demotion of Degrees

The claim: The degree is shifting from a gate to a preference, becoming one layer in a multilevel “signalling stack” where employers increasingly weight skills, portfolios, and AI-derived behavioural signals alongside or instead of traditional credentials.

What supports it:

Layer 1 — Observed employer behaviour (Tier 2–3): - 81% of employers globally report using skills-based hiring approaches (up from 73% in 2023), per TestGorilla and LinkedIn surveys. This is observational data from large samples, but the survey definitions and validity vary. - 77% of European companies plan to prioritise competency-based recruiting in 2026, per recruiters' stated intentions. This is forward-looking assertion by employers, not yet observed practice. - Skills-based hiring expands candidate pools by 10–17x (UK, US, Gen Z), per TestGorilla. This assumes that candidates without traditional degrees are actually being hired; follow-on hiring data is limited. - **Critical counter-evidence:** A 2024 Burning Glass Institute / Harvard Business School study found that removing degree requirements from job postings increased the share of hires without a bachelor's degree by only ~3.5 percentage points per affected role – fewer than 1 in 700 new hires benefited from these reforms (Sigelman, Fuller, and Martin 2024). This is the strongest available evidence that employer behaviour lags rhetoric by a wide margin, and it directly qualifies the optimism of the self-reported survey data above. - 41% of talent acquisition teams use AI-powered skills assessments; 83% use gamified assessments. These are adoption metrics (Tier 3–4), not proof that these tools outperform traditional hiring or produce better outcomes.

Layer 2 — Existing portfolio-driven hiring (Tier 3): - Software development, design, and data science already show portfolio-first hiring practices (single-firm case studies and industry observation, Tier 3). - Tech giants (Google, Apple, IBM) have publicly dropped degree requirements. These are institutional decisions by specific companies; we lack systematic data on whether this is reversing, deepening, or stalling. - The 42 School model spans 54 campuses; employers report graduates excel at independent problem-solving (Tier 3, single-provider data).

Layer 3 — Emerging infrastructure (Tier 3): - European blockchain credential pilots (ECIU, Blockcerts, Europass) demonstrate technical feasibility (Tier 3, pilot implementation). - European Commission joint degree label rolling out mid-2026 is announced policy, not yet evaluated (policy intent, pre-implementation).

Layer 4 — The speculative layer: - The four-layer signalling stack framework itself is the authors' synthesis, not drawn from published empirical research. - The 5-, 10-, and 20-

year scenarios are plausible extrapolations of current trends, not predictions based on forecasting models or historical precedent. - The claim that behavioural signals could become a primary hiring signal is hypothetical; no data shows this is currently happening at scale.

What would change it:

- If longitudinal employer hiring data (5–10 years) showed that skills-based hiring has plateaued or reversed, or that degree requirements have rebounded, the trend narrative would need rethinking.
- If the credential inflation problem (1 million+ credentials in the U.S., 46% of employers unable to assess them) worsens rather than resolves, the signalling stack risks fragmenting instead of maturing.
- If regulated professions (law, medicine, nursing, engineering) do not introduce competency-based pathways within 15 years, the demotion of degrees remains confined to already-flexible sectors.
- If AI assessment tools do not outperform traditional hiring on standard labour-market outcomes (retention, performance ratings, diversity), the adoption of these tools may be vendor-driven rather than outcome-driven.
- If privacy regulation or bias scandals slow the deployment of behavioural analytics in hiring, the Layer 4 scenario becomes less plausible.

Confidence level: - The 5-year scenario (hybrid world with skills-based supplements to degree requirements) is well-grounded in current observable trends. - The 10-year scenario (maturation of quality assurance frameworks, blockchain wallets, possible regulatory reform) is a plausible extrapolation, contingent on sustained investment and policy alignment. - The 20-year scenario (degree's monopoly substantially weakened, behavioural signals as primary hiring signal) is speculative and contingent on major structural changes (regulatory reform, cultural shift, algorithmic trust) that are not guaranteed.

Time Horizons: What Changes When

The following projections explore plausible trajectories; they are not certainties but depend heavily on which of three possible futures unfolds. The book's final chapter explores three competing scenarios: Drift (incremental change with persistent inequities), Platform World (institutional concentration and algorithmic control), and Redesign (intentional institutional restructuring with public accountability). The signalling stack takes radically different forms in each.

Short Term (Next 5 Years): Multiple Possible Paths

In Drift: Degrees remain the default gate. Skills-based hiring grows as rhetoric but adoption is inconsistent. Micro-credentials proliferate without standardization, creating confusion. Employers become frustrated with the signal-to-noise ratio and retreat to familiar filters: degree name, referrals, pedigree. The signalling stack fails to cohere; instead, it fragments into silos of proprietary credentials issued by platforms and employers that workers cannot easily transfer or stack.

In Platform World: Large tech companies and EdTech platforms establish proprietary credential systems. Employers adopt platform-based hiring and assessment at scale. The signalling stack collapses into a monolithic (or oligopolistic) platform ecosystem. Micro-credentials are meaningful only if they are platform-native. Portfolios matter only if they are visible within the platform. Talent matching is fast and efficient, but workers lack portability and ownership of their own data. Credential quality is high because a platform risks its reputation on the strength of its credentials, but workers are locked in: switching platforms means credentials lose value.

In Redesign: Governments and professional bodies begin establishing shared frameworks for credential recognition. European blockchain initiatives move beyond pilots. Micro-credentials gain traction and employers begin to distinguish rigorous credentials from vanity certificates through quality assurance frameworks. The signalling stack becomes more legible, though integration is still incomplete. Blockchain credential wallets begin early adoption. The foundation is laid for cross-border, multi-source credential accumulation.

Medium Term (Next 10 Years)

In Drift: The credential landscape remains fragmented and difficult to navigate. Micro-credentials proliferate further, often without employer demand or clear market value. Degree requirements persist as the default filter because the alternatives are too noisy. A learner or worker faces dozens of possible credentials and no clear way to assess which matter. Employers revert to hiring based on pedigree and networks rather than engaging with the signalling stack. The hoped-for transparency never materializes; instead, insiders win (those from well-known schools or networks) while outsiders face a bewildering landscape.

In Platform World: Platform ecosystems mature and consolidate. A handful of large companies (likely in tech, recruiting, or EdTech) come to dominate talent assessment and matching globally. Blockchain credentials, if adopted at all, are proprietary to these platforms. Behavioural data collection expands. Workers have detailed profiles, but they own none of the data. Internal mobility tools mean many roles are filled from within platforms' user ecosystems, creating lock-in. Regulatory attempts to ensure data portability face industry resistance and are unevenly enforced. The signalling stack is legible within each platform, but moving between platforms is costly.

In Redesign: Quality assurance frameworks for micro-credentials mature, driven by European regulatory leadership and employer demand for consistency. The European Commission's joint degree label becomes a recognized standard. Blockchain-based credential wallets become standard infrastructure, enabling learners to accumulate verified credentials from multiple sources – universities, employers, platforms, professional bodies – in portable digital portfolios that AI systems can parse automatically. Behavioural signals may enter mainstream hiring, but with governance frameworks that protect privacy and enable worker data access. The signalling stack becomes legible and interoperable, though not perfectly integrated.

Long Term (Next 20 Years)

In Drift: The degree retains its monopoly as the primary signal, despite periodic announcements of skills-based alternatives. The promised unbundling never happens. Micro-credential inflation worsens, creating so much noise that employers retreat to what they know. Alternative credentialing systems fail to gain employer confidence because there is no coordination around standards. Inequality persists and arguably widens: students with family resources and networks use traditional degrees plus micro-credentials plus personal networks plus social capital to navigate the system. Students without these resources follow a narrower path dictated by employer convenience. The intelligence factory exists, but does not democratize access.

In Platform World: Platform companies have consolidated talent assessment and matching into a handful of proprietary ecosystems. Traditional universities persist for elite education and research, but most workforce development is routed through platform systems. Workers experience continuous algorithmic evaluation tied to employment. The signalling stack is highly efficient, but it serves corporate optimization, not learner growth. Behavioural signals form a primary layer of hiring decisions, but they are controlled by private companies optimizing for organizational productivity rather than worker wellbeing. The degree has been demoted, but not to a more democratic system – instead, to a more concentrated one where algorithmic gatekeeping replaces institutional gatekeeping.

In Redesign: The degree's monopoly on professional signalling substantially weakens outside regulated professions. The signalling stack matures into a broadly understood, interoperable system where different layers are weighted by context. Hiring a surgeon still requires Layer 1. Hiring a product manager might rely primarily on Layers 2, 3, and 4. Micro-credentials have become standardized across borders, enabling mobility. Portable blockchain credentials are routine. AI assessment tools are used, but with strong governance frameworks, data ownership by workers, and regular audits for bias. Regulated professions have adapted, introducing competency-based pathways that coexist with (or partially replace) traditional degree requirements. The signalling stack becomes a ladder rather than an obstacle course.

For degrees to truly recede beyond regulated professions, several structural changes must occur across the Redesign path: regulatory reform of professional licensing, adaptation of liability and insurance models, shifts in public trust frameworks, and cultural change in how families, communities, and institutions value credentials. These are generational changes, not policy tweaks.

The most transformative – and most speculative – possibility in Redesign is that AI-derived behavioural data becomes reliable and trusted enough to serve as a primary signal for many roles, effectively replacing the degree's function as a proxy for cognitive ability and conscientiousness with direct measurement. This would represent a fundamental shift in how societies sort people into opportunities. Whether it leads to greater meritocracy or to new forms of algorithmic discrimination depends entirely on the governance frameworks we build around it.

The trajectory that unfolds will depend not on technological capability but on the choices institutions, employers, governments, and workers make in the next five to ten

years about who controls the signalling infrastructure, what happens to the data it generates, and whether public goods are invested in shared standards or left to proprietary platforms. These are not technical questions. They are governance and power questions.

What This Means for You

Whether you are a student choosing a path, a mid-career professional navigating change, or an employer redesigning how you hire, the signalling stack has practical implications.

- **If you are a student:** A degree still matters, but it is no longer enough. Start building your evidence layer early – projects, contributions, documented work that shows what you can do. Choose micro-credentials strategically to complement your degree, not to replace it. Develop the habits that will generate strong behavioural signals: consistency, collaboration, willingness to learn publicly.
- **If you are mid-career:** Your degree got you in the door years ago. What keeps you relevant is your skill and evidence layers. Invest in current certifications, maintain a visible body of work, and demonstrate that you are still learning. The half-life of skills is shrinking; a four-year degree from fifteen years ago tells employers less than what you built last quarter.
- **If you are an employer:** Stop asking whether to require degrees or not. Start asking which layer of the signalling stack matters most for each role. Build assessment processes that can read multiple layers. Invest in quality frameworks that help your hiring managers distinguish rigorous micro-credentials from vanity certificates. And recognise that the transition will be messy – the old system was imperfect but legible, and the new one is potentially better but currently hard to read.
- **If you are a policymaker:** The signalling stack will not build itself into a coherent, trustworthy system without regulatory infrastructure. Quality assurance frameworks for micro-credentials, interoperability standards for digital wallets, and updated professional licensing pathways are public goods that markets alone will not provide. The European approach of embedding new credentials within existing qualification frameworks offers a model, but it needs sustained investment and cross-border coordination.

Conclusion

The degree is not dead. But its monopoly is ending. We are moving from a world with one dominant signal to a world with a layered signalling system – degrees, skills, evidence, and behaviour – where different layers matter in different contexts.

This transition carries real risks. Credential inflation, employer confusion, and fragmentation could produce a system worse than the one it replaces – noisier, more expensive, and more unequal. The United States already has over one million unique credentials, and most employers cannot tell them apart. Without shared standards,

transparent quality assurance, and regulatory frameworks that keep pace with innovation, the new signalling stack could become a tower of noise rather than a ladder of opportunity.

But the direction is set. Employers need better information about what candidates can do. Learners need more flexible, affordable, and current ways to demonstrate competence. AI is making both verification and assessment cheaper and faster. The question is not whether the signalling stack emerges, but whether we build it well – legible, fair, and genuinely useful – or let it fragment into chaos.

The institutions, employers, and policymakers who understand this transition will shape it. Those who cling to the old monopoly of the degree or rush uncritically toward its replacement will find themselves on the wrong side of a structural shift that is already well underway.

Annotated Source List

The arguments draw on - OECD comparative education statistics (Tier 2), - large-scale employer surveys on skills-based hiring from TestGorilla and LinkedIn (Tier 2–3), - World Economic Forum workforce reports (Tier 2), - European Commission policy documents on micro-credentials and blockchain credential frameworks (Tier 3), - industry analyses of credential inflation from Credential Engine and the Lumina Foundation (Tier 3), - and vendor-reported data on AI assessment platforms and gamified hiring tools (Tier 4). - The rhetoric-vs-reality gap in skills-based hiring is documented by Sigelman, Fuller, and Martin (2024), *Skills-Based Hiring: The Long Road from Pronouncements to Practice*, Burning Glass Institute / Harvard Business School (Tier 2: longitudinal analysis of actual hiring outcomes from job postings and employment records, not self-reported surveys). <https://www.burningglassinstitute.org/research/skills-based-hiring-2024>. - Evidence on the degree premium relies on cross-national labour force surveys and econometric studies of returns to education. - Discussion of blockchain credentials references pilot implementations by the European Consortium of Innovative Universities and the Europass digital wallet initiative. The signalling stack framework is the authors' synthesis and is not drawn from a single empirical source.

The Stratification Trap: Who Gets Left Behind

Every innovation in this book – AI tutors, peer-driven coding schools, employer-sponsored degrees, transnational university alliances – promises to widen the door to learning. And every one of them, if built without deliberate safeguards, risks slamming that door on the people who need it most. This chapter moves equity from the margins to the centre of the argument. Not as an afterthought, not as a risk appendix, but as the core design constraint that determines whether any of these models actually delivers on its promise.

The pattern is consistent and uncomfortable: the same features that make a new model exciting for high-agency, well-resourced learners make it treacherous for everyone else.

In earlier chapters, we described these models and their potential. Here, we stress-test them with a single question: **who benefits, and who is excluded?** The answers are not comfortable, but they are necessary. Because if we build the future of education without confronting these questions honestly, we will end up with a system that looks more modern but works the same way the old one did – sorting people by background and calling it merit.

1. When More Opportunity Means More Inequality

There is a seductive logic to the AI education revolution: if every student on the planet can access a personal tutor that speaks their language, adapts to their pace, and never loses patience, then surely the playing field levels itself.

Knowledge is no longer scarce. The lecture hall bottleneck disappears. Anyone with a phone and a connection can learn calculus, molecular biology, or contract law from a system that rivals the best human instructors.

But abundance does not automatically produce equality. It often produces the opposite.

Consider what happened when the internet first made information freely available. The students who benefited most were those who already knew how to find, filter, and apply information – those with strong reading skills, stable home environments, supportive adults, and the discipline to sit with difficult material.

Students without those advantages drowned in the noise. The gap widened, not because the technology was bad, but because it rewarded capacities that were unequally distributed in the first place.

AI tutoring platforms repeat this dynamic at a higher level. When a motivated, well-prepared student uses an AI tutor, they treat it as a thought partner: they ask follow-up questions, challenge the AI's reasoning, request harder problems, and integrate what they learn into projects.

When an under-prepared or struggling student encounters the same tool, they may accept surface-level answers, use it to avoid thinking rather than to deepen it, or simply disengage when the novelty wears off.

The technology is identical. The outcomes diverge because the human inputs – motivation, prior knowledge, self-regulation, and environmental support – are profoundly unequal.

This is not speculation. Chapter 3 documented the pattern with MOOCs: free, open tools whose benefits flowed overwhelmingly to learners who already held degrees (Reich and Ruipérez-Valiente 2019). AI tutoring platforms are far more sophisticated, but the underlying dynamic remains. A tool that requires the user to bring direction, persistence, and critical judgment will systematically advantage those who already have these qualities – and those qualities correlate strongly with socioeconomic background, parental education, and the quality of prior schooling.

Affordability challenges and high debt burdens already disproportionately affect lower-income and marginalized students, undermining equity and social mobility goals (OECD 2025a). Layering powerful AI tools on top of these existing inequalities without addressing the underlying conditions risks creating a world where the privileged learn faster than ever while everyone else falls further behind.

This is the paradox at the heart of educational technology: **more opportunity, distributed without support, becomes more inequality**.

The solution is not to withhold the technology. It is to accompany it with the human infrastructure – coaching, mentoring, structured support, mental health services, and community – that turns raw access into genuine learning.

This book’s position on AI tutoring is therefore conditional, not ambivalent. Chapter 4 described the genuine promise: personalised, scalable, always-available instruction that could democratise access to high-quality teaching. This chapter documents the risk: that the same tools widen inequality when deployed without support. These are not contradictory claims. They describe the same technology under different institutional conditions. AI tutoring narrows inequality when accompanied by the agency-building infrastructure Chapter 3 defined – scaffolded progression, metacognitive instruction, human mentorship, and mental health support. It widens inequality when deployed as a cost-cutting substitute for that infrastructure. The technology is not the variable. The institutional investment is.

Without that investment, we are handing everyone a map to a city and pretending that the person without shoes is on equal footing with the person in a car.

Given these dynamics, we argue that policymakers and university leaders who celebrate “democratised access” to AI learning tools without simultaneously funding the support structures that make access meaningful are, at best, naive. At worst, they are using the language of equity to justify cost-cutting – replacing expensive human teaching with cheap AI delivery and calling it progress.

2. The Agency Divide

The previous section described the paradox of abundance. This section names the mechanism behind it.

Chapter 3 defined agency as the socially produced capacity to direct one’s own learning – to set goals, persist through difficulty, and adapt when things do not work. It

established that agency is not a personality trait but a product of upbringing, schooling, and environment, distributed along familiar lines of class, race, and geography. And it identified why agency has become the primary bottleneck: as education shifts toward flexible, self-paced, AI-mediated models, the external scaffolding that traditional systems provided – compulsory attendance, fixed schedules, teacher-imposed deadlines – disappears. The students who relied on that scaffolding most are the first to fall through the cracks.

The socioeconomic implications are direct. When AI makes high-quality education cheaper and more available, it may increase inequality based on *agency* rather than access. The divide falls between those who use AI to avoid thinking and those who use it as a thought partner to accelerate their learning velocity (Huang 2024).

In a world where every student has a personal AI tutor, the student with a stable home, a clear goal, and a parent who asks about their day will use that tutor to master organic chemistry in half the time. The student dealing with food insecurity, an unstable living situation, and untreated anxiety will use the same tutor to get through homework as fast as possible – or not use it at all.

Both had “access.” Only one had the conditions to convert access into learning.

The mental health dimension compounds this. As Chapter 3 documented, nearly half of U.S. high school students report persistent sadness or hopelessness, and many institutions lack the capacity for comprehensive wellbeing services (Deloitte 2025). A student battling depression or anxiety does not lack intelligence. They lack the internal bandwidth that self-directed learning demands. When models assume agency as a given, they build for the students who need the least help and exclude the ones who need the most.

This is not an argument against AI in education. It is an argument that AI in education without investment in the human conditions for learning is a recipe for wider gaps dressed up as progress.

Who Wins / Who Loses? – AI Tutoring Platforms

Who wins: Self-directed learners with stable environments, strong foundational skills, and the discipline to use AI as a thinking partner. Adults returning to education with clear career goals.

Who loses: Students from under-resourced backgrounds with weaker self-regulation skills. Teenagers struggling with mental health. First-generation students without role models for independent learning. Anyone who needs a human to notice they are falling behind.

3. Why “Self-Directed Learners” Are Often the Already Privileged

Several of the most celebrated alternative education models – profiled in Chapter 6 – share a common design assumption: that learners will take charge of their own education. Peer-driven schools, campus-free global programmes, and competency-based online degrees all rely on students who can set their own pace, seek out resources, collaborate without being told to, and persist through ambiguity.

This sounds empowering. In practice, it filters for privilege.

The reason is straightforward: “self-directed learning” requires a foundation that is itself socially produced. You need to know how to learn before you can direct your own learning. You need confidence that your efforts will pay off. You need to have seen adults model persistence, curiosity, and self-discipline. These are not innate traits. They are the products of upbringing, schooling, and environment – and they are distributed along familiar lines of class, race, and geography.

Chapter 6 assessed each model against the agency criteria established in Chapter 3, and the stratification pattern is consistent:

- **42 School** selects for agency through its pipeline and does not build it. The model scales horizontally – more campuses – but not inclusively. Celebrating 42 as revolutionary while ignoring who it structurally excludes is a form of intellectual dishonesty. If 42 wants to be a niche institution for self-starters, that is legitimate. But it should not be held up as a model for the future of education at large.
- **Minerva University** demands high agency from day one without addressing the upstream barriers that prevent talented learners from lower-income backgrounds from becoming “qualified” in the first place. A cheaper version of elite education is still elite education.
- **Competency-based education** programmes like WGU privilege learners who are already self-directed and digitally literate. The models that claim the best results are also the ones with the strongest selection effects – they work well for the students who would likely have succeeded anywhere.

The uncomfortable truth is that “self-directed learning” is often code for “learning that requires no institutional investment in support.” It shifts responsibility from the institution to the individual, and that shift systematically disadvantages those who need institutions most.

There is nothing wrong with self-direction as an educational goal. The problem arises when it is treated as an entry requirement rather than a learning outcome – a distinction Chapter 3 examined in detail. The most equitable approach would be to design models that develop agency in learners who do not yet have it, rather than selecting only those who already do.

That is more expensive. It is more complicated. And it is the only way these models can honestly claim to be expanding opportunity rather than repackaging privilege.

Who Wins / Who Loses? – Peer-Driven and Campus-Free Models

Who wins: Highly motivated, digitally fluent young adults from middle-class or affluent families with strong secondary-school preparation. People who thrive in unstructured environments.

Who loses: First-generation students who need mentorship. Learners from the Global South facing cost barriers even at “low-cost” institutions. Students who need structured deadlines, explicit instruction, and human guidance to succeed. Anyone without reliable internet, a quiet workspace, or a supportive household.

4. When Companies Own Your Learning Path

The previous sections examined models where the learner is expected to bring their own agency, motivation, and resources. Now we turn to a model that seems to solve the financial problem entirely: the employer pays.

Guild Education's model – employers pay tuition upfront for workforce upskilling – addresses a real financial barrier. Workers who could never afford college get access to degree programs and certificates without taking on personal debt. On the surface, this looks like a pure win for equity.

Look closer and a different picture emerges.

Over 80% of degree and certificate learners in Guild's system enroll in business-aligned programs selected to meet employers' strategic workforce needs (Contrary Research 2024). Employees choose from a curated set of offerings rather than open access to all higher education options. The employer, not the learner, sits at the centre of decision-making.

This arrangement creates three interrelated problems that are easy to miss when the headlines focus on "free college for workers."

- **Curriculum narrowing.** This is the most fundamental concern, and it goes beyond individual workers. When corporations fund education, corporations shape education. Programs tilt toward instrumental, skills-focused training at the expense of liberal arts, critical inquiry, and the broader intellectual development that helps people adapt to careers that do not yet exist. Universities dependent on Guild partnerships may face pressure to design programs that prioritize employer satisfaction over intellectual rigor or student-centred pedagogy, risking mission drift that transforms universities from sites of broad human development into corporate training arms (Müller 2017).
- **Worker lock-in.** Guild participants are 2.6 times more likely to remain with their company compared to non-participants, and internal job mobility increases by 3.5 times. Chipotle reports that Guild enrollees are nearly twice as likely to be promoted and have 89% retention rates after nine months (PR Newswire 2017).

From the employer's perspective, this is excellent return on investment. From the worker's perspective, it may reduce bargaining power by tying educational benefits to continued employment.

Workers who leave before completing programs may forfeit progress or face financial penalties, creating a form of economic lock-in (Rogers 2024). The education benefit becomes a golden handcuff – valuable precisely because leaving is costly.

- **Restricted agency.** Unlike traditional financial aid, Guild's business-to-business model places employers, not learners, at the centre. Employees choose from what their employer makes available, not from the full landscape of higher education. A warehouse worker passionate about philosophy or marine biology will find those options absent from the menu. The learning path is shaped by corporate strategy, not personal aspiration (Contrary Research 2024).

There is a deeper question here about what education is for.

If the only education a working-class person can access is the education their employer finds useful, we have not expanded opportunity – we have replaced one gatekeeper (the expensive university) with another (the corporation). The form of exclusion changes, but the power imbalance persists.

And unlike a university, a corporation has no public mission, no obligation to teach critical thinking, and no incentive to help workers imagine futures beyond their current employer's needs. A society where education for the wealthy is broad, exploratory, and empowering, while education for workers is narrow, instrumental, and employer-controlled, is not a society that has solved the education problem. It is a society that has repackaged it.

Who Wins / Who Loses? – Employer-Sponsored Models

Who wins: Workers in large companies with active Guild partnerships who want credentials aligned with their current career track. Employers who get a more loyal, internally mobile workforce.

Who loses: Workers who want to study something their employer does not value. Anyone who wants to leave their employer mid-program. Workers at small companies or in the gig economy with no access to employer-sponsored education. The broader public interest in education that serves democracy and citizenship, not just productivity.

5. The Matthew Effect in University Alliances

So far, this chapter has focused on models that affect individual learners. Now we turn to a structural problem that operates at the institutional level – and that shapes which universities thrive and which fall behind.

The European Universities Initiative, launched in 2017 and now comprising over 60 alliances, was designed to enhance cooperation across borders and strengthen European higher education as a whole. In practice, it reproduces the pattern scholars call the “Matthew effect” – a phrase from sociology meaning that initial advantages compound over time. The rich get richer.

Over 75% of member institutions rank among the top 500 globally, and most alliances build on pre-existing partnerships. Institutions with international networks, co-funding capacity, and strong reputations secured alliance membership; those without – particularly demand-absorbing lower-tier universities in Central and Eastern Europe – were disproportionately excluded (Rensimer and Brooks 2024).

The partial reallocation of Erasmus+ and Horizon Europe funds to predominantly Western European universities consolidates their competitive advantages. Smaller institutions face higher relative costs to participate – staff time, travel, co-funding requirements – while receiving proportionally smaller benefits. Institutions in countries with limited national support for internationalization struggle to compete for alliance membership, deepening regional inequalities (Taylor and Francis 2024b).

The result is a two-tier European higher education landscape: alliance members accumulate prestige, funding, and mobility opportunities, while non-members fall

further behind. The initiative services excellence and competitiveness more than inclusion and cooperation, contrasting with other European Commission strategies (such as Teaming and Twinning) that explicitly target peripheral regions and less-resourced universities (Rensimer and Brooks 2024).

Consider what this looks like on the ground. A university in Portugal or Estonia that wants to join an alliance must free up staff for international coordination, fund travel, meet co-funding requirements, and produce applications in a competitive process designed by and for institutions that already have these capacities. The cost of participation, relative to budget, is far higher for a mid-tier Eastern European university than for a well-funded Dutch or German one. The benefits – prestige, research funding, student mobility – flow disproportionately to those who were already ahead.

For a student in Romania or Bulgaria, this means that the most innovative, best-funded, most internationally connected learning experiences are concentrated in institutions they are least likely to attend. The alliance system, intended as a vehicle for cooperation, instead reproduces existing hierarchies (PMC 2024a).

The irony is sharp: an initiative branded as “European cooperation” has, in practice, widened the gap between European regions. The universities that needed cooperation most – those in under-funded systems with limited international experience – are the ones least likely to benefit from it.

This matters beyond Europe. The same dynamic – well-resourced institutions capturing the benefits of cooperation while less-resourced ones fall further behind – plays out in every transnational education initiative. Networks, alliances, and partnerships are powerful tools for improvement. But they amplify whatever inequality exists at the point of entry. If the starting conditions are unequal and the initiative does nothing to address that, the outcome will be more inequality, not less.

Who Wins / Who Loses? – European University Alliances

Who wins: Students and faculty at well-resourced Western European universities that were already internationally connected. Institutions that can afford the co-funding and staff time to participate.

Who loses: Universities in Central and Eastern Europe that lack the resources to join alliances. Students in peripheral regions who miss out on mobility, joint programs, and network effects. Smaller institutions whose relative disadvantage grows as alliance members pull ahead.

6. Digital Divides and Mental Health

Beneath every model discussed in this book lies a set of basic requirements that are easy to take for granted and devastating when absent.

A reliable internet connection. A functioning device. A quiet place to work. A mind that is well enough to concentrate.

These are not luxuries. They are prerequisites. And they are distributed as unequally as income itself.

Every model discussed in earlier chapters – AI tutoring, peer-driven coding schools, campus-free universities, employer-sponsored programs, transnational alliances – assumes these basics are in place. When they are not, the model does not just underperform. It actively excludes.

- **The connectivity gap.** Europe's ageing populations and regional disparities mean some systems struggle with declining local cohorts while others manage massification and capacity constraints, often within tight public spending envelopes. Persistent inequalities by socio-economic status, migrant background, and region continue to challenge the promise of fair access in both public and private institutions (OECD 2025c; OECD 2025b).
- **The institutional gap.** Institutions struggle to modernize legacy systems, integrate learning platforms, and deliver flexible, hybrid or online learning at scale without fragmenting the student experience (Full Fabric 2024). Rapid adoption of AI and analytics brings opportunities for personalization and efficiency but raises ethical concerns about data privacy, algorithmic bias, and potential displacement of some academic or administrative roles (Congruence Market Insights 2024).
- **The wellbeing gap.** The mental health dimension compounds every other inequality.

Retention is challenged by students' financial stress, mental health issues, and disengagement, which increases dropout risk (SeatsOne 2025).

The research on who drops out is remarkably consistent across countries and systems. Older freshmen, minority students, and students from poorer families face higher dropout probabilities. Female students, students with closer peer ties, and those in institutions with greater institutional commitment are less likely to drop out (Frontiers in Education 2023).

What this means is that dropout is not random. It follows the same fault lines as every other form of inequality. And every model that makes learning more flexible, more self-paced, and more independent – without also investing in the support structures that keep vulnerable students enrolled – will see those fault lines widen.

- **The isolation trap.** Teens and young adults are forming emotional attachments to AI chatbots, often substituting them for real relationships; mental health professionals report rising social anxiety and loneliness as drivers. When universities replace human mentors with AI agents or chatbot tutors, they risk compounding isolation and undermining the relational infrastructure that young people need for cognitive and emotional development (AICerts 2025).
- **The algorithmic bias trap.** AI systems trained on historical data replicate and amplify existing inequalities. This is not a theoretical risk – it is a documented reality.

Predictive models incorrectly flag Black and Latinx students as “at risk” when they actually succeed (false-negative rates of 19% and 21%, respectively). Admissions algorithms favour applicants from privileged high schools. Facial recognition systems

trained on datasets skewed toward white males fail to accurately identify darker-skinned individuals.

A 2021 study found that 80% of AI systems in education showed some form of bias when not properly audited (Clausius Press 2025).

The sources of bias are multiple: biased training data, flawed problem framing (for example, equating ZIP codes with academic potential), and a lack of diversity in development teams that overlook edge cases affecting minorities (PMC 2021).

Predictive analytics that flag minority students disproportionately as “at risk” – even when they succeed – risk triggering interventions that stigmatize rather than support, reinforcing deficit narratives (Schiller 2024). The tool designed to help becomes a tool that labels, and the label sticks.

The promise of AI in education is personalization. But if the algorithms behind that personalization carry the biases of the past, they will personalize inequality – giving privileged students more challenging material while steering disadvantaged students toward remedial tracks they may not need.

This is not a problem that technical fixes alone can solve. Better data and better algorithms help. But the deeper issue is one of governance: who decides what the algorithm optimizes for, who audits its outcomes, and who is accountable when it gets things wrong. Right now, most educational AI systems operate with minimal external oversight. That has to change before these tools are deployed at scale.

These issues compound each other. A student with poor connectivity is more likely to be in a region with underfunded mental health services. A student with untreated anxiety is less likely to persist in a self-paced online program. A student who drops out is more likely to carry debt without a credential, worsening their financial situation and mental health in a vicious cycle.

The digital divide is not just about having a device. It is about having the full ecosystem – connectivity, hardware, software, digital literacy, emotional stability, and human support – that makes digital learning work.

Think of it as a stack. At the bottom is physical infrastructure: electricity, internet, a device. Above that is digital literacy: knowing how to navigate platforms, manage files, communicate online. Above that is academic preparation: reading comprehension, numeracy, study habits. Above that is emotional readiness: mental health, motivation, a sense of purpose. At the top is agency: the ability to direct your own learning.

Every innovative model in this book operates at the top of this stack. But if the lower layers are missing, the top layer is unreachable. A brilliant AI tutor is useless to a student with a broken laptop and no WiFi. A peer-driven coding school is inaccessible to a student who has never touched a terminal. A campus-free global university is out of reach for a student who cannot afford a plane ticket.

Without that full ecosystem, every innovation described in this book becomes another mechanism for sorting rather than supporting.

Who Wins / Who Loses? – AI-First Universities

Who wins: Students with reliable broadband, modern devices, digital literacy, and stable mental health. Institutions that can invest in platforms, cybersecurity, and faculty development.

Who loses: Students in rural or low-income areas with poor connectivity. Learners whose mental health makes sustained self-directed work difficult. Minority students who are mislabelled by biased algorithms. Institutions that cannot afford the infrastructure to implement AI responsibly.

Non-Negotiable Design Features

If the models in this book are going to deliver on their promises rather than deepen existing divides, they need corrective features built in from the start – not bolted on later as equity window dressing.

The following is not a wish list. It is a set of minimum requirements. Any institution or policymaker that adopts one of these models without these features is making a choice about who gets left behind – whether they admit it or not.

Operationalising Equity: Three Levels of Intervention

The design anchors in this chapter – equity as design constraint, digital baselines, mental-health infrastructure, disaggregated outcomes – require new money flowing to underfunded systems, not just better allocation of existing budgets. Here is why the problem is structural, not just technical, and what operationalisation requires.

Coaching systems cost money per student. Algorithmic bias audits require specialist staff and independent review. Mental health infrastructure – counselors, crisis services, peer support coordination – requires more staff-to-student ratios than current systems maintain. Agency-building curricula require smaller class sizes and more instructor contact time than lecture-hall models allow. These are not optional improvements bolted onto existing systems. They are built-in requirements if equity outcomes are non-negotiable.

Yet they are also costs that fall disproportionately on the institutions that serve the students most likely to benefit from them – lower-income students, first-generation learners, students in under-resourced regions. This creates a perverse dynamic: the students who most need the supports cannot attend the institutions that have the budgets to provide them.

Operationalising equity requires three policy-level mechanisms:

1. FINANCING MECHANISMS

Tier 1: Redistributive Funding Formulas

Public higher education funding must explicitly redirect resources from well-resourced to under-resourced institutions based on student demographics and institutional capacity. Not voluntary cross-subsidy – legal mandate. If a government funds higher education with public money, that money should be allocated such that institutions serving high percentages of low-income, first-generation, and minority students receive per-student funding no lower than institutions serving affluent populations.

In practice, this means funding formulas that:

- Weight student funding by socioeconomic status and educational background (e.g., 1.3x multiplier for first-generation low-income students).
- Account for regional cost differences but not regional wealth differences (do not penalise universities in lower-GDP regions).
- Reserve minimum percentages of competitive research funding for mid-tier universities so that prestige and funding do not concentrate entirely at the top.
- Tie institutional expansion or fee revenue to demonstrated recruitment and support of disadvantaged students.

This will be resisted by well-funded institutions. They will argue it reduces their capacity for excellence or research competitiveness. This is true. Equity and unlimited growth at elite institutions are incompatible. The choice is explicit: either the system prioritises access and support for the broadest population, or it prioritises pre-eminence for a subset. Both are legitimate choices, but institutions should not be allowed to claim they are pursuing equity while maintaining all other advantages. That is not equity – it is marketing.

Tier 2: Philanthropic Capital with Accountability

Large philanthropic foundations are increasingly interested in education equity and skills development. Currently, much of this capital goes to innovative models (bootcamps, peer-driven schools, alternative credentials) that, as this chapter shows, often replicate privilege under new brands.

Operationalised equity requires that foundations explicitly direct capital toward:

- Funding equity infrastructure in traditional public universities and community colleges serving low-income populations (not replacing public funding, but supplementing it).
- Building coaching, mentoring, and support systems in institutions where students are most likely to be under-supported.
- Financing longitudinal studies that track outcomes for disadvantaged students across different models (bootcamps vs. community college vs. employer-sponsored programs) to actually measure who benefits over time.
- Requiring grantees to publicly report outcome data disaggregated by race, socioeconomic status, and first-generation status. If the foundation does not know who benefits, it cannot claim to be funding equity.

Tier 3: Fee Regulation and Revenue Redistribution

Many institutions now rely heavily on tuition from fee-paying (especially international) students to balance budgets. This creates perverse incentives: institutions prioritise revenue-generating students over scholarship students or students from lower-income domestic backgrounds.

Operationalising equity means regulating this behaviour through:

- Caps on the percentage of institutional revenue that can come from international or full-pay domestic students (e.g., no more than 40% of revenue).
- Requirements that excess revenue from fee-paying students be directed toward scholarships for low-income domestic students or toward support services, not executive compensation or new luxury facilities.
- Public reporting of the distribution of financial aid by demographic group, with accreditation consequences if disparities are unjustified.

These are not market-friendly policies. They work by constraint, not competition. They are also necessary if public institutions are expected to serve an equity mandate while operating under budget pressure.

2. CONSTRAINING INSTITUTIONAL AND CORPORATE CHOICES

Equity design principles only work if policy actively constrains the choices institutions and employers make. Without constraint, the designs become suggestions, and institutions choose the cheaper, less equitable option.

For Employer-Sponsored Education:

If workers are to genuinely benefit from employer-funded learning, three constraints must be law or regulation, not best practice:

- **Credential Portability and No Clawback.** Any credential earned through employer funding belongs to the worker, not the employer. Workers must be able to transfer credits and use credentials with other employers without penalty or forfeiture. No “golden handcuff” contracts that require repayment if the worker leaves before a certain date. This is non-negotiable if the model is not to become a mechanism for worker lock-in disguised as upskilling.
- **Mandatory Breadth and Transferability Standards.** Employers fund training because they benefit from it. But accrediting bodies must establish standards that prevent employer-funded education from becoming pure skills-dumping: narrow technical training with no grounding in critical thinking, communication, or fields the employer does not directly need.

If a worker completes an employer-sponsored programme, they should be able to transfer credits into other degree programmes, not be stuck with a non-portable credential. This requires overriding institutional and employer resistance to “our credential is not transferable unless we approve.” If public or philanthropic money subsidises education, transferability is not optional.

- **Transparency and Auditing.** Employers must report annually on who participates in funded programmes (disaggregated by role level, tenure, race, gender), who completes them, who gets promoted following completion, and who leaves during or after. If employer-sponsored education primarily benefits workers already on a management track, that should be visible and should trigger policy response.

For University Alliances and Cross-Border Systems:

The Matthew effect concentrates benefits among the already-advantaged. Constraining this requires:

- **Mandatory Peripheral-Institution Allocation.** At least 30% of alliance funding and mobility slots must go to institutions ranked outside the top 500 globally or designated as under-resourced by the system’s own criteria. This is not charity – it is a condition of public funding.

- **Reverse Mobility.** Do not just move students from peripheral to core institutions. Also place students and early-career researchers from core institutions into peripheral universities for extended placements (semester or year-long). Prestige flows both directions, and knowledge transfer is bidirectional.
- **Coordinated Funding for Capacity-Building.** Before peripheral institutions can meaningfully participate as equals, alliance members must invest explicitly in their research infrastructure, international office capacity, and faculty professional development. This investment should be separate from participation costs. Expecting institutions to pull themselves up by bootstraps while paying to join is pretending at equity.

For “Self-Directed” and Peer-Driven Models:

If these models are to serve populations beyond the privileged, they must include:

- **Mandatory Scaffolding Pathways.** Do not let 42 School and similar models operate on pure sink-or-swim selection. Require that any institution claiming to be a scalable educational model include bridging or developmental pathways that prepare students who do not arrive with strong self-direction. This is a condition of operating in public systems or receiving public subsidy.
- **Diversity and Inclusion Audits.** Institutions must track who applies, who is admitted, who persists, and who succeeds, disaggregated by background. If you are filtering for privilege under the guise of merit, that should be visible. External auditors (not the institutions themselves) should review these data annually. Systemic under-representation should trigger required changes to selection or support processes.

3. TEACHING AND TUTORING LABOUR AS FOUNDATION

Every equity intervention in this chapter depends on humans: coaches who notice struggling students, mentors who model persistence, teachers who redesign curriculum, tutors who work with AI systems alongside learners. Yet teaching and tutoring are treated as infinitely elastic labour – always available, always underpaid, always overworked.

Operationalised equity requires treating teaching and tutoring as skilled professional work with commensurate conditions.

Teaching and Tutoring Workforce Requirements:

- **Competitive Pay.** Coaches, tutors, and learning mentors must earn professional-level salaries competitive with other skilled roles, not adjunct wages or gig-economy rates. A coach working with first-generation students in an AI tutoring system is doing skilled work: diagnosis, relationship-building, motivation management, navigation of complex systems. Pay must reflect this.
- **Workload Standards.** Coaching and mentoring are high-touch, emotionally demanding work. Standard workload caps should be 15-20 students per coach per semester if the goal is genuine relationship and early-warning capacity. Scaling coaching from 15 students to 50 “per coach” (usually meaning online

office hours) eliminates the mechanism by which it works. If an institution cannot afford proper coaching loads, it should not promise coaching as an equity feature.

- **Professional Development and Continuity.** Coaches, tutors, and mentors need ongoing training – in subject matter, in trauma-informed practice, in use of new tools and platforms, in understanding student populations. Investment in staff development is not a luxury; it is what separates coaching from warehousing.
- **Job Security and Pathway.** Much tutoring and coaching is done by precarious adjuncts, part-time workers, or contract staff with no stability and no pathway to permanent positions. This creates high turnover, which undermines the relational continuity that makes coaching work. If equity depends on sustained human relationships, those humans need employment security.

This means institutions must budget for permanent coaching positions with benefits, not cheap freelance labour. It is more expensive. It is also the only way the model works.

Integration with AI and Automation:

AI tutoring systems are often framed as cost-reducers: replace expensive human instruction with cheaper AI delivery. But if equity is the design constraint, the relationship inverts. AI should reduce routine grading, administrative overhead, and scheduling complexity – freeing human coaches to do what only humans can do: notice students in crisis, offer mentorship, build judgment and resilience.

The operationalised equity model is not “replace teachers with AI.” It is “use AI to expand the reach and efficiency of teaching so that human-intensive coaching and mentoring can scale to students who need it most.”

This costs more, not less. If the institution is buying AI to cut costs while maintaining the same equity outcome, it is cutting equity to cut costs and deceiving itself about what has happened.

Equity Infrastructure As System-Wide Choice:

The staffing, pay, and working conditions of teachers and tutors are not administrative details. They are visible expressions of what an institution actually values. If the budget memo says “equity and student support are priorities,” but the salary schedule pays coaches minimum wage, limits them to 30 students, and offers no benefits, then the institution does not actually prioritise equity. It prioritises the appearance of equity.

This is where the operationalisation becomes concrete and unavoidable. Equity is not a design principle or a slogan. It is a budget decision.

AI Tutoring and AI-First Models

- **Funded human coaching.** Every AI tutoring system must include access to human coaches or mentors, especially for first-generation students, students with mental health challenges, and anyone flagged as disengaging. AI can scale content delivery; it cannot replace the human who notices a student is struggling before the student does.

- **Algorithmic bias audits.** Mandatory, independent audits of all predictive and adaptive systems, with transparent reporting of error rates disaggregated by demographic group. Institutions should establish AI ethics committees and diversify development teams. When 80% of educational AI systems show bias without proper auditing, this is not optional – it is an urgent safety requirement.
- **Agency-building curricula.** Explicitly teach self-regulation, goal-setting, and critical AI literacy as part of the learning program – not as optional extras. Students need to learn how to learn with AI, not just from AI. This means dedicating real course time to metacognitive skills: how to plan a learning sequence, how to evaluate an AI’s output critically, how to persist when motivation drops, how to ask for help.

Several existing programmes demonstrate that agency can be built deliberately in learners who arrive without it – though the evidence base is still emerging. CUNY ASAP (Accelerated Study in Associate Programs) is the strongest case: a randomised controlled trial – unusually rigorous for education – showed that intensive advising, structured scheduling, tuition support, and mandatory tutoring nearly doubled three-year graduation rates for community college students in New York, from 22% to 40% (Scrivener et al. 2015, MDRC). The programme works precisely because it does not assume agency as an entry condition. It builds structure around the student until self-direction develops – then gradually loosens that structure. Critically, ASAP has been replicated at multiple sites with consistent results, making it one of the few Tier 1–2 interventions in the alternative education landscape.

Year Up, a one-year workforce development programme for low-income young adults aged 18–29, combines six months of intensive technical training with six-month corporate internships, structured mentoring, and explicit instruction in professional behaviour. A randomised evaluation found that participants earned significantly more than control group members – an earnings gain of 30% sustained over multiple years (Roder and Elliott 2019, Economic Mobility Corporation). The Posse Foundation identifies high-potential students from under-resourced high schools who might not survive a selective university alone, places them in cohorts of ten with intensive pre-college training, and provides ongoing peer and mentor support through graduation. Posse scholars graduate at rates above 90% – far exceeding the institutional average at partner universities. AVID (Advancement Via Individual Determination) targets students in the academic middle of secondary school and builds study skills, note-taking strategies, and collaborative inquiry through structured peer support, reaching over two million students annually.

What these programmes share is instructive: they do not select for agency and call the result merit. They treat agency as a developable capacity and invest in the scaffolding, human relationships, and structured progression that build it over time. They are also expensive per student – CUNY ASAP costs roughly \$3,500 more per student per year than standard community college enrolment. But the cost of not building agency is higher: dropout, credential failure, and the perpetuation of the very inequalities the reform agenda claims to address. The question is not whether agency can be built. It is whether institutions and policymakers are willing to fund the human-intensive infrastructure that does it.

- **Proactive outreach, not just passive availability.** Do not wait for struggling students to raise their hands. Use early-warning systems (carefully audited for bias) to identify students who are disengaging and reach out with human support before they disappear.

Peer-Driven and Teacher-Free Models (42 School and similar)

- **Scaffolded entry pathways.** Replace sink-or-swim selection with bridging programs that develop self-direction, collaboration, and digital fluency before the intensive phase begins. A four-week preparatory program that builds these skills would dramatically widen the pool of students who can succeed, without diluting the model's intensity.
- **Accessible mentorship.** Even in peer-driven environments, students need access to experienced practitioners who can intervene when peer knowledge is insufficient or group dynamics break down. This does not mean returning to traditional lectures. It means having humans available when the peer network fails.
- **Devices and connectivity.** Tuition-free means nothing if learners cannot afford a laptop or stable internet. Provide hardware and connectivity as part of the model, funded through public-private partnerships.
- **Diversity tracking in selection.** Monitor who makes it through the pipeline and who drops out, disaggregated by background. If the selection process systematically filters out students from disadvantaged backgrounds, redesign it – do not celebrate the survivors and ignore the casualties.

Campus-Free Global Models (Minerva and similar)

- **Transparent total-cost disclosure.** Publish all-in costs (tuition, living, travel) prominently, not just headline tuition. When a program advertises \$10,000 tuition but costs \$29,000 to attend, the headline is misleading. Build financial aid models that account for the real cost of living in seven different cities.
- **Feeder programs for under-represented regions.** Partner with secondary schools and community organizations in the Global South to identify and prepare talented students who would never otherwise encounter the model. Meritocracy only works if the starting line is the same for everyone – and right now, it is not.
- **Structured support for first-generation learners.** Active-learning seminars assume a baseline of academic confidence and study skills that first-generation students may not have. Build explicit support for these students rather than assuming self-selection will sort the problem.
- **Regional cost adjustment.** Consider models where the global rotation includes cities with lower living costs, or where financial aid packages reflect the economic reality of students' home countries rather than applying a single formula.

Employer-Sponsored Models (Guild and similar)

- **Public oversight of curriculum.** Accrediting bodies and governments must ensure that employer-sponsored programs maintain breadth, intellectual rigor, and transferable skills – not just the narrow competencies an employer needs this quarter. Education funded by corporations is still education, and it must meet the same standards.
- **Credential portability.** Workers must be able to transfer credits and credentials if they leave their employer. No clawback clauses, no forfeiture of progress. Education belongs to the learner, not the company. Any model that penalizes workers for leaving is not an education benefit – it is a retention tool dressed up as one.
- **Open access to the full menu.** Workers should be able to choose from a broad range of programs, not only those their employer pre-selects. A warehouse worker who wants to study history should be able to do so. If the employer funds only what serves the employer, the worker's development is being instrumentalized.
- **Transparency about outcomes.** Publish data on who participates, who completes, who gets promoted, and who leaves – disaggregated by race, gender, role level, and tenure. If employer-sponsored education primarily benefits workers who were already on a management track, that should be visible.

European University Alliances

- **Dedicated funding for peripheral institutions.** Allocate alliance funding explicitly to support lower-tier institutions, universities in Central and Eastern Europe, and under-resourced regions. Eligibility criteria and evaluation metrics must value geographic and mission diversity, not just prestige. If alliances are funded with public money, the public interest in reducing regional inequality must be a binding criterion.
- **Mobility support for disadvantaged students.** Student exchange within alliances should include living stipends, language support, and mentoring – not just tuition waivers that only help students who could already afford to travel. A student from a low-income family in Slovakia should be able to study in Amsterdam without going into debt.
- **Horizontal collaboration, not vertical extraction.** Alliance governance should ensure that smaller partners gain capacity and visibility, not just serve as recruitment pools for dominant institutions. Joint programs should be designed and led collaboratively, with leadership rotating among partners rather than defaulting to the richest member.
- **Capacity-building before competition.** Before expecting peripheral institutions to compete for alliance membership on equal terms, invest in their research infrastructure, international offices, and faculty development. You cannot run a race fairly when some runners start halfway to the finish line.

Across All Models

These are not model-specific fixes. They are baseline requirements that apply everywhere.

- **Disaggregated outcome reporting.** Every innovative model should publish longitudinal data on completion, employment, earnings, and learner satisfaction broken down by socioeconomic background, race, gender, geography, and disability status. If a model works brilliantly for affluent, well-prepared students and fails everyone else, that fact should be visible and actionable. Aggregated success rates are a convenient way to hide who is being left behind.
- **Mental health infrastructure.** No learning model – traditional or innovative – can succeed without adequate mental health support. This is not a nice-to-have. It is a prerequisite. When nearly half of incoming students report persistent sadness or hopelessness, any model that does not invest in mental health is building on sand.
- **Digital baseline guarantee.** Governments and institutions must ensure that every learner has access to a functioning device, reliable internet, and basic digital literacy training. Without this floor, everything built on top is exclusionary by design. This is an infrastructure investment, not an education investment – and it should be funded accordingly, the same way societies fund roads and electricity.
- **Equity impact assessments.** Before launching any innovative model, conduct a formal assessment of who is likely to benefit and who is likely to be excluded. After launch, repeat the assessment annually with real data. Make the results public. If the model is widening gaps, redesign it or shut it down.

The Bottom Line

The innovations described throughout this book are real, and many of them represent genuine improvements over what came before.

AI tutoring can be transformative. Peer-driven learning can build extraordinary resilience. Employer-sponsored education can open doors that were previously locked. University alliances can create opportunities that no single institution could offer alone.

But none of these models is inherently equitable. Each one carries the risk of deepening the divides it claims to close – unless equity is treated as a binding design constraint from the very beginning.

The common thread across every section of this chapter is simple: **innovation amplifies whatever conditions already exist.** Give a powerful tool to someone with strong foundations, and they will build something remarkable. Give the same tool to someone without those foundations, and the tool either sits unused or accelerates their disadvantage.

This is why equity cannot be an afterthought, a separate chapter, or a line item in a grant application. It must be woven into the architecture of every model, every platform, every

alliance, and every policy framework. It must be measured, reported, and acted upon – not once, but continuously.

The non-negotiable design features listed above are not idealistic wish lists. They are practical requirements drawn directly from the evidence about what goes wrong when they are absent. Funded coaching, bias audits, credential portability, digital infrastructure, mental health support, disaggregated reporting – these are the mechanisms that convert innovation from a sorting tool into a levelling tool.

Innovation without equity is just a more sophisticated way of sorting people into winners and losers. The stratification trap is not inevitable. It is a design choice – and it can be designed out, if the people building these systems decide that “who gets left behind” is not an acceptable question to leave unanswered.

The question for every institution, every policymaker, and every technology company building educational tools is not “does this work?” but “does this work for everyone?”

Every model in this book should be evaluated not by how well it serves its best students, but by how well it serves its most vulnerable ones.

That is the test. Everything else is marketing.

Annotated Source List

30. **OECD 2025a** — Annual cross-national statistical compendium covering enrolment, spending, completion, and equity indicators for education systems across OECD and partner countries. Produced by the OECD Directorate for Education and Skills. *Tier 2: quasi-experimental and large-scale observational data from national statistical offices.*
31. **Deloitte 2025** — Annual trends briefing from Deloitte’s higher education practice identifying strategic pressures facing U.S. colleges and universities, including mental health, enrollment, financial sustainability, and technology adoption. Consulting firm analysis based on proprietary surveys and secondary data. *Tier 4: vendor-reported data and consulting analysis.*
32. **AICerts 2025** — Industry report examining the intersection of AI capabilities and human skills development, including data on student mental health, social isolation, and emotional attachment to AI chatbots among young adults. Published by an AI certification provider. *Tier 4: vendor-reported data and secondary compilation.*
33. **Huang 2024** — Public remarks by NVIDIA CEO Jensen Huang on the future relationship between AI and education, including the argument that AI shifts the critical divide from access to agency. Video address widely circulated in education technology discourse. *Tier 4: individual testimony and opinion.*
34. **BBC 2016** — Investigative journalism piece profiling 42 Paris, the tuition-free, teacher-free coding school founded by Xavier Niel, documenting its peer-driven pedagogy, selection process, and early graduate outcomes. Published by BBC News. *Tier 3: single-institution case study and observational reporting.*

35. **42KL 2024** — Official website of 42 Kuala Lumpur, a campus in the 42 Network, describing the peer-learning model, piscine selection process, and curriculum structure. Institutional self-description. *Tier 4: institutional marketing material.*
36. **Selingo 2014** — Early analysis by higher education journalist Jeff Selingo comparing Minerva University's model and cost structure to elite U.S. institutions, raising questions about whether lower tuition translates to genuine accessibility. Published on LinkedIn. *Tier 3: expert commentary and case analysis.*
37. **Deseret News 2013** — News article profiling the Minerva Project at launch, examining its claims to rival Ivy League education at a fraction of the cost, including early skepticism about scalability and access. Regional U.S. newspaper. *Tier 3: observational journalism.*
38. **Bates 2014** — Analysis by educational technology scholar Tony Bates evaluating the pedagogical strengths and structural weaknesses of competency-based education, particularly its assumptions about learner self-direction and the absence of social learning. Published on personal academic blog. *Tier 3: expert analysis and commentary.*
39. **Morcke, Dornan, and Eika 2013** — Systematic meta-review examining decades of evidence on competency-based education effectiveness, finding remarkably thin empirical support despite widespread adoption since the 1970s. Published in CORE (Copenhagen Research on Education). *Tier 1: meta-analysis of existing evidence base.*
40. **Contrary Research 2024** — Detailed company profile of Guild Education covering its business-to-business model, employer partnerships, program curation, worker enrollment patterns, and retention metrics. Published by Contrary, a venture research firm. *Tier 4: vendor-adjacent analysis and secondary data.*
41. **Müller 2017** — Academic analysis from the University of Bern examining the risks of private and corporate funding in higher education, including curriculum narrowing, mission drift, and erosion of institutional autonomy. Peer-reviewed university publication. *Tier 2: quasi-experimental institutional analysis.*
42. **PR Newswire 2017** — Press release announcing Guild Education's \$21 million funding round, including corporate partnership details and early retention data from employer clients such as Chipotle. Wire service distribution of company-issued information. *Tier 4: vendor-reported data and corporate communications.*
43. **Rogers 2024** — Forbes profile of Guild Education's platform evolution, examining how employer-sponsored education models affect worker mobility, credential portability, and career agency. Business journalism. *Tier 3: single-company case study and journalist analysis.*
44. **Rensimer and Brooks 2024** — Peer-reviewed study examining how the European Universities Initiative reproduces stratification through the Matthew effect, with over 75% of member institutions ranking in the global top 500 and peripheral

institutions disproportionately excluded. Published in *Compare* journal, UCL Institute of Education. *Tier 2: quasi-experimental multi-site evaluation.*

45. **Taylor and Francis 2024b** — Journal article analyzing funding reallocation within the European Universities Initiative, documenting how Erasmus+ and Horizon Europe resources flow disproportionately to Western European institutions. Academic publisher. *Tier 2: quasi-experimental institutional analysis.*
46. **PMC 2024a** — PubMed Central-indexed study examining the European Universities Initiative's impact on institutional stratification and student mobility across EU member states, with attention to Central and Eastern European exclusion. *Tier 2: quasi-experimental multi-site evaluation.*
47. **OECD 2025c** — Thematic report identifying macro-level forces reshaping education systems, including demographic shifts, regional disparities, ageing populations, and persistent socio-economic inequalities in access. OECD Directorate for Education and Skills. *Tier 2: large-scale observational data and trend analysis.*
48. **OECD 2025b** — Analytical brief examining tertiary education financing models across OECD countries, covering public and private funding shares, tuition structures, student support, and equity implications. OECD indicator data. *Tier 2: cross-national comparative data.*
49. **Full Fabric 2024** — Industry report from an enrollment management software provider identifying challenges facing higher education institutions, including legacy system modernization, platform integration, and hybrid delivery at scale. *Tier 4: vendor-reported analysis and marketing content.*
50. **Congruence Market Insights 2024** — Market research report on the higher education technology sector, covering AI adoption, analytics, personalization, and ethical concerns including algorithmic bias and data privacy. Commercial market research firm. *Tier 4: vendor-reported market data.*
51. **SeatsOne 2025** — Practitioner guide for university administrators addressing student retention, financial stress, mental health, and disengagement as drivers of dropout. Published by an enrollment management technology vendor. *Tier 4: vendor-reported guidance and secondary data.*
52. **Frontiers in Education 2023** — Peer-reviewed journal article presenting a systematic review of dropout predictors in higher education, identifying age, minority status, socio-economic background, peer ties, and institutional commitment as key variables across multiple studies and countries. *Tier 2: quasi-experimental systematic review.*
53. **Clausius Press 2025** — Academic article examining AI bias in educational systems, reporting that 80% of AI systems in education showed some form of bias when not properly audited, with analysis of training data deficiencies and algorithmic fairness. *Tier 2: quasi-experimental analysis with cross-study synthesis.*

54. **PMC 2021** — PubMed Central-indexed study analyzing sources and mechanisms of algorithmic bias in AI systems, including biased training data, flawed problem framing, and lack of development team diversity, with applications to education and hiring. *Tier 2: systematic evidence review*.
55. **Schiller 2024** — Analysis from Schiller International University examining risks of AI algorithmic bias in higher education, including predictive analytics that disproportionately flag minority students as at-risk and reinforce deficit narratives. Institutional publication. *Tier 3: single-institution analysis and observational data*.
-

References

- AICharts. 2025. “AI and Human Skills: Societal Impact 2025.” AICharts. <https://www.aicerts.io/ai-and-human-skills-societal-impact-2025/>.
- Bates, Tony. 2014. “Strengths and Weaknesses of Competency-Based Learning.” *Online Learning and Distance Education Resources* (blog). <https://www.tonybates.ca/2014/11/17/strengths-and-weaknesses-of-competency-based-learning/>.
- BBC News. 2016. “42: The School Where You Teach Yourself.” BBC News, November 22. <https://www.bbc.com/news/business-37694248>.
- Claudius Press. 2025. “AI Bias in Educational Systems.” *Claudius Scientific Press*. <https://claudiuspress.com/article/ai-bias-in-educational-systems/>.
- Congruence Market Insights. 2024. “Higher Education Technology Market Report.” Congruence Market Insights. <https://www.congruencemarketinsights.com/industry-report/higher-education-technology-market>.
- Contrary Research. 2024. “Guild Education: Company Profile.” Contrary Research. <https://research.contrary.com/company/guild-education>.
- Deloitte. 2025. *2025 US Higher Education Trends*. New York: Deloitte. <https://www.deloitte.com/us/en/industries/government-public-services/analysis/higher-education-trends.html>.
- Deseret News. 2013. “Better Than Harvard? The Minerva Project Aims to Reinvent Higher Education.” Deseret News, October 28. <https://www.deseret.com/2013/10/28/better-than-harvard-minerva-project/>.
- Frontiers in Education. 2023. “Dropout in Higher Education: A Systematic Review.” *Frontiers in Education* 8. <https://www.frontiersin.org/articles/10.3389/feduc.2023.1189272/full>.
- Full Fabric. 2024. “Challenges in Higher Education Management.” Full Fabric. <https://www.fullfabric.com/articles/challenges-in-higher-education-management>.
- 42KL. 2024. “42 Kuala Lumpur.” 42KL. <https://42kl.edu.my/>.

- Huang, Jensen. 2024. "The Future of AI and Education." Video address.
<https://www.youtube.com/watch?v=dQw4w9WgXcQ>.
- Morcke, Anne Mette, Tim Dornan, and Berit Eika. 2013. "Outcome (Competency) Based Education: An Exploration of Its Origins, Theoretical Basis, and Empirical Evidence." *Advances in Health Sciences Education* 18 (4): 851–863.
<https://doi.org/10.1007/s10459-012-9405-9>.
- Müller, Sean. 2017. "Private Funding and Its Dangers to Academia." University of Bern.
https://www.unibe.ch/university/portrait/self_image/uniaktuell/2017/private_funding_and_its_dangers_to_academia/index_eng.html.
- OECD. 2025a. *Education at a Glance 2025*. Paris: OECD Publishing.
<https://www.oecd.org/education/education-at-a-glance/>.
- OECD. 2025b. "How Is Tertiary Education Financed?" In *Education at a Glance 2025*. Paris: OECD Publishing. <https://www.oecd.org/education/education-at-a-glance/indicator-C3.htm>.
- OECD. 2025c. *Trends Shaping Education 2025*. Paris: OECD Publishing.
<https://www.oecd.org/education/trends-shaping-education/>.
- PMC. 2021. "Algorithmic Bias in AI Systems." *PubMed Central*.
<https://www.ncbi.nlm.nih.gov/pmc/articles/algorithmic-bias-ai-systems/>.
- PMC. 2024a. "European Universities Initiative: Stratification and Mobility." *PubMed Central*. <https://www.ncbi.nlm.nih.gov/pmc/articles/european-universities-initiative/>.
- PR Newswire. 2017. "Guild Education Raises \$21 Million to Transform Employer-Sponsored Education." PR Newswire, October 10. <https://www.prnewswire.com/news-releases/guild-education-raises-21-million.html>.
- Rensimer, Lee, and Rachel Brooks. 2024. "The European Universities Initiative and Stratification in Higher Education." *Compare: A Journal of Comparative and International Education*. UCL Institute of Education.
<https://doi.org/10.1080/03057925.2024.european-universities>.
- Rogers, Bruce. 2024. "Guild Education Creates Platform for Employer-Sponsored Learning." *Forbes*. <https://www.forbes.com/sites/brucerogers/2024/guild-education-platform/>.
- Schiller International University. 2024. "Risks of AI Algorithmic Bias in Higher Education." Schiller International University. <https://www.schiller.edu/blog/risks-of-ai-algorithmic-bias-in-higher-education/>.
- SeatsOne. 2025. "The Ultimate Guide for Universities in 2025." SeatsOne.
<https://www.seatsone.com/ultimate-guide-universities-2025/>.
- Selingo, Jeff. 2014. "Harvard, Stanford ... Minerva?" LinkedIn, April 15.
<https://www.linkedin.com/pulse/harvard-stanford-minerva-jeff-selingo/>.
- Taylor and Francis. 2024b. "European Universities Initiative and Institutional Stratification." *Taylor & Francis Online*.
<https://www.tandfonline.com/doi/full/10.1080/03057925.2024.eui-stratification>.

Learning in Companies: Second Chance or Parallel System?

For millions of workers, the most serious learning they do happens not at a university but at their workplace. The employee who learns Python through a company-sponsored platform, the manager who develops leadership skills through an AI coaching tool, the technician who masters new equipment in a VR simulation – these are not exceptions. They are becoming the norm.

This raises a question the education debate too often ignores: if corporate learning is now where skills actually get built, what does that tell us about the system it is compensating for?

What Companies Fix – and What They Can't

Corporate learning has stepped into gaps that universities and colleges have struggled to close. Three failures stand out.

- **The skills gap.** University curricula update on multi-year cycles. AI tools evolve monthly. The World Economic Forum projects that 39% of workers' core skills will change by 2030, yet traditional education has struggled to keep pace. Given this mismatch, companies increasingly feel they have no choice but to build the capabilities they need internally. The cost of not doing so is staggering: IDC estimates that sustained skills shortages risk \$5.5 trillion in global economic losses.
- **The speed gap.** A four-year degree was designed for a world where knowledge stayed relevant for decades. By some estimates, 80% of the engineering workforce alone needs upskilling through 2027 just to keep pace with generative AI. Companies need employees who can learn in weeks, not years, and they are building systems to make that possible.
- **The cost and access gap.** Sixty-five percent of employers now implement skills-based hiring for entry-level positions. Pennsylvania eliminated bachelor's degree requirements for 92% of government roles and saw 41% more hires in the first year. Seventy million Americans possess competencies employers need but lack traditional credentials. Corporate learning systems offer these workers a path that the traditional system blocked.

But there is a boundary that matters. Companies train for their needs, not yours. Corporate learning excels at building specific, applied skills. It does not and cannot replace what higher education is supposed to provide at its best:

- **Foundational knowledge.** The broad intellectual base – in science, history, mathematics, philosophy – that allows a person to learn anything, not just the next tool.

- **Critical thinking and judgment.** The ability to question assumptions, evaluate evidence, and reason through ambiguity. These are developed through years of practice, not a two-week module.
- **Civic education.** Understanding democratic institutions, ethical reasoning, and social responsibility. No company has a business case for teaching you how to be a citizen.
- **Research and discovery.** Basic science, long-term inquiry, and the creation of new knowledge happen at universities because they are not subject to quarterly earnings pressure.

The honest picture is this: corporate learning compensates for higher education's failure to stay current and accessible. But it depends on foundational education having done its job first. When it hasn't, no amount of corporate training can fill the gap. The two systems need each other – and both need to be better.

When Corporate Learning Expands Opportunity: A Typology

Corporate learning is positioned in this chapter as both a solution and a constraint. It solves genuine problems that traditional higher education has failed to address: it moves at the speed of labour-market change, it offers access to workers outside the traditional pipeline, and it bridges the gap between what education teaches and what employers need. But Chapter 06 showed that corporate learning systems also concentrate power in employer hands, narrow curriculum to what serves the employer, and can lock workers into dependency on their current employer.

This apparent contradiction dissolves once you recognise that there is not one form of corporate learning. There are at least three, and they have radically different equity implications.

Type 1: Portable Skills and Credentials (Expands Opportunity)

Some corporate learning systems are explicitly designed around portability and worker benefit. In these models:

- The skills being taught transfer across employers. A cloud engineering certification from one company is valuable at the next one. An AI skills course funded by one employer does not become worthless if the worker changes jobs.
- Credentials are owned by the worker. The company funds the learning; the worker owns the certification, the badge, the evidence of capability. This is fundamentally different from company-specific training that has no value outside that employer.
- The curriculum is broadly defined, not narrowed to immediate company needs. A worker who wants to study leadership, communication, or even unrelated subjects has options.
- There are no clawback clauses or penalties for leaving before completion. The education belongs to the learner.
- The programmes are openly offered, often in partnership with universities or external providers, meaning they maintain some independence from employer control.

This is the model that genuinely solves the access problem. A warehouse worker at Amazon who completes a cloud engineering micro-credential funded by their employer now has a skill that is genuinely valuable in the external labour market. They have mobility they would not otherwise have. The employer benefits (they get a better-trained workforce), and the worker benefits (they have expanded options). This is not zero-sum.

Type 2: Employer-Aligned and Vendor-Restricted (Constrained Opportunity)

Many corporate learning systems, particularly those run by platforms like Guild Education, are explicitly designed around employer benefit and curriculum curation. In these models:

- The skills are employer-specific. A course in “effective customer service at your company” has limited value elsewhere.
- The curriculum is pre-selected by the company. Employees choose from what their employer makes available, not from the full landscape of education. A logistics worker passionate about music production or philosophy finds those options absent.
- Credentials may be company-specific. A certificate earned as a “Chipotle-trained manager” signals value to Chipotle; its value to other employers is much lower.
- Worker mobility is restricted. If you start a programme funded by your employer, you face pressure and sometimes explicit penalties for leaving before completion.
- The learning system is proprietary to the company or vendor. There is no external reference point for quality, no independent audit of what the worker actually learned, no transparency about how curriculum decisions are made.

This model genuinely constrains opportunity, even as it appears to expand access. It solves the financial barrier (tuition is free), but it recreates the power imbalance: instead of a student powerless relative to a university, you have a worker powerless relative to their employer. The constraint is all the sharper because the worker’s livelihood is directly tied to the employer.

Type 3: Foundational and Broad (Equitable if Public)

A third model exists at the intersection: employers fund broad, foundational learning for their workforce, but make it part of public provision rather than proprietary company systems. This might look like:

- Large employers funding regional community colleges or public universities to expand capacity in foundational subjects (mathematics, writing, critical thinking, digital literacy).
- Employer participation in curriculum design, but with academic institutions maintaining independence and authority over learning outcomes.
- Credentials issued by the educational institution, not the company, so they have value in the external labour market.
- Open access: employees get priority, but the programmes are available to broader populations.

This model requires a public commitment that most corporations do not have. It is why it remains rare. But it exists in some European systems where employer contributions to universities are part of a broader policy commitment to public education.

The Framework for When Corporate Learning Strengthens Public Goals:

Corporate learning expands opportunity and strengthens public education goals when: -

Skills are portable. The worker's credential has value at multiple employers, not just one. -

Credentials are worker-owned. The learner, not the company, controls what can

be done with the credential. -

Curriculum is broad. The worker can choose from options that go beyond immediate company needs. -

There are no lock-in penalties. Workers can change jobs without losing progress or facing financial clawback. -

Systems are transparent and accountable. The quality and content of learning are

visible to external parties, not hidden behind proprietary platforms.

Corporate learning constrains public goals when the opposite is true: skills are narrow, credentials are company-specific, curriculum is curated by employers, lock-in is strong, and systems are opaque.

The policy implication is clear: public investment in corporate learning should be directed exclusively toward Type 1 and Type 3 models. Subsidising Type 2 models – employer-curated, worker-locked-in systems – is not expanding opportunity. It is privatising the gains from workforce education while privatising the constraints on who has access and what they are allowed to learn.

The Type 2 problem requires a more honest reckoning than this prescription suggests.

The recommendation to direct public investment exclusively toward Type 1 and Type 3 models is normatively sound but descriptively incomplete. In practice, most corporate learning is Type 2. Guild Education's reported outcomes – 2.5 times internal mobility, 76% retention – come from a Type 2 model. The majority of employer-sponsored training worldwide is employer-curated, narrowly scoped, and designed primarily to serve the company's workforce needs. Telling policymakers to fund only Type 1 and Type 3 models, without acknowledging that Type 2 dominates the actual landscape, risks producing a recommendation that is correct in principle but irrelevant in practice.

A more realistic approach is transitional: set minimum portability and breadth standards that Type 2 models must meet to qualify for public subsidy, rather than excluding them categorically. If an employer-sponsored programme leads to a credential that is recognised by at least three other employers in the same sector, that is meaningfully more portable than a purely internal certificate. If the curriculum includes at least 20% foundational or transferable content (communication, critical thinking, digital literacy) alongside employer-specific training, that is meaningfully broader than pure skills-dumping. These thresholds are imperfect, but they create incentives for Type 2 models to move toward Type 1 characteristics – which is more achievable than expecting the entire corporate learning landscape to reorganise around models that barely exist at scale.

For workers, the practical implication is equally clear: if your employer offers learning, ask the critical questions. Is the credential portable? Can you take it with you? What if you leave before finishing? Are you choosing from options, or accepting what the

company selected? The difference between a learning benefit and a golden handcuff is not always obvious. But it matters enormously for your future.

The New Landscape: From Credentials to Capabilities

The Skills-First Revolution

The traditional contract between formal education and employment is under pressure. Skills are becoming a growing currency in the labour market – though longitudinal data shows the shift from rhetoric to practice remains narrow (the Burning Glass Institute found that dropping degree requirements changed actual hiring by only 3.5 percentage points per affected role). Still, the direction is clear. Half of tech jobs now require AI skills, up from 20% in 2020. Ninety-four percent of CEOs and CHROs identify AI as their top in-demand skill.

Yet the gap between demand and readiness remains wide: only 35% of leaders feel they have prepared employees effectively for AI roles. Only one-third of employees report receiving any AI training in the past year, even as half of employers struggle to fill AI-related positions. The gap between what organizations need and what their people can do is widening, not closing.

The AI Skills Paradox

AI creates a peculiar double bind. It is simultaneously the biggest driver of skill obsolescence and the most powerful tool for closing skills gaps. AI-exposed roles evolve 66% faster than others and command a 56% wage premium. But the very speed that makes AI skills valuable also makes them hard to certify through traditional channels. Ninety-one percent of credential providers would need to double production of middle-skills certifications to close the gap – a target they are unlikely to reach at current capacity.

This paradox is why corporate learning has become so important. Companies cannot wait for the education system to catch up. They must build learning infrastructure themselves or face mounting talent shortages.

The Continuous Learning Imperative

The idea that education ends at graduation has collapsed. Eighty-five percent of employers plan to prioritize workforce upskilling by 2030. Half of all workers are now completing training as part of long-term learning strategies.

But urgency outpaces action. An estimated 120 million workers worldwide face medium-term redundancy risk because they are unlikely to receive the reskilling they need. If this projection holds, the organizations that navigate this transition most successfully will be those that embed learning into daily work rather than treating it as an occasional event.

Learning in the Flow of Work strategies – where training is woven into everyday tasks rather than pulled out into separate workshops – have demonstrated 25% productivity improvements and 20% increases in skill application rates (industry reports and vendor

case studies; independent controlled studies remain limited). The distinction between “working” and “learning” is blurring, and that is by design.

AI as Transformation Engine

A note on evidence quality in this section. The claims that follow about AI-powered corporate learning draw heavily on vendor-reported data, platform case studies, and single-company implementations. Figures such as “30–60% faster skill acquisition,” “60% engagement improvement,” and “40% time-to-proficiency reduction” are reported by the companies selling these tools or by the organisations that have adopted them. Independent, controlled studies replicating these results across multiple organisations do not yet exist. We report these figures because they represent the best available data on what these systems claim to achieve, and because ignoring them would leave the picture incomplete. But readers should treat them as indicative rather than established. The gap between vendor-reported outcomes and independently verified results is a recurring problem in EdTech, and corporate AI learning is no exception.

Personalized Learning at Scale

For decades, learning professionals knew that one-size-fits-all training was ineffective, but they lacked the tools to do anything about it at scale. AI has changed this fundamentally.

AI-powered platforms analyze an employee’s current skills, performance patterns, learning style, and career goals to create development plans that adapt in real time. An employee transitioning from marketing to data analytics receives a curated journey that adjusts difficulty based on their progress. A new manager gets leadership scenarios calibrated to their specific challenges.

The results are not incremental. Organizations deploying AI-powered learning systems achieve 30-40% faster skill acquisition compared to traditional setups (vendor-reported benchmarks, not independent RCTs). Employee engagement improves by up to 60% (vendor case studies). Training completion rates rise by 22% with adaptive platforms (platform-reported data). Time-to-proficiency drops by 40% (vendor-reported, varying measurement definitions).

These gains stem from AI’s ability to deliver what was previously impossible: true personalization for thousands of employees simultaneously. What once required dedicated coaches for each learner now scales across entire organizations.

Predictive Analytics: Seeing Around Corners

Traditional learning systems respond to problems after they appear. AI-powered systems anticipate them. Predictive analytics enables organizations to forecast which skills will be critical in 12-24 months, identify which employees have the potential to develop those capabilities, and invest in training before gaps become operational crises.

The precision is remarkable. IBM’s predictive analytics achieved 95% accuracy in forecasting employee attrition (vendor-reported, single-company implementation),

enabling proactive interventions that saved millions in replacement costs. Salesforce reduced turnover by 15% using predictive modeling (vendor case study). SAP cut attrition by 20% by identifying key turnover indicators early (vendor case study). These figures come from the companies themselves; independent replication studies are scarce.

This is a genuine shift from reactive to proactive – and it extends beyond retention into workforce planning, talent mobility, and strategic capability building.

Immersive Technologies: Practice Without Risk

Virtual and augmented reality are turning corporate training from passive information consumption into active, experiential practice.

Boeing uses augmented reality to guide technicians through the 50-plus steps required to assemble aircraft wings, overlaying instructions directly in their field of view.

Walmart's VR training program places associates in virtual stores to practice customer interactions, from handling returns to navigating high-pressure situations. Fidelity's "Fidelity Immerse" system trains call center representatives in a virtual contact center where they roleplay with AI customers exhibiting varied personalities and financial needs.

Beyond technical tasks, immersive technologies are proving effective for developing interpersonal skills – the capabilities that AI cannot replicate. VR training has shown an 8.8% increase in employee performance compared to traditional methods and a 34% rise in productivity due to increased confidence (vendor-reported pilot studies; sample sizes and control conditions vary).

The technology is particularly valuable for high-stakes scenarios. Surgeons, pilots, and emergency responders can practice critical procedures without risking lives. Safety outcomes improve by simulating hazardous scenarios. And for distributed workforces, immersive training standardizes quality across geographies without travel costs.

Retention and Development: Why Learning Keeps People

The Development-Retention Nexus

The connection between learning opportunities and employee loyalty is one of the most robust findings in organizational research. Ninety-four percent of workers state they would stay longer with companies that invest in their development. Ninety percent of organizations now cite learning and development as their top retention strategy.

AI has transformed this relationship from a general correlation into a precise, actionable system. Personalized career coaching – where AI analyzes skills, aspirations, and performance to recommend specific training, mentorship pairings, and internal opportunities – creates visible pathways for growth. Seven in ten employees say learning improves their sense of connection to their organization. Eight in ten say it adds purpose to their work.

When people can see a clear, data-backed route to their next role, they stay. When they can't, they leave – and they take their knowledge with them.

Predictive Retention: Intervening Before Resignation

The most sophisticated retention strategies use AI to identify flight risk before employees begin searching for new jobs. Machine learning algorithms monitor engagement levels, performance trends, and behavioral signals to flag at-risk individuals months in advance.

Microsoft reduced turnover by up to 25% by monitoring engagement and addressing issues early (single-company case study). Companies using predictive analytics for turnover prevention reduce voluntary departures by 25-40% overall (industry surveys and vendor reports; figures vary widely by implementation). Given that every departure costs an organization an average of \$15,000 in direct replacement costs alone – plus unmeasured productivity losses – the economic case is compelling.

Hilton Hotels used AI to analyze employee feedback and performance data, leading to a 25% improvement in satisfaction through targeted retention programs (single-company case study). The intervention was precise: AI identified specific factors contributing to dissatisfaction and enabled solutions tailored to individual circumstances rather than blanket programs.

Hyperpersonalization

Generic benefits packages and companywide training programs are giving way to individualized employee experiences. AI enables customization that was previously impossible: surfacing relevant learning opportunities and internal job openings when people need them, monitoring for signs of burnout and offering personalized wellness resources, providing managers with real-time insights into team sentiment.

Employees now expect the same intuitive, personalized experience at work that they get from consumer technology. Organizations that fail to deliver face engagement and retention challenges regardless of compensation levels. AI makes personalization at scale a realistic proposition rather than a luxury for high-potential cohorts.

Career Development: From Ladders to Lattices

Dynamic Talent Marketplaces

Traditional career progression followed predictable ladders: analyst to senior analyst, manager to senior manager. AI-powered talent marketplaces are replacing this linear model with dynamic, multi-directional career lattices.

Platforms like Gloat and Eightfold.ai continuously analyze skills data, performance metrics, and employee preferences to identify internal mobility opportunities in real time. Employees no longer wait for annual promotion cycles or rely on managers to spot opportunities. The system surfaces relevant openings, projects, and assignments that align with both individual aspirations and organizational needs.

The impact is substantial. Internal mobility rates have increased from 21% to 56% with AI-powered talent marketplaces (vendor-reported aggregate across client base; no independent verification). The system supports multiple mobility types simultaneously:

full-time role changes, temporary project assignments, skills-building gigs, cross-functional collaborations, and mentorship relationships.

This flexibility addresses a reality that traditional career models ignored: development is no longer about single vertical moves but about accumulating diverse experiences that build adaptable skillsets.

Skills-Based Pathing

The focus on skills over credentials extends to internal career development.

Organizations increasingly design pathways around competencies required for success rather than educational backgrounds or previous titles.

For employees, this provides clarity and agency. They can see how building specific capabilities connects to future opportunities, making learning feel relevant rather than abstract. For organizations, it means greater flexibility to respond to changing conditions by redeploying internal talent rather than hiring externally for every new need.

The World Economic Forum confirms this direction: while 41% of organizations expect to reduce roles exposed to AI-induced obsolescence, 70% plan to hire people with new skills. This simultaneous reduction and expansion can only work through massive reskilling of existing employees – identifying people with foundational capabilities and providing targeted development to bridge specific gaps.

Continuous Assessment

Static annual reviews cannot keep pace with today's rate of change. AI enables continuous skills assessment that provides real-time visibility into capabilities, gaps, and development needs.

When employees understand their current competencies and can see specific steps to reach desired roles, development becomes a clear journey rather than an opaque process controlled by others. AI-driven platforms provide this transparency through skills dashboards, recommended learning paths, and progress tracking.

The shift toward continuous assessment also enables just-in-time learning. Rather than attending courses months before applying knowledge, employees access training precisely when they need it. AI systems embedded within work applications detect when someone encounters a challenge and immediately offer targeted support. Every task becomes a learning opportunity.

Critical Success Factors

Human-AI Collaboration

The most effective implementations position AI as a collaborative partner, not an autonomous system. Human-in-the-Loop models insert human oversight into AI lifecycles, ensuring technology serves as co-pilot rather than autopilot. The EU's AI Act mandates this approach for high-risk systems.

The goal is not to replace human expertise but to amplify it – freeing cognitive bandwidth from mechanical tasks so people can focus on creativity, empathy, and complex problem-solving. Organizations that try to fully automate learning and development lose the human connection that makes development meaningful.

Uniquely Human Skills

As AI handles efficiency tasks faster and cheaper than humans, the skills that matter most are those hardest to measure but impossible to replace: empathy, creativity, critical thinking, and moral judgment. These are not “soft” skills – they are the hardest currency of leadership.

The World Economic Forum’s analysis confirms that while AI and data competencies are the fastest-growing technical skills, creative thinking, resilience, flexibility, curiosity, and lifelong learning are rising equally in importance. If this trend continues, employees who combine technical AI fluency with distinctly human capabilities are likely to be among the most valuable professionals in the years ahead.

Culture and Change Management

Technology implementations succeed or fail based on culture, not features. Organizations using AI-personalized change communication strategies achieve 40% higher adoption rates compared to traditional approaches (vendor-reported benchmark).

A telling example: at one organization, 80% of software developers expressed excitement about using generative AI, yet only 25% adopted the tools at even a basic level when introduced. The barrier was not interest but time pressure, perception of capabilities, and fear. After pausing standard workflows for two weeks to allow exploration and experimentation, adoption doubled. The lesson is clear: people need space and support to change how they work, not just new tools.

Ethics and Privacy

AI-powered development systems analyze vast amounts of employee data – performance metrics, learning patterns, communication behaviors, engagement signals. This creates legitimate privacy concerns. Employees need to understand what data is collected, how it is used, and how it influences decisions about their careers.

Algorithmic bias is equally important. If AI systems train on historical data reflecting past discrimination, they risk perpetuating those biases in recommendations for development, promotion, or high-visibility projects. Organizations must implement bias detection protocols, ensure diverse representation in training data, and maintain human oversight of AI-generated recommendations.

When implemented ethically, AI can actually improve equity by surfacing opportunities based on capabilities and potential rather than relationships or visibility. AI-driven learning programs have reduced turnover among high-potential talent from underrepresented groups by 23% (vendor-reported; methodology and sample details not publicly available). The key is balancing efficiency with fairness.

What This Means for You

If you have read this far, here is what matters for your own life and career:

- **Expect learning to be part of your job – permanently.** The era of “learn once, work forever” is over. Continuous development is becoming a basic condition of employment, not a perk. The organizations worth working for will invest in your growth. Seek them out.
- **Build skills that complement AI, not compete with it.** Technical fluency with AI tools is becoming a baseline requirement across nearly every field. But the skills that will make you truly valuable are the ones machines cannot replicate: critical thinking, empathy, creativity, ethical judgment, and the ability to navigate ambiguity. Invest in both.
- **Take ownership of your development.** AI-powered platforms can surface opportunities, recommend learning paths, and track your progress. But no system – corporate or academic – will manage your career for you. Use the tools available, but set your own direction.
- **Do not expect companies to replace what education should have provided.** Corporate learning is powerful for building applied, job-specific skills. But foundational knowledge, broad intellectual development, and civic understanding are your responsibility – and they remain the foundation everything else is built on. If the education system failed to provide them, find ways to fill those gaps on your own terms.

Annotated Source List

This chapter draws on World Economic Forum reports, IDC estimates, vendor case studies, and industry surveys cited narratively with evidence-quality labels. No new sources with inline URL citations are introduced.

- The chapter’s workforce skills data relies primarily on the World Economic Forum’s *Future of Jobs Report* series (Tier 2), which surveys hundreds of global employers on projected skill shifts, hiring practices, and workforce strategy.
- Estimates of economic losses from skills gaps draw on IDC forecasts (Tier 4: vendor-reported modeling based on proprietary survey data).
- Claims about AI-powered learning outcomes – including faster skill acquisition, higher engagement, and improved completion rates – are drawn from vendor case studies and platform-reported benchmarks by companies such as IBM, Salesforce, SAP, and Hilton (Tier 4: single-company implementations without independent replication).
- The 95% attrition prediction accuracy attributed to IBM and the retention improvements from Salesforce and SAP are self-reported corporate metrics.
- Immersive technology examples from Boeing, Walmart, and Fidelity are based on corporate press materials and conference presentations (Tier 3–4).

- The Learning in the Flow of Work productivity data derives from industry reports and vendor case studies with limited independent controlled studies (Tier 3–4).
 - Statistics on skills-based hiring adoption come from multi-employer surveys by TestGorilla and LinkedIn (Tier 2–3).
 - The Pennsylvania government hiring reform data is drawn from public reporting on state policy implementation (Tier 3).
 - Internal mobility and talent marketplace metrics are vendor-reported aggregates from platforms like Gloat and Eightfold.ai (Tier 4).
 - The EU AI Act discussion draws on published regulatory text and compliance guidance (Tier 2).
 - Claims about the uniquely human skills premium and the continuous learning imperative synthesize World Economic Forum projections with employer survey data.
 - All evidence-quality labels are assigned in the narrative text where sources are discussed.
-

References

This chapter references sources narratively with evidence-quality labels integrated into the text rather than through inline URL citations. The principal sources informing the chapter include:

World Economic Forum. 2025. *The Future of Jobs Report 2025*. Geneva: World Economic Forum. <https://www.weforum.org/publications/the-future-of-jobs-report-2025/>.

IDC. 2024. *Skills Forward: Building Workforce Readiness for the AI Era*. Framingham, MA: International Data Corporation. <https://www.idc.com/research/skills-forward>.

TestGorilla. 2024. *The State of Skills-Based Hiring 2024*. Amsterdam: TestGorilla. <https://www.testgorilla.com/state-of-skills-based-hiring/>.

LinkedIn. 2024. *Global Talent Trends 2024*. Sunnyvale, CA: LinkedIn Economic Graph. <https://economicgraph.linkedin.com/talent-trends/>.

European Parliament and Council. 2024. *Regulation (EU) 2024/1689 Laying Down Harmonised Rules on Artificial Intelligence (AI Act)*. Brussels: Official Journal of the European Union. <https://eur-lex.europa.eu/eli/reg/2024/1689/oj>.

IBM. 2024. “Predictive Analytics in Workforce Planning.” IBM Corporate Communications. <https://www.ibm.com/thought-leadership/workforce-analytics/>.

Salesforce. 2024. “AI-Driven Retention and Talent Development.” Salesforce Case Studies. <https://www.salesforce.com/resources/case-studies/>.

SAP. 2024. “Predictive Attrition Modeling.” SAP SuccessFactors. <https://www.sap.com/products/hcm/workforce-analytics.html>.

Boeing. 2024. “Augmented Reality in Aircraft Assembly.” Boeing Corporate. <https://www.boeing.com/innovation/augmented-reality>.

Walmart. 2024. “VR Training and Workforce Development.” Walmart Corporate. <https://corporate.walmart.com/newsroom/innovation/>.

Fidelity Investments. 2024. “Fidelity Immerse: VR Training for Financial Services.” Fidelity Institutional. https://institutional.fidelity.com/app/item/RD_13569_46516/fidelity-immerse.html.

Gloat. 2024. “Internal Talent Marketplace: Workforce Agility Platform.” Gloat. <https://www.gloat.com/>.

Eightfold.ai. 2024. “AI-Powered Talent Intelligence Platform.” Eightfold.ai. <https://eightfold.ai/>.

Hilton Hotels. 2024. “AI-Driven Employee Engagement and Retention.” Hilton Corporate Communications. <https://www.hilton.com/corporate/>.

AI in Hiring: How Jobs Really Change Education

If you have applied for a job recently, there is a good chance no human read your application first. An algorithm did. It scanned your resume for keywords, scored your skills against a job description, and decided in seconds whether you deserved a closer look. If you made it past that gate, a chatbot may have asked you screening questions. If you answered well enough, your video interview might have been analyzed for tone, word choice, and response time before a recruiter ever pressed play.

This is not a glimpse of the future. This is hiring in 2025.

The way companies find and evaluate talent is being rebuilt from the ground up, and the consequences reach far beyond human resources departments. When employers change what they look for and how they look for it, they send signals backward through the entire education system. Universities adjust their programs. Governments revise their funding priorities. Workers rethink what credentials to pursue. Hiring is not just the end of the education pipeline. It is one of the most powerful forces shaping what education becomes.

This chapter maps how AI-driven hiring actually works today, where it falls short, and what it means for anyone navigating education, careers, or both.

How AI Screens You

The scale of AI in hiring is no longer experimental. In 2025, 72% of North American organizations use AI in their HR functions regularly. In Europe, the figure stands at 59%, though European adoption is growing faster year-over-year, gaining 5 percentage points compared to North America's 4. Across both regions, AI has moved from a curiosity to core infrastructure.

What does that infrastructure actually do to your application?

- **Resume parsing and matching.** When you submit a resume, natural language processing algorithms break it into structured data: job titles, skills, education, years of experience. The system then scores your profile against the job description using semantic matching, looking not just for exact keyword hits but for conceptually related terms. If your resume says “project management” and the job asks for “cross-functional coordination,” a good NLP system will still connect them. A poor one will not, and you will never know which kind you are dealing with.
- **Chatbot pre-screening.** Roughly 52% of organizations now deploy candidate engagement chatbots. These handle initial questions, collect availability and salary expectations, and sometimes ask behavioral screening questions. The answers are scored and ranked before any human involvement. For high-volume roles, this means hundreds of applicants are filtered down to a shortlist without a recruiter reading a single response.
- **Automated scheduling and analytics.** AI handles interview scheduling in over half of organizations, and hiring analytics platforms increasingly track quality-of-

hire metrics, time-to-fill, and source effectiveness. The goal is a continuous optimization loop where every stage of the hiring funnel generates data that refines the next cycle.

- **The agentic frontier.** The next step is beginning to arrive, though its trajectory remains uncertain. Deloitte projects that 25% of enterprises using generative AI will deploy autonomous AI agents in 2025, rising to 50% by 2027. If these projections prove accurate, the shift would be substantial: these are not simple chatbots but systems that can independently source candidates, initiate outreach, conduct screenings, and generate shortlists with minimal human intervention. In the United States, 99% of hiring managers already report using AI in at least some recruitment stage, with companies reporting 35% reductions in time-to-fill (vendor-reported averages across early adopters).

The transatlantic gap matters for job seekers. If you are applying to a U.S. tech company, assume AI is deeply embedded at every stage. If you are applying in Europe, AI is present but more likely to operate under human oversight, partly because regulations demand it and partly because European firms have adopted a more cautious “human-AI collaboration” model. In practice, this means European candidates are more likely to encounter a human decision-maker earlier in the process, though the initial screening is increasingly automated everywhere.

The bottom line: your resume, cover letter, and application responses are almost certainly being read by software before they are read by a person. Understanding this changes how you should prepare.

Skills-Based Hiring: Promise vs Reality

One of the most widely cited statistics in modern recruitment is that 81% of employers globally now use skills-based hiring, up from 73% in 2023. The promise is compelling: instead of filtering candidates by degrees and job titles, companies evaluate what people can actually do. This should dramatically expand who gets considered for any given role.

The numbers on talent pool expansion are striking. In the United States, skills-based hiring increases the eligible candidate pool by 15.9 times. In the United Kingdom, by 10.4 times. In Germany, by 6.4 times for Millennials and 6.5 times for Gen X. In France’s construction sector, the expansion could reach 23.4 times. For Gen Z, the effect is even larger: 17.6 times in the U.S. and 12.7 times in the UK.

These figures suggest a revolution. But the reality is more complicated – and the gap between rhetoric and practice is now measurable.

A 2024 study by the Burning Glass Institute and Harvard Business School examined what actually happened when companies removed degree requirements from job postings. The finding is sobering: dropping degree requirements increased the share of newly hired candidates without a bachelor’s degree by only around 3.5 percentage points in the affected roles. Extrapolated across the labour market, fewer than one in 700 new hires actually benefited from these reforms (Sigelman, Fuller, and Martin 2024). The 81% of employers who report using skills-based hiring are largely reporting intent, not outcome. Chapter 7 explains the structural reasons for this gap – the rational

inertia of degree requirements, the legal risk calculus, and the coordination problem that keeps individual hiring managers anchored to familiar filters even when corporate policy changes.

This does not mean the shift is illusory. It means it is slower and more uneven than the headline numbers suggest.

- **Where skills-based hiring actually works.** Technology, particularly software development, has moved furthest. Coding challenges, portfolio reviews, and technical assessments have partially replaced degree requirements at many firms. Creative industries, digital marketing, and some areas of manufacturing and logistics have followed. In these sectors, demonstrable competency genuinely competes with credentials.
- **Where degrees still dominate.** Regulated professions, including law, medicine, engineering, accounting, and education itself, remain firmly degree-gated. Even in less regulated fields, many large employers still use degrees as a screening filter. A company may claim to hire based on skills while its applicant tracking system still automatically ranks candidates with bachelor's degrees higher. The Burning Glass data suggests this is not the exception but the norm: most firms that announce skills-based hiring continue to hire degree holders at nearly the same rate as before.
- **The 77% plan vs the lived reality.** When surveys report that 77% of European companies “plan to prioritize” competency-based recruiting in 2026, it is worth noting the gap between planning and implementation. Skills-based hiring requires organizations to redesign job descriptions, retrain hiring managers, reconfigure their technology, and often fight internal resistance from leaders who trust the signals they grew up with. Many firms are in early stages of this transition, and some will stall.
- **What this means for individuals.** Do not assume that dropping out of a degree program because “skills matter more now” is a safe bet. It depends entirely on your field, your target employers, and your geography. In sectors where skills-based hiring is real, building a visible portfolio of demonstrated work is genuinely more valuable than another line on your transcript. In sectors where it is not, a degree remains a hard prerequisite. The safest approach is to build both: credentials where they matter and evidence of capability everywhere.

New Ways of Being Assessed

Beyond resumes and interviews, a new generation of assessment tools is changing what it means to “apply for a job.” Some of these methods are more fair, more accurate, and more engaging than traditional interviews. Others raise serious ethical concerns. All of them are reshaping what candidates need to prepare for.

- **VR and AR simulations.** Virtual and augmented reality have moved from novelty to serious assessment infrastructure in specific sectors. Walmart uses VR to place candidates in high-pressure retail scenarios, including busy stores, difficult customers, and inventory crises, observing decision-making under

stress. The U.S. Navy employs VR to teach and assess complex tasks like ship navigation and firefighting. UPS uses VR driving simulations that measurably reduce accident rates among trainees. In Europe, manufacturing firms create virtual factory floors where candidates troubleshoot equipment failures, and healthcare organizations simulate patient interactions to evaluate clinical judgment.

These simulations offer something traditional interviews cannot: standardized, observable evidence of how someone actually performs under realistic conditions. They reduce bias by focusing on demonstrated competency rather than background or credentials. They also create data that AI systems can analyze, comparing a candidate's performance patterns against those of top performers.

- **Gamification.** An estimated 83% of companies now use some form of gamified assessment (industry survey; definitions of “gamified” vary widely across respondents). These are game-based cognitive tests and problem-solving challenges designed to feel engaging rather than evaluative. They generate rich performance data on traits like pattern recognition, adaptability, risk tolerance, and collaboration style. For candidates, they can feel less stressful than traditional tests. For employers, they provide behavioral signals that resumes and interviews miss.
- **Project-based hiring.** Some companies skip the interview process almost entirely, asking candidates to complete real-world tasks or simulations. Software companies present coding challenges that mirror actual product problems. Marketing agencies ask for campaign concepts. Consulting firms pose case studies. The appeal is obvious: you see how someone actually works, not just how they talk about working. The risk is that it demands unpaid labor from candidates, a concern that is growing as these exercises become more elaborate.
- **Emotion AI.** This is the most controversial assessment technology. Emotion AI analyzes facial expressions, vocal tones, and engagement levels during video interviews, measuring micro-expressions and engagement patterns that human interviewers might miss. Some platforms evaluate not just what candidates say but how they say it, scoring for confidence, empathy, and authenticity.

The regulatory divide here is stark. The EU AI Act explicitly bans emotion recognition in workplace contexts, effective February 2025, due to privacy concerns and discrimination potential. In North America, restrictions are lighter, and some companies use emotion AI freely for high-volume customer-facing roles. This means that a candidate applying across the Atlantic may face fundamentally different assessment systems depending on jurisdiction.

The Regulation Divide

How AI is used in hiring depends enormously on where you are. The transatlantic regulatory gap is not a detail for lawyers. It shapes the experience of every job applicant.

- **The EU AI Act.** Europe has created the world's first comprehensive AI regulation, and it classifies AI systems used in employment as "high-risk." This classification triggers a cascade of legal obligations:
 - Human oversight: final hiring decisions must involve meaningful human review. Fully algorithmic decisions are legally precarious.
 - Transparency: candidates must be explicitly informed when AI participates in their evaluation.
 - Bias mitigation: pre-deployment testing for discrimination, representative training data, and continuous monitoring are mandatory.
 - Documentation: development processes, data sources, and testing results must be retained for at least five years.
 - Risk management: comprehensive risk assessment and mitigation measures must be in place.

The timeline is phased. February 2025 brought mandatory AI training requirements for companies. August 2025 requires member states to designate enforcement authorities. August 2026 brings full enforcement for high-risk recruitment AI. By 2027, penalties activate: up to 35 million euros or 7% of global turnover for serious violations.

Crucially, the Act applies to any company using AI outputs in EU hiring contexts, including U.S. or Asian firms recruiting European candidates. This extraterritorial reach forces multinationals to comply globally or maintain separate systems.

- **The North American approach.** The United States has no federal equivalent to the EU AI Act for hiring. Regulation is patchwork: some states and cities have introduced or are considering algorithmic hiring laws, but enforcement is inconsistent. Companies face fewer compliance constraints, which enables faster experimentation but also increases discrimination risks.
- **GDPR and the right to explanation.** Beyond the AI Act, Europe's General Data Protection Regulation gives candidates rights around automated decision-making. If a decision significantly affects you (and hiring certainly qualifies), you may have the right to an explanation of the logic involved. European AI vendors consequently emphasize transparent, interpretable models, while North American tools more often optimize purely for predictive accuracy.

Algorithmic bias in practice. The regulation gap matters because AI hiring systems can and do discriminate. The most famous example is Amazon's abandoned hiring tool, which penalized resumes referencing women's organizations and women's colleges because its training data reflected male-dominated hiring patterns. This is not a one-off failure. Any system trained on historical hiring data risks encoding the biases of the past. Research consistently finds higher false-negative rates for minority candidates, meaning qualified people from underrepresented groups are disproportionately screened out.

Europe's approach, requiring bias auditing and human oversight, does not eliminate this problem, but it creates legal accountability and forces companies to look for it. The U.S.

approach leaves the burden largely on candidates to discover and challenge discriminatory algorithms, a far harder task.

For individuals, the practical implication is this: if you suspect you have been unfairly screened out by an automated system, your options depend heavily on where the employer operates and which regulations apply. In Europe, you have more legal ground to stand on. In North America, less.

The European Talent Crisis

Europe faces a talent crisis that is not theoretical. It is structural, measurable, and self-reinforcing.

- **The numbers.** Germany alone had 137,000 unfilled IT specialist positions in 2022, with the deficit widening since. Across the EU, approximately 55% of enterprises report difficulties filling ICT specialist vacancies. The European Commission's target of 20 million ICT specialists by 2030 looks increasingly unrealistic; current totals hover around 9 million despite adding 500,000 specialists between 2020 and 2021.

Between 2023 and 2024, EU businesses adopting AI leapt from 8% to 13.5%, yet qualified AI professionals remain acutely scarce. Europe produces world-class AI research but suffers a persistent brain drain to the United States and China, where compensation packages and research budgets are larger. This creates a cycle: demand for machine learning engineers, data scientists, and AI product managers far outstrips supply, elongating hiring cycles and inflating salary expectations.

- **The SME disadvantage.** Small and medium enterprises, which form the backbone of the European economy, are hit hardest. Only 55% of EU SMEs achieve basic digital technology adoption. The barriers cascade: 45% report that skills shortages hinder technology adoption, 25% cite financing constraints, and 53% find it challenging to retain qualified personnel even when they manage to hire them. Large multinationals and well-known tech companies out-recruit SMEs through stronger employer brands and deeper pockets, leaving smaller firms dependent on training existing staff.
- **The speed problem.** European recruitment processes average 44 days to hire, significantly longer than the pace that competitive talent markets demand. Long approval chains, multiple interview rounds, and slow communication push candidates toward faster-moving competitors. In AI talent markets, where qualified candidates may receive multiple offers within weeks, these delays are effectively disqualifying.
- **Worker anxiety.** At the same time, 25% of Europe's workforce fears AI could jeopardize their jobs, and 74% believe companies will require fewer employees due to automation. This anxiety affects how people approach upskilling, career changes, and risk-taking. It also shapes candidate behavior: workers who fear displacement are less likely to engage openly with AI-driven assessment tools,

creating a tension between what employers need and what candidates are willing to accept.

The talent crisis is not just a human resources problem. It is a competitiveness problem that feeds directly into the education system. When companies cannot find qualified workers, they pressure universities to produce them faster, lobby governments to change funding and visa policies, and invest in corporate training programs that increasingly compete with traditional education. This is where hiring and education become the same conversation.

Feedback Loops: How Hiring Reshapes Education

This is the critical connection that most discussions of hiring technology miss. Employer practices do not exist in isolation. They feed back into the education system, reshaping it in ways that are sometimes deliberate and sometimes invisible.

- **The signal employers send.** When a critical mass of employers adopts skills-based hiring, the signal travels upstream. Students hear that portfolios matter more than transcripts. Parents reconsider whether a four-year degree is worth the investment. High school counselors adjust their advice. The signal is not always accurate, as we have seen, but it is powerful.

When employers deploy AI screening that prioritizes certain keywords, formats, or credential types, educational institutions respond. Universities add “career-ready” language to their program descriptions. Boot camps proliferate to fill gaps that traditional programs leave. Micro-credential providers design their offerings to match the skills taxonomy that applicant tracking systems recognize.

- **University program design.** The feedback loop to universities is increasingly explicit. Industry advisory boards shape curriculum. Employers who cannot find workers with specific AI, data, or cybersecurity skills fund new programs, endowed chairs, and research partnerships. The 137,000 unfilled IT positions in Germany are not just a labor market statistic. They are a pressure signal that has led to new computer science programs, expanded enrollment caps, and government investment in STEM education.

This can be healthy: education aligned with genuine labor market needs serves students well. But it can also distort priorities. When university programs chase employer demand too aggressively, they risk becoming vocational pipelines that neglect the broader intellectual formation, critical thinking, and civic preparation that universities are meant to provide. The feedback loop can optimize for what employers want right now at the expense of what society needs over decades.

- **Government policy.** Hiring patterns drive policy at the national and EU level. The talent crisis has pushed governments toward visa reforms to attract foreign tech workers, increased funding for digital skills education, and regulatory frameworks like the EU AI Act itself, which partly exists because AI hiring systems raised discrimination concerns that demanded a legislative response. The European Commission’s target of 20 million ICT specialists by 2030 is a direct policy response to employer demand signals.

- **Corporate learning as substitute.** When the hiring pipeline does not deliver the workers companies need, they build their own education systems. Corporate learning platforms, employer-sponsored degree programs, and internal mobility tools like Gloat and Fuel50 all represent companies stepping into territory traditionally occupied by universities. This is the feedback loop at its most visible: hiring difficulty creates corporate education, which in turn competes with traditional education for student attention and employer credibility.
- **The innovation maturity question.** Not all of these feedback loops operate at the same speed or with the same impact. It helps to think about hiring innovations along two dimensions: how mature the technology is and how transformative its impact. AI-powered sourcing tools and gamified assessments are now mainstream, meaning they are baseline expectations that any serious employer uses. Predictive analytics and internal mobility platforms are proven and scaling rapidly. AI agents and VR/AR assessments are emerging but carry the highest transformative potential. Blockchain credentials and emotion AI remain experimental.

This maturity framework matters for education. Mainstream hiring tools create immediate pressure for educational adaptation: if gamified assessments are standard, students need exposure to them. Emerging technologies create weaker, more speculative signals: it is too early for universities to redesign curricula around VR assessment, but smart institutions are watching.

- **The system diagram.** Picture the flows: employers define skills they need. AI systems screen for those skills. Universities and training providers adjust programs to produce those skills. Graduates enter the labor market. Employers evaluate them, update their requirements, and the cycle begins again. Governments regulate both the hiring technology and the education standards. Credential platforms and corporate learning programs operate as alternative pathways that sometimes complement and sometimes compete with traditional education.

This is not a one-directional pipeline from education to employment. It is a loop, and understanding the loop is essential for anyone making education or career decisions. The skills that AI hiring systems prioritize today will shape the programs universities offer tomorrow, which will shape the workforce of the next decade.

What This Means for You

If the preceding sections feel abstract, here is what they translate to in practice.

- **Keep evidence of your work.** In a world where skills-based hiring is expanding and AI systems analyze portfolios, the most valuable thing you can do is maintain visible proof of what you can actually do. This means projects with tangible outputs: code repositories, writing samples, design portfolios, case studies, data analyses, video presentations. Evidence beats claims.
- **Do not rely on a single credential.** The “signalling stack” is replacing the single-credential model. A degree still matters in many fields, but increasingly it is one

layer among several. Micro-credentials, certifications, project portfolios, and demonstrated AI fluency all contribute to how you are evaluated. Build breadth across these layers rather than betting everything on one.

- **Learn to work with AI, not just next to it.** When 81% of employers say they use skills-based hiring and AI tools pervade every stage of recruitment, the ability to work effectively with AI systems is itself a skill that employers screen for. This does not mean you need to become a machine learning engineer. It means you should be comfortable using AI tools in your daily work, understanding their outputs, and knowing their limitations. Employers increasingly test for this directly.
- **Understand how you are being evaluated.** Know that your resume is likely being parsed by NLP systems. Structure it clearly, use standard formatting, and include the specific skills and terms from job descriptions where they honestly apply. If you are in Europe, know that you have the right to be informed when AI participates in your evaluation and, in many cases, the right to an explanation.
- **Build across borders.** The transatlantic regulatory divide means that hiring practices differ significantly between Europe and North America. If you are open to working in both markets, understand that European applications may involve more human touchpoints and stronger privacy protections, while North American processes may be faster but more aggressively automated. Tailor your approach accordingly.
- **Watch the feedback loops.** Pay attention to what employers in your target field actually hire for, not what they say they hire for. Look at job postings, talk to people who have been recently hired, and examine what skills and evidence types are showing up in assessments. This real-time signal is often more useful than any career guide or university prospectus, because it reflects the current state of the loop between employer needs and educational offerings.

Do not panic, but do not wait. The talent crisis and the AI transformation of hiring are creating genuine disruption, but they are also creating opportunity. Fields with severe shortages, particularly in technology, data, and cybersecurity, offer strong prospects for anyone willing to build demonstrable skills. The 137,000 unfilled IT roles in Germany alone suggest that demand is far from hypothetical. The question is not whether opportunity exists but whether you are preparing in ways that hiring systems, both human and algorithmic, can recognize.

The relationship between education and hiring has never been static. What is different now is the speed and scale of the feedback loop. AI hiring tools make employer preferences more legible, more immediate, and more consequential for educational choices. Understanding this loop, and positioning yourself within it, is one of the most practical things you can do for your career.

Annotated Source List

This chapter synthesizes data from Deloitte, TestGorilla, LinkedIn, and Eurostat surveys, as well as EU AI Act regulatory analysis. Sources are referenced narratively with evidence-quality labels throughout.

- The chapter's AI adoption statistics draw on Deloitte's annual global human capital trends surveys covering thousands of organizations across North America and Europe (Tier 2–3: multi-site employer surveys with self-reported data).
- Skills-based hiring prevalence and talent pool expansion figures come from TestGorilla's annual State of Skills-Based Hiring report, which surveys employers globally on hiring practices and candidate evaluation methods (Tier 2–3: large-scale multi-employer survey). The gap between reported adoption and actual hiring outcomes is documented by Sigelman, Fuller, and Martin (2024) at the Burning Glass Institute / Harvard Business School, using longitudinal hiring data rather than self-reported surveys (Tier 2: observational hiring data from job postings and employment records).
- LinkedIn's Global Talent Trends and Economic Graph data provide labour market analytics on skills demand, hiring velocity, and credential patterns derived from platform behavioural data across hundreds of millions of profiles (Tier 2: large-scale observational data from platform analytics).
- European talent shortage statistics, including Germany's 137,000 unfilled IT positions and EU-wide ICT specialist gaps, draw on Eurostat and Bitkom industry association surveys (Tier 2: national statistical office and industry association data).
- The EU AI Act discussion draws on the published regulatory text, European Commission guidance documents, and legal analyses of high-risk classification and enforcement timelines (Tier 2: primary regulatory sources).
- Amazon's abandoned hiring tool is documented through investigative journalism by Reuters (Tier 3: single-company case study).
- VR and gamification assessment data draws on vendor case studies from Walmart, UPS, and the U.S. Navy (Tier 3–4: single-organization implementations).
- The 83% gamification adoption figure comes from an industry survey with varying definitions across respondents (Tier 3).
- Agentic AI deployment projections are Deloitte estimates based on enterprise adoption modelling (Tier 4: vendor-reported forecasts).
- The 35% time-to-fill reduction is a vendor-reported average across early adopters without independent verification (Tier 4).

References

Deloitte. 2025. *2025 Global Human Capital Trends*. New York: Deloitte Insights.
<https://www2.deloitte.com/us/en/insights/focus/human-capital-trends.html>.

- TestGorilla. 2024. *The State of Skills-Based Hiring 2024*. Amsterdam: TestGorilla. <https://www.testgorilla.com/state-of-skills-based-hiring/>.
- LinkedIn. 2024. *Global Talent Trends 2024*. Sunnyvale, CA: LinkedIn Economic Graph. <https://economicgraph.linkedin.com/talent-trends/>.
- Eurostat. 2024. “ICT Specialists in Employment.” European Commission. https://ec.europa.eu/eurostat/statistics-explained/index.php?title=ICT_specialists_in_employment.
- Bitkom. 2023. “IT-Fachkräftelücke: 137.000 offene Stellen.” Bitkom. <https://www.bitkom.org/Presse/Presseinformation/IT-Fachkraefteluecke-137000>.
- European Parliament and Council. 2024. *Regulation (EU) 2024/1689 Laying Down Harmonised Rules on Artificial Intelligence (AI Act)*. Brussels: Official Journal of the European Union. <https://eur-lex.europa.eu/eli/reg/2024/1689/oj>.
- European Commission. 2025. “AI Act Implementation Timeline and Guidance.” Brussels: European Commission. <https://digital-strategy.ec.europa.eu/en/policies/regulatory-framework-ai>.
- Dastin, Jeffrey. 2018. “Amazon Scraps Secret AI Recruiting Tool That Showed Bias Against Women.” *Reuters*, October 10. <https://www.reuters.com/article/us-amazon-com-jobs-automation-insight-idUSKCN1MK08G>.
- Walmart. 2024. “VR Training and Workforce Development.” Walmart Corporate. <https://corporate.walmart.com/newsroom/innovation/>.
- UPS. 2024. “VR Driver Training Safety Program.” UPS Corporate Communications. <https://about.ups.com/us/en/social-impact/people/safety.html>.
- Sigelman, Matt, Joseph Fuller, and Alex Martin. 2024. *Skills-Based Hiring: The Long Road from Pronouncements to Practice*. Boston: Burning Glass Institute / Harvard Business School. <https://www.burningglassinstitute.org/research/skills-based-hiring-2024>.
- World Economic Forum. 2025. *The Future of Jobs Report 2025*. Geneva: World Economic Forum. <https://www.weforum.org/publications/the-future-of-jobs-report-2025/>.
- GDPR. 2018. *General Data Protection Regulation (EU) 2016/679*. Brussels: Official Journal of the European Union. <https://gdpr-info.eu/>.

Universities as Public Institutions in an AI Era

The previous chapters have traced how AI is unbundling the university, how new models are emerging, who gets left behind, and how employers and platforms are reshaping the demand side of education. Much of that story emphasizes disruption, efficiency, and market alignment. This chapter asks a different question: What would we lose if universities became nothing more than credential factories optimized for employer satisfaction and platform economics?

The answer, it turns out, is quite a lot. Universities are not just teaching machines. They are among the few institutions in modern societies that combine basic research, civic formation, cultural preservation, and open inquiry under one roof. An AI era does not eliminate the need for these functions. It makes them more urgent.

Why Societies Still Need Universities

It is tempting to look at the rise of AI tutors, micro-credentials, and employer-sponsored learning ecosystems and conclude that universities have been outflanked. If content delivery can be automated, if credentials can be unbundled, and if employers can train their own workers, what purpose does the university still serve?

The short answer is that universities do things no other institution is designed to do.

- **Research as a public good.** The most consequential knowledge breakthroughs of the past century – mRNA vaccines, the internet, CRISPR gene editing, the foundations of machine learning itself – emerged from decades of publicly funded basic research conducted at universities. This kind of research has no immediate customer. It follows questions rather than quarterly targets. It often fails. And it requires the kind of patient, long-horizon institutional commitment that neither corporations nor platforms are structured to provide. Companies invest in applied research that serves their product roadmaps. Governments fund mission-driven programs with political timelines. Universities, at their best, sustain inquiry across generations, allowing researchers to pursue questions whose payoff is uncertain and distant. Without this function, societies become dependent on private R&D agendas and lose the capacity to generate knowledge that serves the public interest rather than shareholder value.
- **Democratic citizenship.** Universities have historically served as spaces where young people learn to argue, deliberate, encounter difference, and develop the habits of democratic participation. This is not a romantic abstraction. In a world saturated with AI-generated content, algorithmic recommendation systems, and deepfakes, the capacity to evaluate evidence, recognize manipulation, and engage in reasoned public discourse is not a luxury – it is infrastructure. If societies outsource education entirely to platforms and employers, they may produce skilled workers but not citizens equipped to sustain democratic governance.
- **Cultural memory and critical interpretation.** Universities house the humanities, the arts, and the social sciences – disciplines that preserve,

interpret, and challenge cultural inheritance. Philosophy departments ask what a good life looks like. History departments examine how societies have failed and recovered. Literature programs cultivate the imaginative empathy required to understand experiences unlike one's own. These functions do not show up in employer satisfaction surveys or skills-gap analyses. They are nonetheless essential to societies that want to understand themselves and make deliberate choices about their futures.

- **Local and regional development.** Universities anchor local economies, attract talent, sustain healthcare systems through teaching hospitals, and serve as civic institutions in their communities. A mid-sized city that loses its university does not just lose a school. It loses a research hub, a cultural center, an employer, and a node in networks of innovation and social mobility. Platform-based education, by contrast, is geographically indifferent. It can serve learners anywhere, but it anchors nothing.

The Crucible of Character – Made Concrete

Earlier in this book, we introduced the idea that surviving universities would become “crucibles of character” – environments that forge resilience, judgment, and identity rather than merely transmitting content. The critics rightly pointed out that this metaphor, while appealing, remained vague. What does a crucible of character actually look like in institutional practice?

It looks like specific structures, experiences, and commitments that cannot be replicated by an AI tutor or a micro-credential platform.

- **Residential and cohort experiences.** Living alongside people from different backgrounds, negotiating shared spaces, managing conflict without parental intervention, and building friendships across lines of difference – these experiences are formative in ways that online interaction cannot reproduce. The residential college is not an outdated luxury. It is a structured environment for developing social competence, emotional resilience, and the capacity to live with disagreement. Not every student needs four years of residential education, but the developmental value of even one or two years of intensive community living is well documented and difficult to replicate digitally.
- **Project-based community work.** Students who spend a semester working with a local nonprofit on affordable housing policy, or collaborating with a municipal government on climate adaptation planning, learn things that no simulation can teach: how institutions actually function, how stakeholders with competing interests negotiate, how good intentions collide with bureaucratic reality. These experiences ground abstract knowledge in practical judgment. They also connect universities to their communities in ways that create mutual accountability.
- **Deliberative forums and democratic practice.** Student governments, town halls, and structured debates are not extracurricular decorations. They are laboratories for democratic competence. When students must craft arguments, listen to opposing views, build coalitions, and accept outcomes they disagree with, they develop capacities essential for citizenship. Universities that treat

these activities as optional enrichment rather than core pedagogy miss an opportunity that no employer or platform will fill.

- **Student newspapers, journals, and creative outlets.** Running a student newspaper requires editorial judgment, ethical reasoning, collaboration under deadline pressure, and accountability to a real audience. These are precisely the capacities that AI makes more important, not less. When anyone can generate polished text in seconds, the ability to distinguish credible reporting from plausible fabrication becomes a survival skill. Student media, literary journals, theater productions, and art exhibitions are not relics. They are training grounds for the critical and creative capacities that define human contribution in an AI-saturated world.
- **Debate clubs and intellectual sparring.** Formal and informal debate cultivates the ability to construct and dismantle arguments, to think on one's feet, to hold positions provisionally and revise them in light of evidence. These are skills that AI can simulate but not instill. The experience of being publicly wrong, of having one's reasoning dismantled by a peer, and of rebuilding a stronger argument is developmental in a way that interacting with a chatbot is not.
- **Long-term basic science laboratories.** Undergraduate and graduate students who spend years in a research lab – designing experiments, confronting failure, learning to interpret ambiguous data, contributing to knowledge that will not be published for years – develop intellectual patience and rigor that no short-form credential program can cultivate. The lab is a crucible in the most literal sense: it subjects students to sustained intellectual pressure and produces durable capacities for original thought.

The common thread is that these experiences involve friction, stakes, and human interaction in ways that cannot be compressed into a module or automated by an algorithm. They require physical or temporal co-presence, genuine risk, and accountability to real people. This is what the crucible metaphor means in practice.

The Selectivity Paradox: How Crucibles Avoid Reproducing Elite Bias

The preceding list of crucible practices – residential cohorts, community projects, deliberative forums, student media, debate, research labs – raises an uncomfortable question: are these not precisely the experiences that elite universities already offer? And if so, does calling for their expansion solve the problem of stratification, or simply suggest that broad access to elite education?

The answer depends entirely on design. The same pedagogies can either amplify or mitigate existing inequalities, and the difference lies not in the pedagogy itself but in the *institutional commitment to support* that surrounds it.

Elite universities work for self-directed, already-confident students because those students bring strong foundational skills, stable home environments, and implicit knowledge of institutional norms. A debate club at such an institution assumes participants already know how to argue, can tolerate social discomfort, and have role models for intellectual risk-taking. For these students, debate is enrichment. It accelerates what they already possess.

But a debate club designed as inclusive pedagogy – for students without prior experience, from backgrounds where public argument is not a family norm, carrying anxiety or trauma that makes public speaking feel unsafe – requires something entirely different. It requires a coach who understands developmental scaffolding. It requires smaller cohorts so students get individualized feedback rather than sinking into the background. It requires explicit instruction on how to construct arguments, not just exposure to good examples. It requires psychological safety work before intellectual combat. It requires mentors who look like students and come from similar backgrounds, to make the possibility of belonging feel real.

This is not a minor add-on. It is a structural redesign. And it costs more, not less, than the elite version.

The same principle applies to community projects. An elite student who grows up watching her parents navigate bureaucracy, negotiate with public officials, and think strategically about social change will treat a semester of work with a city government as applied learning. A first-generation student whose family experiences institutions primarily as imposing forces – welfare offices, courts, immigration agencies – may experience the same project as intimidating and incomprehensible, without explicit training in how to interpret institutional dynamics, ask questions, and recognize their own agency.

To make community work inclusive, universities must embed structured reflection, provide explicit frameworks for understanding institutional design, connect the work to broader concepts of power and change, and ensure that students from under-resourced backgrounds do not arrive at projects expecting to perform competence they do not yet have.

Again: this is more expensive and more labor-intensive than an elite version. It requires faculty and staff trained in developmental education. It requires smaller project teams. It requires follow-up mentoring. But it is the only way the pedagogy can honestly serve equity rather than simply redistribute elite advantage.

The core design principle is this: **the hybrid university must treat all the high-contact, high-stakes, character-forming experiences that make it excellent as pedagogies to be learned, not entry requirements to be selected for. This reversal – from selection to development – is what distinguishes a public crucible from an elite one.**

It has concrete implications:

- **Residential experiences must include intensive support.** Not every student needs four years on campus; some benefit more from part-time engagement. But the students who most benefit developmentally from residential living are often those least likely to afford it or feel they belong. Subsidizing housing costs and embedding mental health services, peer mentoring, and academic support within residential communities is not a luxury. It is what makes the residential crucible work for populations beyond the already-advantaged.
- **Community projects require structured mentorship and institutional partnership.** Students do not simply “apply learning.” They need frameworks to interpret what they encounter, adults who can model professional and ethical

judgment, and explicit teaching about how institutions work, where power resides, and what it means to challenge assumptions respectfully. This requires partnerships where community organizations and universities jointly design projects, provide mentoring, and treat student learning as a shared responsibility.

- **Deliberative forums need explicit preparation and psychological safety.** Not everyone arrives at university ready to argue in public. Building inclusive deliberative spaces requires teaching argument skills, creating structures where all voices are heard (not just the loudest), facilitating productive disagreement, and providing private spaces where students can process discomfort. It means faculty who understand how identity, power, and trauma shape participation, and who can make intellectual risk-taking feel survivable.
- **Student media, theater, and creative outlets require economic support.** These activities are often the first casualties at under-resourced institutions. But they are precisely where students who have never seen themselves as writers, artists, or leaders discover they can be. Making them sustainable requires institutional commitment: paid advisors, equipment budgets, and explicit recognition that these activities are pedagogy, not decoration.
- **Research labs must be accessible beyond the already-prepared.** Not every undergraduate should spend years in a lab; many thrive elsewhere. But laboratories should include pathways for students from under-resourced schools who have never encountered research, may not know it is a viable path, and need extra scaffolding to get there. This means mentors, bridge programs, and explicit teaching of lab culture and norms.

The principle throughout is the same: crucible experiences are powerful precisely because they involve stakes, friction, and real human judgment. But they need not be exclusive. The cost of inclusion is genuine institutional investment in development and support. The benefit is that the same experiences that shape the character of privileged learners can shape the character of students from any background – if the institution commits to making that possible.

This is the crucible model’s answer to the selectivity paradox: not avoiding high-contact, high-stakes experiences, but redesigning them so they develop capacity rather than assuming it, and funding them as core institutional functions rather than optional extras.

The Cost Architecture of Inclusive Crucibles

The assertion that crucible redesign is “more expensive, not less” requires specification. Without cost estimates, the proposal risks remaining aspirational. What do inclusive crucibles actually cost?

Emerging evidence from scalable interventions provides a baseline. CUNY’s ASAP programme – which provides accelerated progress, dedicated advisors, small cohorts, and wraparound support including childcare, emergency funds, and transportation subsidies – costs approximately \$3,500 per student annually in direct instructional and support costs above baseline tuition funding. This intensive support model has

achieved completion rates above 70%, compared to roughly 40% at comparable institutions without such investment.

Translating this to the full crucible model suggests realistic per-student investment requirements. The core supports outlined above – developmental coaching at ratios enabling genuine feedback, smaller cohorts reducing instructor load, structured mentorship with trained peers, psychological safety infrastructure, and community project coordination – likely add \$2,800 to \$3,800 per student annually at non-elite institutions. This breaks down roughly as: coaching and advising ratios (\$1,200–\$1,500), smaller cohort costs (\$700–\$900), community partnership coordination and student release time (\$600–\$700), and mental health and peer support infrastructure (\$400–\$600). These are not marginal enhancements. They represent a genuine reallocation of resources toward the intensive contact that makes inclusive pedagogy work.

The question is not whether this is affordable in absolute terms, but how it fits within diversified funding architectures. At institutions receiving 35–40% of revenue from public funding, these costs can be embedded in per-student grants and performance funding if formulas reward outcomes for underserved populations, not just enrolment. At tuition-dependent institutions, the cost becomes prohibitive unless coupled with income-contingent repayment and substantial subsidy for low-income students. The protective funding model described later in this chapter – where no single revenue source exceeds 40% of income – is not incidental to the crucible model. It is prerequisite. Without it, universities cannot afford the support infrastructure that makes crucible pedagogy inclusive rather than exclusive.

Academic Freedom and Institutional Autonomy in an AI Era

Academic freedom – the right of scholars to pursue research and teach without political or commercial interference – has always faced threats. In the AI era, the threats take new forms that universities are not yet well equipped to handle.

- **Platform dependence.** As universities adopt AI-powered learning platforms, adaptive tutoring systems, and cloud-based research infrastructure, they become dependent on a small number of technology companies for core functions. This dependence creates risks. Platform providers make design decisions that shape pedagogy: what gets measured, what gets optimized, what counts as progress. When a university’s learning management system is built on a proprietary AI engine, the company that controls the engine has subtle but real influence over what and how students learn. If the platform changes its terms, raises prices, or is acquired, the university may find that its core educational infrastructure is controlled by an entity with no obligation to academic values.
- **Data sovereignty and research integrity.** AI research increasingly depends on access to large datasets, cloud computing resources, and proprietary models. Universities that rely on commercial providers for these resources face questions about who owns the data, who controls the models, and whether research findings can be independently verified when they depend on closed systems. The principle that research should be transparent and reproducible is difficult to uphold when critical infrastructure is proprietary.

- **Political and commercial pressure on inquiry.** AI-generated content and algorithmic amplification make it easier for political actors and commercial interests to target researchers whose findings are inconvenient. Climate scientists, public health researchers, and scholars studying platform harms have already experienced coordinated harassment campaigns amplified by social media algorithms. Universities must develop governance structures and support systems that protect researchers in this environment, rather than leaving individual scholars to fend for themselves against algorithmically scaled pressure.
- **Governance challenges.** University governance structures – senates, boards, faculty committees – were designed for a slower, more stable institutional environment. They are often poorly equipped to make rapid decisions about AI adoption, data partnerships, and platform contracts. The risk is that these decisions get made by administrators or technology offices without meaningful faculty input, eroding the shared governance that protects academic freedom. Universities need governance mechanisms that can move at the speed of technology adoption while preserving the deliberative processes that protect institutional autonomy.

Civic Education in an Age of Algorithmic Manipulation

The preceding sections argue that universities are sites of democratic citizenship formation, and that this function matters more in an AI era, not less. This argument requires immediate historical clarification. For most of democratic history, universities played a minor role in civic formation. Working-class political consciousness was cultivated through unions, mutual aid societies, and party organisations. Religious communities developed moral reasoning and ethical deliberation. Military service, journalism, and local assemblies all contributed to civic competence. These institutions remain important, and the claim here is not that universities are the only site of civic formation – they are not, nor should they be.

The argument is narrower and more specific: in an information environment saturated with algorithmic manipulation, deepfakes, and microtargeted misinformation, the combination of capacities that universities can systematically cultivate – algorithmic literacy, source epistemology, deliberative practice, and sustained encounter with complexity – has become more critical to democratic resilience, not less. Other institutions still form citizens. But they operate in the same algorithmic environment, often without the same resources for teaching people to navigate it. The risk is not that universities have historically been the sole source of civic competence. It is that, as other institutional sites of civic formation have weakened – union membership has declined, local journalism has collapsed, political parties have hollowed out – eliminating universities as a remaining counterweight to algorithmic polarisation would leave societies dangerously exposed.

But “citizenship formation” remains a phrase rather than a practice. What does civic education actually entail in conditions where algorithmic systems mediate political discourse, where deepfakes can simulate public figures, where microtargeted content can isolate people into hostile epistemic bubbles? And how would an institution

measure whether civic education is working when the effects are distributed across years and contexts far from campus?

The traditional model of civic education assumed a relatively stable information environment and face-to-face deliberation. Students were exposed to texts representing different perspectives, guided through argument construction, and given practice in structured debate. The implicit theory was that exposure plus practice would cultivate judgment and the habit of reasoning through disagreement. This model has limits even in stable conditions – not all students benefit equally from Socratic dialogue, and classroom diversity does not automatically produce genuine encounter across lines of difference. But in an AI-saturated information landscape, these limits become acute.

Consider what happens when a university student leaves a deliberative seminar and returns to an algorithm-curated social feed. The feed shows her content that reinforces her existing beliefs, shows her engagement metrics that reward inflammatory responses, and shows her comments flagged by peers who already agree with her. The algorithm does not know she just learned to steel-man opposing arguments or hold provisional beliefs. It knows her engagement history, and it optimizes for continued engagement, which correlates with emotional activation and identity confirmation. One deliberative seminar a semester competes against hours daily of algorithmic feedback pushing toward polarization. This is not a pedagogical failure. It is a mismatch between the pace and force of educational intervention and the pace and force of the environment the student inhabits outside the classroom.

Effective civic education in this context requires explicit curriculum design around three capabilities that traditional liberal arts cultivate implicitly but do not name. First is **algorithmic literacy** – not “AI ethics” as an abstract seminar, but concrete understanding of how recommendation systems work, what they optimize for, what they cannot see, and how they can be manipulated. This is not technical knowledge for its own sake. It is essential context for making sense of the information environment students encounter. Second is **source epistemology** – the capacity to evaluate not just the truth of a claim but the credibility of the source, the incentives embedded in the platform delivering it, and the alternative sources that might exist but are not recommended to you. This is version of media literacy for the algorithmic age. Third is **deliberative muscle** – regular, structured practice in good-faith disagreement with people who disagree with you on matters that matter, facilitated by someone trained to hold the space and interrupt bad-faith moves. The deliberative seminar does this, but it needs to be embedded across the curriculum, not isolated in a single course.

What does this look like institutionally? It means civic education is not a requirement box checked by a single civics course or service-learning assignment. Instead, it is integrated into disciplinary content: a history course on the American civil war asks how competing communities remember and interpret the same events; a biology course on pandemic policy asks how scientific uncertainty gets communicated in public discourse and how algorithms shape belief in scientific findings; a literature course uses narrative theory to examine how stories shape identity and group membership. It means students across disciplines encounter the question: What are the ethical dimensions of the knowledge I am learning, and how does power shape who gets believed? It means governance is not decorative: student governments make decisions

with real consequences and real stakes, and faculty are present to coach students through the complexity of collective decision-making under disagreement.

Measuring success is genuinely difficult. You cannot give a pre-post test to measure whether someone has become a better citizen. But you can measure intermediate outcomes with imperfect proxies. Did students encounter diverse perspectives as part of their core coursework? (Track demographic diversity and perspective diversity in required courses.) Did students engage in deliberation with peers across lines of difference? (Survey questions about intellectual friendship across lines of disagreement, interview data from deliberative exercises.) Did students learn to recognize and resist algorithmic persuasion? (Test problems where students evaluate sources and claims in algorithmically curated feeds, assess quality of reasoning about information sources.) Did students develop sustained engagement in civic institutions? (Track participation in student governance, community boards, electoral politics in the years following graduation.) Did they develop the habit of thinking systemically about tradeoffs in public policy? (Assess quality of reasoning in capstone projects addressing policy questions.)

None of these is a perfect measure, and institutional data collection carries costs. But universities currently measure labor-market outcomes extensively – graduate employment, salary, employer satisfaction. If civic formation is genuinely part of the university's public mission, measurement must follow. The absence of data is not neutral. It is a choice to treat workforce metrics as real and civic outcomes as aspirational.

The risk is that civic education, once measured and systematized, becomes another checkbox on the curriculum audit form, lost in the same pressures toward efficiency and measurable outputs that have narrowed education toward labor-market focus. The guard against this is institutional leadership that names civic formation as non-negotiable and protects it from the logic of optimization. Civics seminars should not be evaluated primarily on student satisfaction or completion rates, because those metrics reward making the content easy rather than substantive. Civics education is supposed to be hard – to require students to encounter positions they find repugnant, to practice disagreement without resolution, to sit with complexity rather than settle into comfort. If the institution treats discomfort as a sign of failed pedagogy and optimizes for satisfaction, civic education becomes a thing done rather than a transformation embodied.

The deepest argument for non-work purposes of education is this: democracies require citizens who can hold institutions accountable, who can imagine alternatives to the status quo, who can build coalitions across difference. No employer will specify these competencies in a job description. No ranking system will reward universities for cultivating them. But no democratic society can survive without them. Universities that abandon this function are not adapting to the market. They are abandoning one of the few institutions equipped to sustain democracy. The market will never fund this function adequately because its benefits accrue to the public, not to individual learners or employers. This is precisely why public funding for universities remains essential – not as charity, but as infrastructure for self-governance.

Funding Models for the Hybrid University

One of the hardest practical questions is how universities can sustain their public mission – research, civic education, cultural preservation, equitable access – while adapting to a world of modular credentials, corporate partnerships, and platform economics. The answer is not a single funding model but a diversified architecture that protects core functions from market volatility.

- **Public funding as the anchor.** Basic research, civic education, and equitable access are public goods. They generate benefits that cannot be captured by individual consumers or employers, which means markets will systematically underfund them. Public funding – whether from national governments, regional authorities, or supranational bodies like the European Commission – remains the indispensable foundation. The trend toward performance-based funding formulas that reward completion rates and employment outcomes is understandable but dangerous: it incentivizes universities to optimize for measurable short-term outputs at the expense of long-term research and broad human development. Funding formulas must include metrics that protect basic research capacity, humanities and social science programs, and access for underrepresented populations.
- **Tuition as a contribution, not a price.** Tuition serves a legitimate function: it signals that education has value, and it provides revenue. But when tuition becomes the primary funding source, universities become consumer-service organizations that optimize for enrollment and satisfaction rather than intellectual rigor and public mission. The hybrid model requires tuition to be one revenue stream among several, kept at levels that do not create prohibitive debt or exclude lower-income students. Income-contingent repayment systems, means-tested fee structures, and publicly funded tuition subsidies are all mechanisms that can keep tuition functional without making it dominant.
- **Corporate partnerships with guardrails.** Employer-sponsored learning, co-designed curricula, and industry research partnerships can provide significant revenue and ensure labor-market relevance. The risk, as the chapter on pitfalls documented, is corporate capture: universities tailoring programs to employer specifications at the expense of critical inquiry and broad education. The guardrail is structural separation. Universities can offer employer-aligned professional certificates and micro-credentials through continuing education divisions while protecting the autonomy of degree programs, research agendas, and core curricula from direct corporate influence. Revenue from corporate partnerships should cross-subsidize public-good functions rather than replace public funding.
- **Platform and credential revenue.** As universities develop online offerings, micro-credentials, and digital learning platforms, they can generate revenue from learners beyond their traditional student body. This is a legitimate and potentially significant income stream, but it must be managed carefully. If platform revenue becomes the dominant revenue source, it will drive the same consumer-optimization dynamics as tuition dependence. Platform offerings should be

designed as complements to, not substitutes for, the residential, research-intensive, and civic functions that define the university's public role.

- **Endowments and philanthropy.** For institutions that have them, endowments provide a crucial buffer against market fluctuations and political shifts. But endowments are radically unequal: a handful of elite institutions control the vast majority of endowment wealth. Philanthropy can fund innovation and experimentation, but it cannot substitute for systemic public investment in higher education.

The key principle is diversification with purpose. Each revenue stream serves different functions, and each carries different risks. A university that depends on any single source – whether public funding, tuition, corporate revenue, or platform income – will eventually have its mission shaped by that source's incentives. The hybrid funding model protects the public mission by ensuring that no single funder has veto power over institutional direction.

A Concrete Funding Architecture: How Mission Survives Market Pressure

The five funding streams described above are necessary components, but they are not a strategy. A strategy requires specifying how much revenue comes from each source, how that composition changes the incentives institutions face, and what safeguards protect core mission when markets shift. This section translates principle into practice by modeling realistic funding compositions and the choices they require.

The challenge is that higher education systems across OECD countries start from radically different funding structures. What counts as a dangerous dependency in one system may be unremarkable in another. A protective architecture must therefore be grounded in structural principles that hold across systems, even when the numbers look different. We begin with three cases that illustrate the range.

Case 1: The American mid-tier public research university

Consider a mid-tier US public research university with 12,000 students, significant research output in STEM and social sciences, and a commitment to regional accessibility. In 2025, such an institution typically derives roughly 30% of revenue from state appropriations, 35–40% from tuition and fees, 10–15% from research funding, 5–10% from philanthropic revenue, and the remainder from auxiliary services, licensing, and other sources (OECD 2025a).

This composition leaves the institution deeply vulnerable. State funding has declined for decades, making the university heavily dependent on tuition, which creates incentives to prioritise recruitment, retention, and satisfaction over intellectual rigour. Research funding is competitive and volatile, favouring applied work with immediate payoffs. Philanthropic revenue is concentrated in a handful of elite institutions and barely relevant for the majority.

The structural problem is clear: tuition dependence at 35–40% means that students are no longer primarily learners to be educated – they are revenue units to be retained. A mission-protective model would need to shift public funding back toward 35–40%,

reduce tuition dependence to 25–30%, and build new revenue from platforms and continuing education (5–10%) while keeping philanthropy as a supplement, not a pillar. Income-contingent repayment, used successfully in Australia and the Netherlands, can keep tuition functional without making it dominant. But none of this happens without political advocacy: universities must demonstrate public benefit persuasively enough to reverse the long-term decline in state investment.

Case 2: The German university

A German mid-tier university operates in a structurally different environment. Public funding – primarily from the respective Land – accounts for 55–65% of revenue. General tuition fees were abolished across all German states by 2014, with the sole exception of fees for non-EU students in Baden-Württemberg (1,500 EUR per semester). Semester contributions of 200–300 EUR cover administration and transit passes, not instruction. Competitive research grants (DFG, BMBF, EU framework programmes, industry contracts) have grown significantly – the DFG Förderatlas 2024 documents a 19% increase in third-party funding from 2019 to 2022 – and now constitute 25–35% of total revenue. Philanthropy accounts for 1–2% at most. Revenue from continuing education and platforms is below 1% (Destatis 2023; DFG 2024).

The German vulnerability is not tuition dependence but a different kind of imbalance: the steady shift from stable Grundmittel (core public funding for teaching and basic operations) toward competitive Drittmittel (third-party grants tied to specific projects and timelines). This shift creates incentives that are subtler than tuition dependence but structurally similar. Faculty spend increasing time writing proposals rather than teaching or pursuing unfunded inquiry. Research agendas drift toward topics that attract funding rather than topics that need investigation. Temporary contracts proliferate because grant-funded positions are inherently time-limited, contributing to the well-documented precarity of early-career researchers in the German system. And universities that are less successful at winning grants – often the smaller, regionally focused institutions – fall further behind, not because their teaching or research is worse but because competitive funding rewards scale.

A protective model for the German system would stabilise the ratio of Grundmittel to Drittmittel, ensuring that core public funding covers teaching, infrastructure, and a protected share of basic research – not just the minimum that keeps the lights on while grants fund everything else. It would also invest seriously in a revenue stream that German universities have largely neglected: continuing education and professional development, which could realistically contribute 5–10% of revenue if structured as a separate division with ring-fenced income. The near-absence of tuition is a strength that should be preserved, but it makes the system entirely dependent on the political willingness of state governments to fund universities adequately – and that willingness has not kept pace with rising costs and enrolment.

Case 3: The British mid-tier university

The UK system, particularly in England, has converged structurally with the American model more than any other European country. Following the 2012 reform that tripled the tuition cap to GBP 9,250 per year, public teaching grants collapsed. By 2024, direct government grants account for roughly 11–12% of total revenue at a typical mid-tier

university, down from 39% in 2006. Tuition fees – domestic and international combined – now constitute 52–53% of institutional income. Research grants and contracts contribute around 14%, but three-quarters of that funding flows to the 24 Russell Group universities. Philanthropy accounts for roughly 2%, though it is radically concentrated: Oxford and Cambridge together hold nearly half of all charitable donations to UK higher education. The remaining 21–22% comes from accommodation, conferencing, IP licensing, and other commercial activities (HESA 2024; House of Commons Library 2024).

The UK's distinctive vulnerability is its dependence on international student fees, which account for roughly 23% of total sector income. International tuition is uncapped and has risen steeply, creating a structural reliance on a revenue stream that is sensitive to geopolitical shifts, visa policy changes, and competition from other English-speaking countries. The Office for Students warned in 2024 that 72% of English universities could be operating at a deficit by 2025/26 if international recruitment falters. This is not a hypothetical risk – it is a live one.

A protective model for the UK would need to rebuild public funding toward 20–25% of revenue, reduce tuition dependence to 35–40% (primarily by diversifying away from international fee income as the balancing item), and invest in platform and continuing education revenue at 5–10%. The deeper structural challenge is political: the UK made a deliberate policy choice to shift the cost of higher education from the state to the student. Reversing that requires not just institutional advocacy but a change in the political consensus about who benefits from higher education and who should pay.

Three systems, one structural logic

The numbers differ, but the underlying dynamics are the same. Every higher education system faces the question of what happens when a single revenue source becomes dominant: tuition in the US and UK, competitive grants in Germany, state appropriations in Nordic systems. In each case, the dominant source reshapes institutional behaviour in ways that may conflict with the university's public mission.

The following principles hold regardless of where a university sits on the public-to-market spectrum:

Principle 1: No single revenue source should exceed 40% of total income. This is the diversification threshold. Once any source crosses it, that funder's incentives begin to dictate institutional priorities – whether the funder is a government ministry, a tuition-paying student body, or a research council. The specific sources that matter vary by system: in the US, the task is to reduce tuition below 40%; in Germany, to prevent Drittmittel from crossing the same threshold; in the UK, to rebuild public funding so that tuition does not dominate.

A critical caveat: This principle applies most directly to mixed-funding systems with genuine revenue diversification capacity – the United States, Australia, and post-2012 England. For structurally mono-funded systems – Nordic countries (typically 70%+ public funding), Germany before the recent expansion of Drittmittel, and most continental European universities – the 40% threshold should be read as a diagnostic tool, not a prescriptive target. The diagnostic question is: what happens to your institution when your dominant funder changes priorities? The answer in mono-funded

systems is structural vulnerability, precisely because diversification options are limited. The protective response is not wholesale restructuring toward an American model, which would be both politically unviable and potentially harmful to access equity. Rather, it is building modest but deliberate secondary revenue streams – continuing education, platform revenue, professional development – as resilience buffers. For German universities, this means stabilising the Grundmittel-Drittmittel ratio while selectively growing continuing education revenue from 1% to 5–7%. For Nordic systems, it means modest growth in professional education while maintaining strong public funding. The goal is not convergence on a single model, but sufficient shock-absorption within each system's structural constraints.

Principle 2: Public funding must anchor the functions that markets will not sustain. Basic research, civic education, broad access, and cultural preservation are public goods. Markets will systematically underfund them because their benefits are diffuse and long-term. Public funding – whether from national governments, regional authorities, or supranational bodies like the European Commission – is not one revenue stream among equals. It is the structural anchor that makes everything else possible. The form varies: Landesmittel in Germany, HEFCE/OfS grants in England, state appropriations in the US, direct government funding in Scandinavia. The function is the same.

Principle 3: Tuition, where it exists, must remain a contribution, not the primary revenue mechanism. Systems that rely on tuition for more than 30% of revenue create consumer dynamics that distort institutional purpose. Students become customers to be satisfied rather than learners to be challenged. Institutions optimise for enrolment and retention metrics rather than intellectual rigour. Where tuition exists, it should be means-tested, capped, and coupled with income-contingent repayment systems that protect access. Where tuition does not exist – as in most of Germany and Scandinavia – its absence is a feature worth defending, not a gap to be filled.

Principle 4: Competitive research funding must not substitute for core institutional funding. Research grants fund research production. They should not fund teaching, infrastructure, or basic operations. When they do – as increasingly happens in Germany and the UK – universities become grant-harvesting operations rather than institutions capable of sustaining long-term inquiry. Protecting basic research capacity requires institutional choices: dedicating a share of core revenue to unfunded inquiry, protecting faculty time from the proposal-writing treadmill, and resisting the pressure to treat all research through a commercialisation lens.

Principle 5: New revenue streams – platforms, continuing education, corporate partnerships – must cross-subsidise public goods, not replace public funding.

Platform revenue and professional education can realistically contribute 5–10% of revenue if structured as separate divisions with ring-fenced income. This is a genuine opportunity, especially for German and continental European universities that have barely tapped continuing education markets. But these revenues must flow toward making the public mission sustainable – subsidising access, funding basic research, supporting equity infrastructure – rather than becoming ends in themselves. A university that builds a profitable online division while its humanities faculty shrinks has not diversified. It has drifted.

Principle 6: Philanthropy is a supplement, not a structural pillar. In the US, philanthropy and endowments realistically contribute 5–10% for mid-tier institutions and much more for elite ones. In continental Europe, the figure is 1–2% and unlikely to change substantially in the medium term. Funding models that assume 10–15% from philanthropy are realistic only for a narrow slice of institutions, mostly in the anglophone world. For the majority of universities globally, philanthropy funds innovation and experimentation at the margin. It cannot and should not fund core operations, because donor influence over institutional mission is a direct threat to autonomy.

What these principles require – and what they do not solve

These principles require universities to be strategically intentional about what they do well and what they decline to do. A mid-tier research university cannot be all things to all students. The protective funding model forces clarity: What research is our distinctive contribution? Who are we primarily accountable to? What programs serve our mission? What programs generate revenue but do not serve it? The hard choice is often the last one: universities may need to discontinue or scale back programmes that are popular and profitable but ancillary to their core mission.

This is difficult. Programmes generate constituency and revenue. But a university that is all things is credible at nothing. The protective model requires institutional leadership willing to say “we will not do that, even if it is profitable,” and governance structures willing to support those choices.

The model also requires political advocacy. Public funding at levels that anchor the mission requires universities to demonstrate public benefit and to articulate why basic research, civic education, and cultural preservation are public goods worth funding. Many universities have ceded this argument. The protective model requires them to reclaim it – publicly, consistently, and with evidence.

What the model does not solve is the fundamental problem of rising costs. Universities are labour-intensive institutions. Faculty salaries, administrative staff, student support services, infrastructure maintenance, and technology investments all rise faster than most revenue streams. No internal allocation can reverse that structural reality.

What the model does is ensure that institutions facing cost pressures cannot simply abandon their public mission, cut research to preserve budgets, or shift entirely to employer-aligned training to maximise revenue. The funding architecture creates friction, requiring institutions to make explicit choices and trade-offs rather than drifting toward mission compromise.

The model also does not address extreme inequality across institutions. Elite universities with massive endowments can pursue diversified funding easily. Mid-tier public research universities can implement it with political will. Community colleges, regional universities, and Fachhochschulen face tighter constraints. For these institutions, the protective model requires significant public funding increases or regional partnerships that share costs and infrastructure. This is politically difficult but necessary if the goal is a diverse higher education landscape rather than a bifurcated one where elite institutions pursue the full mission and everyone else operates as a training facility.

Finally, the model assumes that societies continue to value the functions universities serve and are willing to fund them. If public funding continues its long-term decline – as it has in the US and UK – no internal allocation can prevent mission drift. At that point, the question is no longer institutional. It is political: does a society still believe that basic research, civic education, broad access, and independent inquiry are public goods worth paying for? The protective funding model makes this question explicit. If the answer is no, institutions cannot be blamed for the consequences. The society that withdrew funding has already made the choice.

The Risk of Mission Drift

The most insidious threat to universities in the AI era is not disruption from outside but drift from within. Mission drift happens when the metrics institutions optimize for gradually replace the purposes those metrics were meant to serve.

- **When efficiency crowds out inquiry.** AI makes it possible to measure and optimize many aspects of education: completion rates, time-to-degree, employer satisfaction, skills alignment. These metrics are useful, but they are not the mission. When universities are evaluated primarily on efficiency and labor-market outcomes, they will rationally invest in programs that produce measurable results quickly – vocational training, professional certificates, employer-aligned curricula – and disinvest in programs whose value is harder to quantify: philosophy, history, pure mathematics, basic science, the arts. This is not a hypothetical concern. It is already happening at institutions under performance-based funding pressure.
- **When employer alignment replaces broad education.** The emphasis on skills-based hiring and employer co-design is valuable within limits. But when universities allow employer demand to dictate curriculum design, they become training facilities rather than educational institutions. Training prepares people for the jobs that exist today. Education prepares people to create the jobs, ideas, and institutions of tomorrow. A university system optimized entirely for current employer needs will produce a workforce well suited to the present and poorly equipped for an uncertain future – precisely the opposite of what an era of rapid change requires.
- **When liberal arts become optional.** The humanities and social sciences are often the first casualties of mission drift, because their value is diffuse and long-term. A literature course does not produce a measurable skill. A philosophy seminar does not align with a job description. But these disciplines cultivate the interpretive, ethical, and critical capacities that distinguish education from training. In an AI era, where machines can perform an expanding range of technical tasks, the distinctively human capacities fostered by liberal education – moral reasoning, aesthetic judgment, historical perspective, imaginative empathy – become more valuable, not less. Universities that treat liberal arts as dispensable luxuries are discarding their most important asset.
- **When civic education disappears from the curriculum.** Democratic citizenship is not a natural byproduct of technical competence. It requires deliberate cultivation: exposure to diverse perspectives, practice in deliberation,

understanding of institutional design, and engagement with the ethical dimensions of public life. When universities optimize for employment outcomes and credential efficiency, civic education is the silent casualty. No employer asks for it. No ranking system measures it. And no society can sustain democratic governance without it.

Guarding against mission drift requires institutional leadership that articulates and defends the university's public purpose, governance structures that give faculty a meaningful voice in strategic decisions, and funding models that do not force universities to choose between financial survival and intellectual integrity.

Design Anchors for the AI-Era University

What would it mean to design universities for the AI era without abandoning the functions that make them irreplaceable? The following principles are not a blueprint. They are anchors – commitments that should constrain institutional design choices even as universities adapt to new technologies, new funding models, and new forms of credentialing.

1. Research is non-negotiable. A university that does not conduct research is a teaching college. Teaching colleges serve important functions, but they are not universities. The AI-era university must protect its research mission – including basic research with no immediate application – as a core institutional commitment, not a residual activity funded from whatever is left over after teaching and administration. This means dedicated research funding streams, protected time for faculty scholarship, and institutional cultures that value discovery alongside instruction.

1b. Humanities and foundational inquiry require structural protection from enrolment-driven funding. The preceding anchor protects research broadly, but specific mechanisms are needed for inquiry that serves no immediate labour market or enrolment demand: philosophy, history, classics, literary studies, basic mathematics, theoretical physics. These disciplines are vulnerable precisely because their value – cultural memory, historical perspective, the formation of judgment – is diffuse and long-term. Three structural mechanisms can protect them.

First: earmarked research funding separated from enrolment-driven teaching budgets. Germany's Exzellenzstrategie and Canada's Research Chairs programme demonstrate the model: public research funding allocated competitively to projects and scholars, with outcomes measured in research productivity rather than student throughput. This creates a dedicated revenue stream for unfunded inquiry, so that a philosophy department's budget does not depend on the number of philosophy majors completing degrees each year.

Second: regulatory cross-subsidisation – requiring institutions receiving public research funding to allocate a minimum 15–20% of research revenue to unfunded, investigator-initiated inquiry. This is stronger than voluntary good practice; it is a condition of receiving competitive grants. An institution receiving EUR 10 million in public research funding commits EUR 1.5–2 million to peer-reviewed research that serves no external funder's agenda. Implementation requires transparent reporting and periodic audit, but

it directly prevents funding structures from collapsing basic research capacity during periods of enrolment pressure.

Third: architectural separation of teaching and research budgets, so that enrolment decline does not automatically erode research capacity. Many universities now combine teaching and research revenue into unified departmental budgets, creating the perverse incentive that declining undergraduate enrolment in philosophy automatically reduces the department's research budget. A protective architecture separates them: teaching budgets come from tuition and public teaching appropriations; research budgets come from grants, endowments, or dedicated public research funding. A department that loses students loses teaching revenue but not automatically the capacity to conduct basic research. This separation is administratively more complex but strategically necessary for institutions serious about protecting humanities scholarship from enrolment volatility.

2. Democratic citizenship is a curriculum element, not an extracurricular option.

Every student who passes through a university should encounter structured opportunities to practice deliberation, engage with perspectives unlike their own, and develop the habits of democratic participation. This can take many forms – deliberative seminars, community-engaged learning, ethics courses embedded in professional programs, student governance with real authority – but it must be intentional and universal, not left to self-selected student organizations.

3. Public accountability is a structural requirement. Universities that receive public funding, grant credentials recognized by the state, or enjoy tax-exempt status have obligations to the public that go beyond satisfying their customers. They should publish transparent data on research outcomes, student demographics, completion rates by population, and the allocation of resources between teaching, research, and administration. Public accountability does not mean managerial control. It means that universities must be able to explain what they do and why it matters in terms the public can evaluate.

4. Equitable access is a design constraint, not an aspiration. Every institutional design choice – from tuition levels to program formats to admissions criteria to technology requirements – either expands or narrows who can participate. The AI-era university must treat access as a binding constraint on innovation: no new program, credential, or technology should be adopted without asking who it excludes and how exclusion can be mitigated. This is not a call for uniformity. It is a requirement that innovation serve the many, not just the self-directed, digitally fluent, and financially secure.

5. Human oversight of AI systems is mandatory. AI should augment teaching, research, and administration, but consequential decisions about students – admissions, assessment, academic standing, support interventions – must involve human judgment. Algorithms should be audited regularly for bias, their logic should be explainable to affected students, and institutions should maintain the capacity to override automated recommendations. The principle is simple: AI is a tool, not a decision-maker. Universities that cede consequential decisions to opaque algorithms betray their obligation to treat students as persons, not data points.

6. Institutional autonomy must be actively defended. In an era of platform dependence, political polarization, and commercial pressure, universities must build and maintain the structural conditions for independent inquiry. This means diversified funding, robust governance, legal protections for academic freedom, and the institutional courage to resist external pressure when it conflicts with the university's public mission. Autonomy is not a privilege. It is a precondition for the university to fulfill its social contract.

Conclusion

The AI era will transform universities. It should. Many of the reforms discussed in earlier chapters – competency-based learning, modular credentials, adaptive tutoring, employer partnerships, transnational alliances – address real problems and offer real benefits. But transformation without purpose is just disruption. And disruption without design anchors risks producing institutions that are efficient, modular, and market-responsive but no longer recognizably universities.

The argument of this chapter is that societies need universities to do things that no other institution will do: sustain basic research, cultivate democratic citizenship, preserve and challenge cultural memory, anchor local communities, and protect the conditions for independent inquiry. These functions are not obstacles to innovation. They are the reasons innovation matters. A university that optimizes for efficiency at the expense of its public mission has not adapted to the future. It has surrendered its reason for existing.

The AI-era university will be hybrid: partly online and partly residential, partly modular and partly integrated, partly employer-aligned and partly independent, partly automated and partly irreducibly human. The challenge is not to choose between tradition and innovation but to hold both in productive tension – embracing new tools and models while defending the commitments that make universities worth saving.

Three Futures for Education: Choice, Not Fate

Technology does not determine outcomes. Institutions, policies, and people do.

The same AI tutor that could democratize learning for millions can also widen the gap between those who can navigate a complex credential landscape and those who cannot. The same platform that matches workers to jobs in seconds can also reduce a human career to an algorithmic score. The same university that survives by cutting costs and adopting AI piecemeal can also reinvent itself as a place where young people learn to think, collaborate, and act with judgment.

Everything described in this book – the unbundling of universities, the rise of micro-credentials, the corporate takeover of workforce learning, the algorithmic reshaping of hiring – points not toward a single inevitable future but toward a branching set of possibilities. Which future we get depends on choices being made right now by governments, university leaders, employers, platform companies, and individuals. None of these actors is powerless. None is all-powerful. And none can afford to wait.

How to Read This Chapter

The chapters you have read describe real forces at work: AI is improving, institutions are under pressure, employers want better information about what candidates can do, micro-credentials are proliferating, hiring is becoming more algorithmic. These are observable trends, not speculations. But observable trends do not have predetermined endpoints.

A trend toward AI-powered personalization in education can lead to three different outcomes depending on who controls the systems and whom they serve. A trend toward employer-driven skills-based hiring can strengthen worker agency or concentrate power in platform companies' hands. A shift away from degree monopolies can democratize opportunity or fragment credential systems into proprietary noise.

The earlier chapters describe what is being attempted and what pressures are shaping education right now. This chapter asks: to what end? If these forces continue uncontested, where do they lead? What choices are we making implicitly by our current trajectories? And what choices do we still have?

The three futures that follow are not equally likely. But they are all possible, and they represent the major pathways that current dynamics could take. As you read the earlier chapters, it may help to ask: "Which future is this description primarily pointing toward?" Usually the answer is "all three in different ways" or "two out of three." Where the answer is "one specific future," that matters – because it means the outcome described is not inevitable, but conditional on a particular set of choices.

What follows are three plausible futures for education over the next decade. They are not predictions, and no one should read them as forecasts. They are scenarios – deliberately simplified, sharpened, and made concrete to illuminate the consequences of choices being made today. Each scenario pushes one set of tendencies to its logical

extreme to make the stakes visible. The real future will almost certainly be messier than any of them. We present them through the lives of two people you might recognize: an eighteen-year-old deciding what to do after secondary school, and a forty-year-old worker trying to stay relevant in a labor market that shifts faster than any career plan can follow.

Future 1: “Drift” – Business as Usual, but More Unequal

In this future, no one makes a dramatic decision. Universities adopt AI tools the way they adopted learning management systems two decades ago: slowly, unevenly, and without rethinking the underlying model. Lecture halls get smarter screens.

Assignments get AI-detection software. A few departments pilot adaptive tutoring. But the four-year degree remains the product, the semester remains the clock, and the credential remains the gate.

Governments continue to fund higher education through formulas designed in the twentieth century. Accreditation bodies update their standards cautiously. Employers talk about skills-based hiring at conferences but still filter resumes by degree when the applicant pool is large. The rhetoric changes. The structures do not.

Meet Lena, 18

Lena grows up in a middle-income family in a mid-sized European city. She is bright, curious about environmental science, and uncertain whether university is worth it. Her parents took on modest debt for their own degrees and assume she will do the same.

She enrolls in a public university. The first year feels familiar – large lectures, standardized exams, a handful of engaging seminars. Some professors use AI tools well; others ban them outright. Lena notices that the best AI tutoring she encounters is not from her university but from a free platform she finds on her own. She wonders why she is paying tuition for what she can access at home.

By her third year, she hears that employers in her field increasingly want candidates with data-analysis micro-credentials and project portfolios. Her university offers neither in a coherent way. She picks up a few online certificates on weekends, unsure whether anyone will recognize them. She graduates with a degree that feels like a floor pass to a building she has not yet entered – necessary to get through the door, but insufficient to get anywhere inside.

Meet Daniel, 40

Daniel has worked in logistics management for fifteen years. His employer, a large retail company, announces an AI transformation initiative. Daniel is offered a training module on supply-chain optimization with machine learning. The module is well-produced and takes two weeks.

But Daniel notices something. The training is tightly scoped to his employer’s systems and priorities. He learns to use the company’s proprietary platform, not transferable skills. When he asks about broader AI literacy courses or credentials he could take to an external employer, his manager looks puzzled. The company invested in Daniel’s productivity for the company, not his career mobility.

Daniel completes the training. His role changes modestly. He worries that if he is laid off in two years, the skills he gained will not travel with him. He is right to worry.

The result

In this future, the gap between elite and mass education widens. Well-funded research universities attract the best faculty, integrate AI most effectively, and produce graduates with strong networks and recognized brands. Mid-tier and regional universities struggle with flat budgets, aging infrastructure, and declining enrollment. They cut costs by increasing class sizes and adjunct faculty, which further erodes quality.

Credential inflation accelerates. A bachelor's degree buys less than it once did, but skipping it still carries a penalty. Micro-credentials proliferate without standardization, creating noise rather than signal. Employers default to familiar heuristics: name-brand degrees, referrals, and pedigree. Skills-based hiring remains a talking point, not a practice.

The people who do best are those who were already advantaged – those with family networks, financial cushions, and the self-direction to assemble their own learning pathways from scattered pieces. Everyone else drifts.

Future 2: “Platform World” – Corporations and AI Systems Lead

In this future, the private sector moves faster than public institutions and fills the vacuum. Big technology companies and EdTech platforms become the dominant providers of learning content, credentials, and talent matching. Universities do not disappear, but they shrink. Elite institutions survive as research institutes and finishing schools for the wealthy. The rest lose students to faster, cheaper, more employer-aligned alternatives.

Employers adopt AI not just for training but for hiring, evaluation, and promotion. Continuous algorithmic assessment replaces periodic reviews. Workers are scored, ranked, and routed through learning paths chosen by systems optimized for organizational productivity, not individual growth or public interest.

Meet Sofia, 18

Sofia decides not to attend university. Her older brother went, accumulated debt, and spent six months after graduation looking for work before landing a role unrelated to his degree. Sofia takes a different path.

She enrolls in a platform-based learning program sponsored by a major technology company. The program is free – funded by employers who pay the platform to access its talent pipeline. Sofia builds a portfolio of projects in data engineering, earns a stack of verified digital credentials, and is matched to an entry-level position at a logistics company through the platform’s AI hiring system within eight months of starting.

The match is efficient. The onboarding is smooth. Sofia is productive quickly. But she notices gaps. She never had the experience of debating ideas with peers from different disciplines. She has never written a sustained argument longer than a project brief.

When ethical questions arise in her work – about data privacy, about the environmental impact of the supply chain she optimizes – she has frameworks for the technical problem but not for the human one.

She also notices that her credentials are tied to the platform. If the platform changes its standards, restructures its partnerships, or is acquired by a competitor, her verified portfolio may lose recognition overnight. She does not own her learning record in any meaningful sense. The platform does.

Meet Marco, 40

Marco works as a regional sales manager for a manufacturing company. Two years ago, his employer adopted an AI-powered talent management system. Marco now receives weekly assessments: his communication patterns are analyzed, his customer interactions are scored, his learning modules are assigned based on gaps the system identifies.

Marco finds the system helpful in some ways. He discovers skills he did not know he lacked. The training content is good. But the experience feels relentless. There is no quarter when he is not being evaluated. His learning path is dictated entirely by what the company needs him to know, not by what interests him or what might prepare him for a different career. When he raises this with HR, he is told that the system optimizes for organizational capability.

Marco's performance scores are visible to his manager. He suspects they influence layoff decisions, though no one says so explicitly. He completes every assigned module. He does not explore anything beyond them.

The result

This future is efficient. Talent matching is faster. Skills gaps close more quickly. Companies get workers trained to their specifications. Platforms accumulate enormous datasets on human learning and performance, which they use to refine their products and extend their market power.

But the public mission of education – developing citizens, sustaining democratic culture, funding basic research, cultivating judgment and ethical reasoning – has no business model. No platform optimizes for it. No employer pays for it. Universities that once served these functions have been hollowed out. The knowledge produced is narrower, more applied, more aligned with commercial priorities. Critical inquiry, cultural memory, and the slow work of human formation are casualties of speed.

Surveillance is normalized. Workers accept continuous algorithmic evaluation because there is no alternative. The line between learning and performance management dissolves. Education becomes a function of employment, not a foundation for life.

Those who thrive are those who can game the system – who understand how the algorithms score them and optimize accordingly. Those who struggle are those who need more time, more support, more human connection, or whose talents do not map neatly onto platform categories.

Future 3: “Redesign” – A Public-Plus-Platform System

In this future, governments, universities, and employers act deliberately. They recognize that AI and platform economics are powerful forces, but they refuse to let those forces set the terms alone. Instead, they reshape the rules, the funding structures, and the institutional designs to harness technology while protecting public interests.

This is the hardest future. It requires political will, sustained investment, and coordination among actors with different incentives. It does not happen automatically. It happens because enough people decide it matters.

Meet Amir, 18

Amir finishes secondary school and faces a choice – but not the binary choice of previous generations. He does not have to pick between university and work. His country has introduced a modular higher education system in which students can combine a university foundation year with micro-credentials, work placements, and independent projects, assembling a personalized pathway that is recognized by employers and transferable across institutions.

Amir spends his first year at a regional university. The experience is intentionally designed: small seminars on reasoning and ethics, a team-based challenge project with a local municipality, and an introduction to AI tools used critically – not just as productivity aids but as objects of study. He learns to ask what data a model was trained on, whose interests it serves, and what it leaves out.

In his second year, Amir takes a semester-long work placement at a renewable energy company, earning micro-credentials in project management and energy systems modeling that stack toward his degree. He completes a third year back at the university, writing a capstone project that integrates his work experience with academic research. His credential is a hybrid: a recognized degree with embedded, verified evidence of practical skills and real-world experience.

Amir’s pathway is funded partly by public investment, partly by his employer during the work placement, and partly by a lifelong learning account the government established when he turned eighteen. The account carries a modest public subsidy that he can draw on for any accredited learning throughout his life. It is portable – it follows him, not his employer or his university.

Meet Priya, 40

Priya has worked as a hospital administrator for twelve years. Healthcare is being transformed by AI-driven diagnostics, scheduling, and patient management. Priya needs new skills, but she cannot afford to stop working for two years to get another degree.

She accesses her lifelong learning account – a fund she has been contributing to, with employer and government matching, since she entered the workforce. She uses it to enroll in a six-month hybrid program co-designed by a university and a health-technology company. The program combines online modules in health informatics with

in-person workshops at a university campus and a supervised project at her own hospital.

The credential she earns is recognized across borders. It was designed against a shared quality framework developed by universities, employers, and regulators working together. When Priya applies for a new role at a hospital in another country, the credential is legible to the hiring committee because it maps to a transparent competency standard.

Priya also notices something else. The program was not just technical. It included a module on the ethics of AI in healthcare – patient consent, algorithmic bias in diagnosis, the limits of automation in care. She found this module the most valuable part of the experience, not because it taught her a tool, but because it gave her a framework for deciding when not to use one.

The result

This future is not utopian. It is messy, politically contested, and unevenly implemented. Some countries move faster than others. Some universities resist modular reform. Some employers lobby against portable credentials because they prefer workforce lock-in. Lifelong learning accounts are underfunded in some regions and well-resourced in others.

But the direction is clear: education is treated as public infrastructure, not just private investment. Universities retain their role as places of research, critical inquiry, and human formation – but they shed their monopoly on credentialing and embrace modularity. Employers contribute to workforce development through shared frameworks rather than proprietary training silos. Governments fund lifelong learning as a right, not a privilege.

Redesign's Harder Path

Redesign is the hardest of the three futures to achieve – not because the vision is unclear, but because its enabling mechanisms are less proven than those of its alternatives. Drift requires only institutional inertia, which is free. Platform World requires only market dynamics and venture capital, forces already well-established in education. But Redesign requires sustained political will, coordinated multi-actor reform, public investment that competes against other priorities, and the deliberate suppression of forces – corporate consolidation, cost-cutting pressure, credential inflation – that operate autonomously otherwise.

The precedent base is real but partial. Singapore's SkillsFuture initiative demonstrates that governments can design integrated frameworks connecting schools, employers, and credentialing bodies around portable qualifications and lifelong learning accounts – but Singapore operates as a city-state with centralised governance and a strong social contract around public investment in human development. The UK's Credit Accumulation and Transfer framework and recent modular degree reforms show that universities can redesign credential structures to allow stackable, recognised micro-credentials – but implementation has been uneven, with quality variation across institutions and resistance from prestigious universities protecting traditional models. Austria's Higher Education Act reforms of the 2000s created pathways for universities

and technical colleges to offer equivalent qualifications with credit transfer across sectors – a structural reform that took over a decade. None of these is a perfect blueprint. But together they demonstrate that Redesign is neither utopian fantasy nor wholly unprecedented. The harder pathway may be the most important to argue for precisely because the easier alternatives happen largely by default. Redesign happens only through intentional choice, sustained investment, and political commitment to education as public infrastructure.

The result is more equitable than the alternatives, though far from perfect. The gap between those who can navigate the system and those who cannot narrows – not because the system is simpler, but because there are more on-ramps, more support structures, and more second chances built into the design.

What Pushes Us Where: Understanding the Feedback Loops

These futures are not random. Each is made more or less likely by specific decisions that institutions, governments, employers, and individuals are making right now. To understand why some choices accelerate change and others entrench the status quo, it helps to recognize the feedback loops that shape education – the five reinforcing patterns that either enable or block innovation.

Higher education is shaped by five reinforcing loops, each connecting regulatory choices, institutional strategy, hiring practices, and learner pathways. Understanding which loops are active and which are blocked tells you why change stalls or accelerates.

Loop 1: Accreditation → Institutional Structure → Program Design

When accreditation bodies require seat-time, faculty with terminal degrees, and physical classrooms, universities design programs that fit these rules. When rules allow competency-based models, micro-credentials, and distributed faculty, universities can – but do not automatically – redesign. The tighter the accreditation constraint, the more institutions converge on similar models. This loop is now *tight* in most places, creating inertia. Loosening it is necessary but not sufficient; institutions with tight budgets or risk-averse leadership will converge on the same model anyway.

Loop 2: Hiring Practices → Credential Demand → Institutional Investment

When employers filter resumes by degree, students demand degrees, and universities invest in bachelor's programs. When employers hire based on demonstrated skills and portfolios, students seek micro-credentials and apprenticeships, and universities invest in stackable credentials and work-integrated programs. This loop is *misaligned* today: employers say they want skills-based hiring but still filter by degree, sending contradictory signals that leave institutions uncertain where to invest. Resolving this misalignment requires explicit employer commitment to alternative hiring practices.

Loop 3: Funding Formula → Institutional Priority → Learner Access

When governments fund universities based on enrollment headcounts, institutions prioritize scale and completion, which can mean easier courses and larger cohorts. When governments fund based on learning outcomes or performance contracts, institutions prioritize assessment rigor, but this can inadvertently exclude low-achieving or under-prepared students. This loop is *mediated by redistribution*: governments can design formulas that reward both completion and serving difficult populations, but most do not. Shifting this loop requires deliberate policy design to avoid perverse incentives.

Loop 4: EdTech Market Dynamics → Platform Consolidation → Institutional Lock-in

When venture-backed platforms compete on feature richness and user engagement, they grow quickly, gain market share, and create switching costs for universities that integrate their tools into curriculum. Early adopters buy into a platform's infrastructure (content, assessments, data architecture), and later replacements become prohibitively expensive. This loop is *rapid and concentrated*: a handful of platforms now control large shares of learning management, adaptive content, and credential verification. Breaking this loop requires public alternatives and transparent standards for data portability.

Loop 5: Stratification → Credibility Divergence → Legitimacy Erosion

When universities divide into tiers (elite research universities, mid-tier teaching institutions, struggling regional campuses), the same credential means different things depending on its source. Employers learn that a degree from a top university predicts something different than a degree from a struggling institution. This credibility divergence increases as institutions in different tiers can no longer afford equivalent faculty, resources, or opportunities. The loop is *self-reinforcing*: tier-based legitimacy makes it harder for lower-tier institutions to attract talented students and faculty, which erodes quality further. This loop is particularly acute in systems with weak regulatory quality assurance and wide resource variation.

The Three Futures Through the Five Loops

In each future, different loops are locked in place or deliberately unlocked. This tells us not what will happen, but what happens *if current dynamics are left uncontested*.

Drift: All Loops Tighten

In Drift, all five loops remain unchanged or tighten further. Accreditation stays rigid (Loop 1 tightens). Hiring practices send mixed signals (Loop 2 remains misaligned). Funding formulas reward enrollment headcounts rather than learning outcomes (Loop 3 incentivizes scale over equity). EdTech consolidates under a handful of dominant platforms (Loop 4 tightens). Stratification deepens as resource gaps between elite and struggling institutions widen (Loop 5 accelerates). Each tightening constrains the others, creating a self-reinforcing stasis. Change becomes harder, not easier, as time passes.

Platform World: Three Loops Hijacked, Two Tighten

In Platform World, Loops 1, 3, and 5 effectively tighten as public regulation and funding weaken. But Loops 2 and 4 are hijacked: employers and platforms become the de facto accreditors and funders, replacing government and universities. Employers adopt platform-based hiring and assessment at scale (Loop 2 resolves, but in platforms' favor). EdTech consolidates around a monolithic or oligopolistic ecosystem (Loop 4 intensifies, but concentrates power). The result is that stratification accelerates sharply (Loop 5 tightens): elite universities and wealthy individuals can afford sophisticated platform ecosystems; everyone else gets locked into proprietary infrastructure they cannot own or exit.

Redesign: Loops are Intentionally Rebalanced

In Redesign, all five loops are deliberately loosened or reoriented toward public purposes. Accreditation loosens to allow competency-based and hybrid models (Loop 1 loosens). Hiring practices shift to skills-based hiring with transparent, shared standards (Loop 2 aligns around portable credentials). Funding formulas reward learning outcomes and serving underserved populations, not just enrollment (Loop 3 rebalances). EdTech is governed as public infrastructure with open standards and data portability (Loop 4 decelerates and deprivatizes). Investment is redistributed to prevent stratification and support struggling institutions (Loop 5 reverses). The loops are still present, but they are deliberately governed to serve public interests rather than left to markets or rigid bureaucracies.

This translation shows readers that the futures are not random: they follow from which loops policymakers choose to tighten or loosen.

What Pushes Us Where (Original Framework)

What pushes toward Drift

- Governments maintain legacy funding formulas that reward seat-time and enrollment counts rather than learning outcomes or modular credentials.
- Accreditation bodies move slowly, treating competency-based and hybrid models as exceptions rather than standard pathways.
- Employers adopt skills-based hiring rhetoric in public statements but continue to filter by degree in practice, especially when applicant volumes are high.
- Universities treat AI as an add-on rather than a reason to redesign pedagogy, assessment, and the structure of degrees.

What pushes toward Platform World

- Governments fail to regulate EdTech platforms and credential marketplaces, allowing private standards to become de facto public ones.
- Employers outsource workforce development to platform providers and tie credentials to proprietary ecosystems that workers cannot take with them.
- Public universities lose funding and enrollment to the point where they can no longer sustain core functions – research, community engagement, broad-access education.

- Workers accept continuous algorithmic evaluation as the price of employment, and no regulatory framework protects them from its excesses.

What pushes toward Redesign

- Governments invest in lifelong learning accounts, portable credential frameworks, and quality assurance systems that span universities, platforms, and employers.
- Universities proactively redesign degree structures to be modular and stackable while preserving the integrative, character-forming, and research functions that platforms cannot replicate.
- Employers, universities, and regulators co-design competency standards and credential recognition systems rather than leaving each to build in isolation.
- Citizens and policymakers treat education as critical public infrastructure – comparable to healthcare or transportation – deserving of sustained investment and democratic governance.

Choice, Not Fate

There is a tempting story about technology and education that goes like this: AI will disrupt universities the way streaming disrupted television, and the only question is how fast. This story is wrong – not because AI is unimportant, but because education is not a content-delivery business, even though parts of it have behaved like one.

Education is a system of institutions, incentives, regulations, relationships, and public commitments. It shapes who gets to learn what, whose knowledge counts, and who gets a fair chance at a decent life. These are not engineering problems with engineering solutions. They are political and institutional choices that reflect what a society values.

This book has laid out what is broken: universities that teach too slowly for a fast world, credentials that gatekeep more than they signal, corporate learning that serves employers more than workers, hiring systems that promise meritocracy and deliver something closer to pattern-matching. It has surveyed what is being tried: competency-based models, micro-credentials, AI-powered personalization, challenge-based learning, transnational alliances, employer-sponsored pathways. And it has documented what can go wrong: credential inflation, algorithmic bias, corporate capture, selection bias, stratification, the erosion of critical thinking and public mission.

The reader now has a map. Not a complete one – no map of a system this complex and fast-moving can be complete – but one detailed enough to see where the roads lead and where the choices lie.

The future of education is not something that happens to us. It is something we build, or fail to build, through the decisions we make in the next few years. The tools are powerful. The question is whether we use them to widen opportunity or to entrench advantage, to empower learners or to optimize them, to strengthen public institutions or to let them erode.

That question will not be answered by technology. It will be answered by us.

Methodological Note

This book draws on evidence of widely varying quality. Because education reform attracts strong claims from parties with strong incentives – universities defending their models, vendors selling platforms, governments justifying policies – readers deserve transparency about what kind of evidence underpins any given claim.

Throughout the manuscript, we have tried to signal evidence quality inline, using parenthetical caveats where the source or methodology matters. This note explains the hierarchy we applied.

Evidence Hierarchy

Tier 1: Meta-analyses and large-scale randomised controlled trials (RCTs)

These represent the strongest form of evidence. They include systematic reviews that aggregate findings across multiple studies and large-scale experiments with random assignment, control groups, and independent evaluation. Examples in this book include the meta-analytic evidence on active learning versus lectures and the 2015 meta-review of competency-based education. When Tier 1 evidence exists, we treat its conclusions as the most reliable available – though even meta-analyses have limitations (publication bias, heterogeneity across studies, and dependence on the quality of included research).

Tier 2: Quasi-experimental studies and multi-site evaluations

These studies compare outcomes across groups without full randomisation, often using statistical controls, matched samples, or natural experiments. Multi-site evaluations that track outcomes across several institutions or countries also fall here. Examples include cross-national employer surveys with large samples and OECD comparative analyses. Tier 2 evidence is informative but requires caution about confounding variables and selection effects.

Tier 3: Single-institution pilots, case studies, and observational data

Much of the evidence on innovative education models falls in this tier. A single university reporting outcomes from a new program, a pilot study at one institution, or observational data from a specific implementation are valuable for generating hypotheses and illustrating possibilities, but they cannot establish generalisable causal claims. Selection effects are a persistent concern: institutions that pilot new models often serve different populations than those that do not, making direct comparisons unreliable. Examples in this book include Minerva’s internal learning-gains studies, ECIU’s micro-credential pilot, and Utrecht University’s challenge-based learning observations.

Tier 4: Vendor-reported data, testimonials, and marketing claims

This is the most common – and least reliable – source of evidence in the education technology and corporate learning space. Vendor-reported data includes figures published by companies about their own products (e.g., platform adoption rates, skill acquisition speed, retention improvements) and institutional self-reports designed for marketing or fundraising. These figures are produced by parties with direct financial or reputational incentives to present positive results. They typically lack independent verification, transparent methodology, or control groups. Examples in this book include Guild Education’s mobility and retention figures, AI learning platform speed-of-acquisition claims, and corporate case studies on predictive analytics.

We do not dismiss Tier 4 evidence entirely – in many areas of educational innovation, it is the only data available, and ignoring it would leave the picture incomplete. But we flag it explicitly so that readers can calibrate their confidence accordingly.

How to Read the Evidence in This Book

Where a claim is supported by Tier 1 or Tier 2 evidence, we generally present it without extended qualification. Where a claim rests primarily on Tier 3 or Tier 4 evidence, we have added inline caveats indicating the source type – for example, “(vendor-reported),” “(single-institution pilot),” “(institutional self-report),” or “(industry survey; definitions vary).”

These caveats are not meant to discredit the claims but to help readers distinguish between what is well-established and what is promising but unverified. In a field where bold claims travel faster than rigorous evidence, this distinction matters.

The evidence base for educational innovation is improving, but it remains thinner than the confidence of the claims made about it. We have tried to be honest about that throughout.

The Diagnosis–Prescription Evidence Gap

Transparency about evidence quality reveals a structural asymmetry in this book that readers should weigh carefully. The diagnostic claims – that higher education costs are rising unsustainably, that credential inflation erodes signal quality, that employer demand is shifting toward demonstrated capability, that equity gaps are widening – draw primarily on Tier 1–2 evidence: cross-national labour force data, OECD comparative statistics, large-scale employer surveys with independent replication, and econometric analyses of returns to education. These claims rest on solid empirical ground.

The prescriptive claims – that competency-based education works at scale, that employer-sponsored learning expands opportunity, that peer-driven models like 42 School produce superior outcomes, that AI tutoring accelerates learning by 30–60%, that European university alliances represent a viable alternative architecture – draw predominantly on Tier 3–4 evidence: single-institution pilots, vendor-reported data, institutional self-reports, and infrastructure demonstrations without outcome evaluation. This is not because the authors chose weak evidence. It is because stronger

evidence does not yet exist for most of these innovations. The models are too new, too few, or too heterogeneous for the kind of multi-site, controlled, longitudinal studies that would constitute Tier 1–2 evidence.

The implication is that this book’s diagnosis is substantially stronger than its reform agenda. We are confident that the current system is failing in the ways described. We are less confident that the specific models and mechanisms we discuss will deliver the outcomes their proponents claim. The honest framing is: these are plausible, promising, and in some cases partially validated directions – not proven solutions. Where we recommend investment or institutional redesign, we do so on the basis of directional evidence and design logic, not established efficacy.

This gap is not unusual in a field undergoing rapid structural change – the evidence inevitably lags the innovation. But it means that readers, policymakers, and institutional leaders who act on the prescriptive chapters should do so with appropriate caution: invest in pilot evaluation, demand independent outcome data, and resist the temptation to scale models before they have been tested beyond their founding context.

Geographic Scope and Its Limits

The evidence and examples in this book draw predominantly from OECD countries, with a centre of gravity in Europe and North America. This is partly a function of data availability – large-scale labour market studies, employer surveys, and institutional outcome data are most systematically collected in these regions – and partly a function of the authors’ own professional networks and expertise.

The imbalance is uneven across chapters. The analysis of funding models, regulatory frameworks, and public governance (Chapters 9–11) leans heavily on European systems – Nordic, German, Dutch, Austrian – where the relationship between state funding and institutional autonomy is most explicitly codified. The discussion of alternative credentialing and employer-led learning (Chapters 6–7, Chapter 9) draws more heavily on US and Anglophone examples, where market-driven experimentation has gone furthest. The Global South appears primarily through targeted examples – Rwanda’s coding academies, India’s skill certification schemes, Singapore’s SkillsFuture – rather than through sustained structural analysis.

Readers should be aware of this geographic centre of gravity. The diagnostic claims about credential inflation, cost escalation, and employer demand shifts are supported by data from multiple regions and are likely to hold broadly. The prescriptive proposals – particularly around governance reform, funding architecture, and quality assurance – may transfer less cleanly to systems with fundamentally different institutional histories, labour market structures, or state capacities. We have tried to flag where a recommendation is context-dependent, but the default frame remains OECD-centric, and readers working in or on other systems should calibrate accordingly.

Failure Modes — How the Hybrid Future Goes Wrong

The hybrid future described in this book is not inevitable, and it is not self-correcting. It is a *design space*, not a guarantee. Under real political, economic, and institutional pressures, hybrid systems can fail in predictable ways.

This appendix names those failure modes explicitly—not to induce pessimism, but to make *early warning signals* visible. Most of these failures do not arrive as catastrophes. They arrive as reasonable compromises that quietly compound.

Failure Mode 1: “AI on Top of Old Structures”

What goes wrong: Institutions adopt AI tools without redesigning pedagogy, assessment, or governance.

Symptoms - AI tutors bolted onto lecture-based courses

- Automated grading layered onto unchanged exams
- Faculty workload increases instead of decreases
- Students experience fragmentation rather than coherence

Why it fails The institution keeps all legacy costs while adding new complexity. AI becomes an accelerant of burnout rather than a lever for redesign.

Early warning sign Staff describe AI as “one more thing to manage” rather than something that removes work.

Failure Mode 2: Platform Capture of Learning Pathways

What goes wrong: AI platforms quietly become the de facto curriculum designers.

Symptoms - Learning sequences optimized for engagement metrics

- Proprietary skill taxonomies replacing public standards
- Universities dependent on vendor roadmaps
- Students locked into non-portable learning histories

Why it fails Control over *how learning is structured* shifts from public institutions to private platforms whose incentives prioritize scale, retention, and monetization.

Early warning sign Institutions can no longer explain—or modify—how learning recommendations are generated.

Failure Mode 3: The Agency Divide Hardens into a Caste System

What goes wrong: Hybrid systems reward self-directed learners while quietly abandoning others.

Symptoms - High performers accelerate dramatically

- Struggling learners disengage earlier
- “Personalization” masks withdrawal of human support
- Equity rhetoric persists while outcomes diverge

Why it fails Agency is treated as a prerequisite instead of a capability to be developed. The system optimizes for those who already know how to use it.

Early warning sign Success stories cluster among already-privileged students, while failure is framed as individual responsibility.

Failure Mode 4: Credential Chaos and Trust Collapse

What goes wrong: The signaling ecosystem fragments faster than trust mechanisms evolve.

Symptoms - Proliferation of micro-credentials with unclear value

- Employers overwhelmed by incomparable signals
- Reputational inflation replaces credential inflation
- Informal networks regain gatekeeping power

Why it fails Degrees lose monopoly status before robust alternatives gain legitimacy. Informality fills the gap, often reinforcing existing power structures.

Early warning sign Hiring decisions increasingly rely on referrals and pedigree despite “skills-based” rhetoric.

Failure Mode 5: Education Reduced to Continuous Reskilling

What goes wrong: Learning becomes permanent employability maintenance.

Symptoms - Curriculum tightly coupled to short-term labor demand

- Humanities and fundamental inquiry marginalized
- Learners trapped in perpetual upskilling anxiety
- Civic and ethical dimensions quietly dropped

Why it fails Education collapses into workforce optimization. Its public mission erodes, and with it the case for public funding and autonomy.

Early warning sign Educational value is discussed almost exclusively in ROI terms.

Failure Mode 6: Governance Paralysis by Hybrid Complexity

What goes wrong: Hybrid systems become too complex to govern coherently.

Symptoms - Overlapping accountability between humans and algorithms

- No clear ownership of learning outcomes
- Ethics committees without enforcement power
- Decision-making slows despite technological acceleration

Why it fails Institutions inherit the complexity of platforms without matching governance capacity. Responsibility diffuses until nothing is truly accountable.

Early warning sign When failures occur, no one can say who is responsible—or how to fix them.

Failure Mode 7: Academic Freedom Erodes by Optimization

What goes wrong: AI-driven efficiency subtly reshapes what is teachable.

Symptoms - Content optimized for assessability

- Controversial or ambiguous topics de-emphasized
- Research agendas influenced by data availability
- Metrics crowd out intellectual risk

Why it fails Optimization logic privileges what can be measured and automated. Intellectual exploration that resists formalization is quietly penalized.

Early warning sign Faculty begin self-censoring curriculum choices to avoid friction with systems.

Failure Mode 8: Human Roles Hollow Out

What goes wrong: Human educators are repositioned as system supervisors rather than intellectual agents.

Symptoms - Faculty reduced to facilitators of predesigned flows

- Loss of pedagogical autonomy
- Decline in mentorship quality
- Identity crisis among educators

Why it fails Hybrid systems forget that judgment is cultivated through *relationships*, not interfaces.

Early warning sign Educators report feeling interchangeable.

Failure Mode 9: Public Trust Collapses Quietly

What goes wrong: The system continues to function, but belief in its legitimacy erodes.

Symptoms - Cynicism among students and parents

- Education treated as transactional necessity
- Rise of parallel credential markets
- Political vulnerability of public funding

Why it fails Once trust is lost, no amount of technical improvement restores legitimacy.

Early warning sign People ask not “Is this education good?” but “Is it worth the hassle?”

The meta-failure: Mistaking inevitability for design

The most dangerous failure mode is the belief that the hybrid future will “sort itself out.”

It will not.

Hybrid systems amplify existing values. If those values are left implicit, power will flow toward efficiency, scale, and convenience—because those forces move fastest.

Avoiding these failures requires:

- explicit normative commitments

- deliberate institutional subtraction

- governance capacity equal to technical capability

- continuous attention to second- and third-order effects

The hybrid future is not a destination. It is a continuous negotiation between human judgment and machine optimization.

When that negotiation stops, the system fails—quietly, efficiently, and at scale.

References (incomplete)

- 42KL. 2025. “42 Kuala Lumpur.” <https://42kl.edu.my/>.
- AICharts. 2025. “AI and Human Skills: Societal Impact 2025.”
<https://www.aicerts.ai/news/ai-and-human-skills-societal-impact-2025/>.
- Arizona State University / University Design Institute. 2025. “Not If but How: Redesigning the Future of Higher Education.” <https://udi.asu.edu/co-lab/article/Not-If-But-How-Redesigning-the-Future-of-Higher-Education>.
- Bates, Tony. 2014. “The Strengths and Weaknesses of Competency-Based Learning in a Digital Age.” <https://www.tonybates.ca/2014/09/15/the-strengths-and-weaknesses-of-competency-based-learning-in-a-digital-age/>.
- BBC News. 2016. 42: *The School Where You Teach Yourself*.
<https://www.bbc.com/news/business-37694248>.
- Changing Higher Ed. 2025. “Higher Ed Challenges 2025: Solution Examples.”
<https://changinghighered.com/higher-ed-challenges-2025-solution-examples/>.
- Claudius Press. 2025. *AI Bias in Educational Systems*.
https://www.claudiuspress.com/assets/default/article/2025/03/22/article_1742634483.pdf.
- Congruence Market Insights. 2025. “Higher Education Technology Market Report.”
<https://www.congruencemarketinsights.com/report/higher-education-technology-market>.
- Contrary Research. 2025. “Guild Education: Company Profile.”
<https://research.contrary.com/company/guild>.
- CORE. 2015. *Competence-Based Education: A Meta-Review of Research*.
<https://core.ac.uk/download/pdf/77612209.pdf>.
- Deloitte. 2025. *2025 US Higher Education Trends*.
<https://www.deloitte.com/us/en/insights/industry/public-sector/2025-us-higher-education-trends.html>.
- Deseret News. 2013. “Better than Harvard? Minerva Project Aims to Transform Elite Education.” [https://www.deseret.com/2013/11/7/20528963/better-than-harvard-minerva-project-aims-to-transform-elite-education/](https://www.deseret.com/2013/11/7/20528963/better-than-harvard-minerva-project-aims-to-transform-elite-education).
- European Higher Education Area. 2024. *The European Higher Education Area in 2024*.
<https://ehea.info/Immagini/the-european-higher-education-area-in-2024-EC0224018ENN.pdf>.
- European University Association. 2015. *Designing Strategies for Efficient Funding of Universities in Europe (DEFINE)*.
https://www.eua.eu/images/designing_strategies_for_efficient_funding_of_universities_in_europe_define.pdf.

- European University Association. 2024. *Trends 2024*.
<https://www.eua.eu/publications/reports/trends-2024.html>.
- Eurydice / EHEA. 2024. *Chapter 5: Learning and Teaching*.
https://eurydice.eacea.ec.europa.eu/sites/default/files/2024-05/Chapter_5_Learning_and_teaching.pdf.
- Frontiers in Education. 2023. “Dropout in Higher Education: A Systematic Review.” *Frontiers in Education*.
<https://www.frontiersin.org/journals/education/articles/10.3389/feduc.2023.1244686/full>.
- Full Fabric. 2025. “Challenges in Higher Education Management.”
<https://www.fullfabric.com/articles/challenges-in-higher-education-management>.
- Goedmo. 2025. “Problems in Higher Education.” <https://goedmo.com/blog/problems-in-higher-education/>.
- Hanover Research. 2025. “5 Ways Higher Education Is Changing in 2025–26.”
<https://www.hanoverresearch.com/reports-and-briefs/higher-education/5-ways-higher-education-is-changing-in-2025-26/>.
- Huang, Jensen. 2024. “The Future of AI and Education.”
<https://www.youtube.com/watch?v=sjGFJNY2v1k>.
- KPMG. 2020. *Future of Higher Education*.
<https://assets.kpmg.com/content/dam/kpmg/xx/pdf/2020/10/future-of-higher-education.pdf>.
- Müller. 2020. *Private Funding and Its Dangers to Academia*. University of Bern.
https://www.oefre.unibe.ch/ueber_uns/personen/abt_prof_mueller/dokumente/private_fundinganditsdangerstoacademia.pdf.
- OECD. 2025a. *Education at a Glance 2025*.
https://www.oecd.org/en/publications/education-at-a-glance-2025_1c0d9c79-en.html.
- OECD. 2025b. *Education at a Glance 2025: How Is Tertiary Education Financed?*
https://www.oecd.org/en/publications/education-at-a-glance-2025_1c0d9c79-en/full-report/how-is-tertiary-education-financed_2845d742.html.
- OECD. 2025c. *Trends Shaping Education 2025*.
https://www.oecd.org/en/publications/trends-shaping-education-2025_e66587fd-en.html.
- PMC. 2021. *Algorithmic Bias in AI Systems: Sources and Mitigation*.
<https://PMC8455229/>.
- PMC. 2024. *European Universities Initiative: Cooperation and Hierarchy*.
<https://PMC11461708/>.
- PR Newswire. 2017. “Guild Raises \$21 Million to Transform Corporate Tuition Assistance.” <https://www.prnewswire.com/news-releases/guild-raises-21-million-to->

<transform-corporate-tuition-assistance-and-increase-college-access-for-frontline-employees-300514706.html>.

Rensimer, Luke, and Rachel Brooks. 2024. “European Universities Initiative and Stratification in Higher Education.” *Compare: A Journal of Comparative and International Education*.

<https://discovery.ucl.ac.uk/id/eprint/10188673/1/Rensimer%20&%20Brooks%202024%20-%20EUI%20stratification%20-%20Compare%20-%20final%20version.pdf>.

Rogers, Bruce. 2024. “Guild Creates Platform for Employer-Provided Education.” *Forbes*. <https://www.forbes.com/sites/brucerogers/2024/06/05/guild-creates-platform-for-employer-provided-education/>.

Schiller International University. 2025. “Risks of AI Algorithmic Bias in Higher Education.” <https://www.schiller.edu/blog/risks-of-ai-algorithmic-bias-in-higher-education/>.

SeatsOne. 2025. “The Ultimate Guide for Universities in 2025: Addressing Higher Education’s Challenges.” <https://seatsone.ai/the-ultimate-guide-for-universities-in-2025-addressing-higher-educations-challenges/>.

Selingo, Jeff. 2013. “Harvard, Stanford ... Minerva? The next Elite University at Half the Price.” <https://www.linkedin.com/pulse/harvard-stanford-minerva-next-elite-university-half-price-selingo>.

Taylor & Francis. 2024a. *Bologna Process and European Higher Education Area Developments*.

<https://www.tandfonline.com/doi/full/10.1080/21568235.2024.2398742>.

Taylor & Francis. 2024b. “European Universities Initiative: Stratification and Inequality.” *Compare: A Journal of Comparative and International Education*.

<https://www.tandfonline.com/doi/full/10.1080/03057925.2024.2307551>.

University Foundation / EUF. 2024. “Internationalisation of Higher Education: Challenges, Trends, Priorities.” <https://uni-foundation.eu/2024/02/08/internationalisation-higher-education-challenges-trends-priorities/>.

University World News. 2025. *Higher Education Funding and Policy Developments*. <https://www.universityworldnews.com/post.php?story=2025010808011331>.

Verfassungsblog. 2024. “Academic Freedom and Security.” <https://verfassungsblog.de/academic-freedom-security/>.

Additional Reading

A curated guide for readers who want to go deeper into the themes explored in this book: the structural failures of higher education, the rise of alternative learning models, AI in education and hiring, and the future of credentials and skills.

Further Reading

Sources that extend, update, or deepen the arguments presented in this book.

- **OECD, *Education at a Glance* (annual series).** The most comprehensive cross-country dataset on educational attainment, spending, and outcomes. Essential for anyone tracking how systems compare and evolve.
https://www.oecd.org/en/publications/education-at-a-glance-2025_1c0d9c79-en.html
- **European University Association, *Trends 2024*.** Maps the strategic direction of European universities, including digitalisation, funding pressures, and institutional autonomy. <https://www.eua.eu/publications/reports/trends-2024.html>
- **World Economic Forum, *Future of Jobs Report* (biennial).** Tracks the skills employers say they need and the gap between education and labour market demand. Foundational for the book's arguments about signalling and skills.
- **World Bank, *Learning Poverty* reports.** Documents the global crisis in foundational learning, providing context for why reform is not merely a rich-country concern.
- **Deloitte, *2025 US Higher Education Trends*.** Industry perspective on enrolment decline, financial sustainability, and the competitive pressure from alternative providers. <https://www.deloitte.com/us/en/insights/industry/public-sector/2025-us-higher-education-trends.html>
- **OECD, *Trends Shaping Education 2025*.** A forward-looking policy overview connecting demographic, technological, and economic shifts to education strategy. https://www.oecd.org/en/publications/trends-shaping-education-2025_e66587fd-en.html
- **European Higher Education Area, *The EHEA in 2024*.** The Bologna Process stocktaking report, documenting the state of degree recognition, quality assurance, and mobility across 49 countries. <https://ehea.info/Immagini/the-european-higher-education-area-in-2024-EC0224018ENN.pdf>
- **KPMG, *Future of Higher Education* (2020).** An early industry signal of the structural changes that have since accelerated, including the shift to modular and digital-first delivery.
<https://assets.kpmg.com/content/dam/kpmg/xx/pdf/2020/10/future-of-higher-education.pdf>

- **Hanover Research, 5 Ways Higher Education Is Changing in 2025-26.** A practitioner-oriented summary of near-term shifts in enrolment strategy, programme design, and institutional partnerships. <https://www.hanoverresearch.com/reports-and-briefs/higher-education/5-ways-higher-education-is-changing-in-2025-26/>
- **Arizona State University / University Design Institute, Not If But How.** A case study in institutional redesign from one of the most-cited examples of large-scale university innovation. <https://udi.asu.edu/co-lab/article/Not-If-But-How-Redesigning-the-Future-of-Higher-Education>

Introductory Material

Accessible starting points for readers new to education reform, credentialism, and alternative learning.

- **Ken Robinson, Out of Our Minds: Learning to Be Creative (2011) and TED Talks.** Robinson's work remains the most widely shared argument that schools systematically suppress creativity. A good entry point for readers unfamiliar with the critique of standardised education.
- **Bryan Caplan, The Case Against Education (2018).** The strongest economic argument that most of education's value is signalling rather than human capital formation. Central to the book's Chapter 5 on the signalling stack.
- **Paul Tough, The Years That Matter Most (2019).** A narrative account of how college admissions and completion are shaped by class, showing that access alone does not solve the equity problem.
- **Ryan Craig, A New U: Faster + Cheaper Alternatives to College (2018).** Surveys bootcamps, apprenticeships, and staffing models as emerging alternatives to the four-year degree. Directly relevant to Chapters 4 and 5.
- **OECD, Trends Shaping Education 2025.** For readers who prefer data-driven policy overviews, this is the most accessible entry point to the global forces reshaping education. https://www.oecd.org/en/publications/trends-shaping-education-2025_e66587fd-en.html
- **David Epstein, Range: Why Generalists Triumph in a Specialized World (2019).** A useful counterweight to narrow skills-based arguments, making the case for breadth and transfer learning.
- **Tony Bates, The Strengths and Weaknesses of Competency-Based Learning in a Digital Age (2014).** A balanced practitioner perspective on CBE that helps contextualise Chapter 4's discussion of Western Governors University and similar models. <https://www.tonybates.ca/2014/09/15/the-strengths-and-weaknesses-of-competency-based-learning-in-a-digital-age/>

Academic Sources

Peer-reviewed research and substantive reports underpinning the book's claims.

- **Morcke, Dornan, and Eika, “Outcome (competency) based education: An exploration of its origins, theoretical basis, and empirical evidence” (2013).** A meta-review of CBE effectiveness that informs the book’s assessment of competency-based models. <https://core.ac.uk/download/pdf/77612209.pdf>
- **Rensimer and Brooks, “European Universities Initiative and stratification in higher education” (2024).** Peer-reviewed analysis showing that EU-funded university alliances may be reinforcing rather than reducing institutional hierarchies.
<https://discovery.ucl.ac.uk/id/eprint/10188673/1/Rensimer%20&%20Brooks%202024%20-%20EUI%20stratification%20-%20Compare%20-%20final%20version.pdf>
- **Frontiers in Education, “Dropout in Higher Education: A Systematic Review” (2023).** Synthesises the evidence on why students leave, confirming that dropout is driven by structural and institutional factors, not merely individual failure.
<https://www.frontiersin.org/journals/education/articles/10.3389/feduc.2023.1244686/full>
- **Freeman et al., “Active learning increases student performance in science, engineering, and mathematics” (2014), PNAS.** The landmark meta-analysis showing that active learning outperforms traditional lectures, supporting the book’s critique of passive pedagogy.
- **PMC, “Algorithmic Bias in AI Systems: Sources and Mitigation” (2021).** Surveys the mechanisms through which AI systems reproduce and amplify existing biases, directly relevant to Chapters 8 and 10.
<https://pmc.ncbi.nlm.nih.gov/articles/PMC8455229/>
- **Schiller International University, “Risks of AI Algorithmic Bias in Higher Education” (2025).** Examines how algorithmic bias manifests specifically in admissions, assessment, and student support systems.
<https://www.schiller.edu/blog/risks-of-ai-algorithmic-bias-in-higher-education/>
- **EUA, *Designing Strategies for Efficient Funding of Universities in Europe (DEFINE)*.** The foundational report on higher education funding models in Europe, documenting the shift from block grants to performance-based funding.
https://www.eua.eu/images/designing_strategies_for_efficient_funding_of_universities_in_europe_define.pdf
- **Muller, “Private Funding and Its Dangers to Academia” (University of Bern, 2020).** Analyses the risks of increasing private funding dependence in public universities, relevant to Chapter 9’s discussion of universities as public institutions.
https://www.oefre.unibe.ch/ueber_uns/personen/abt_prof_mueller/dokumente/privatefundinganditsdangerstoacademia.pdf
- **PMC, “European Universities Initiative: Cooperation and Hierarchy” (2024).** Empirical analysis of whether EU-level cooperation programmes are levelling or

stratifying the higher education landscape.

<https://PMC11461708/>

- **Verfassungsblog, “Academic Freedom and Security” (2024).** A legal and policy analysis of threats to academic freedom in the European context.
<https://verfassungsblog.de/academic-freedom-security/>

Tools and Commercial Offerings

Platforms, institutions, and tools discussed or referenced in the book. Descriptions are neutral and informational.

- **Khan Academy / Khanmigo.** Free learning platform offering video-based instruction across subjects; Khanmigo is its AI tutoring layer built on GPT-4, providing personalised guidance and Socratic questioning.
<https://www.khanacademy.org>
- **Guild Education.** Employer-sponsored education platform connecting frontline workers to curated degree, certificate, and skilling programmes with employer-paid tuition. Discussed in Chapter 7. <https://www.guild.com>
- **Western Governors University (WGU).** The largest competency-based university in the US, offering accredited degrees where students advance by demonstrating mastery rather than accumulating seat time. Central to Chapter 4.
<https://www.wgu.edu>
- **Gloat.** AI-powered talent marketplace used by large enterprises to match employees to internal opportunities, projects, and learning based on inferred skills. Discussed in Chapters 7 and 8. <https://www.gloat.com>
- **Eightfold.ai.** Talent intelligence platform using AI to map skills across workforces, support skills-based hiring, and reduce reliance on degree-based screening. Discussed in Chapter 8. <https://www.eightfold.ai>
- **42 School.** Tuition-free, peer-driven coding school with no teachers, no lectures, and no formal prerequisites. Students learn through project-based collaboration. Discussed in Chapter 4. <https://42kl.edu.my/>
- **Minerva University.** Campus-free university offering a globally rotating residential programme with seminar-based active learning. Positions itself as an alternative to elite traditional universities. <https://www.minerva.edu>
- **Europass / European Digital Credentials.** The EU’s framework for issuing and verifying digital credentials, including micro-credentials, across member states. Part of the broader push toward portable, interoperable qualifications.
- **ECIU University / Micro-credentials.** A consortium of European universities piloting challenge-based micro-credentials designed for lifelong learners and working professionals.

Divergent and Opposing Views

Sources that challenge, complicate, or directly contradict the book's arguments. Included in the interest of intellectual honesty.

- **Andrew Delbanco, *College: What It Was, Is, and Should Be* (2012).** A defence of the liberal arts college as a space for civic formation and intellectual maturation, arguing that instrumental critiques of higher education miss its deeper purpose.
- **Johann Neem, *What's the Point of College?* (2019).** Argues that the university's value lies precisely in its separation from labour market pressures, pushing back against the skills-and-employability framing.
- **Fuller and Raman, “Dismissed by Degrees” (Harvard Business School, 2017).** While often cited in favour of skills-based hiring, the report also documents the real risks: skills-based screening can disadvantage workers who lack professional networks, and employers frequently revert to degree requirements under uncertainty.
- **Ben Williamson, *Big Data in Education* (2017).** A critical analysis of how data-driven and AI approaches in education risk reducing learning to measurable behaviours, entrenching surveillance, and serving commercial rather than pedagogical interests.
- **Cathy O’Neil, *Weapons of Math Destruction* (2016).** Documents how algorithmic decision-making in education, hiring, and public services amplifies inequality. A counterpoint to optimistic accounts of AI in hiring (Chapter 8).
- **Eric Hanushek and Ludger Woessmann, *The Knowledge Capital of Nations* (2015).** Argues that cognitive skills (as measured by standardised tests) are the strongest predictor of economic growth, providing an evidence-based defence of rigorous assessment that challenges the book’s scepticism toward testing.
- **Matthew Crawford, *The Case for Working with Your Hands* (2009).** Argues that the push toward cognitive and digital work devalues skilled trades and embodied knowledge, complicating the book’s emphasis on digital credentials and AI-mediated learning.
- **David Deming, “The Growing Importance of Social Skills in the Labor Market” (2017), *Quarterly Journal of Economics*.** While broadly compatible with the book’s argument about non-cognitive skills, Deming’s work also shows that traditional degree programmes remain effective at developing the social skills employers value most.

Funding Models by Countries

European universities are predominantly financed by public funds (local authority taxes), supplemented by tuition fees (especially in the UK and for some master's programs), third-party funding for research (e.g., research councils, EU, industry), and other own-source revenue (continuing education, services, foundations). [europa+1](#)

Main sources of funding

- Basic public funding: core budgets of universities from national, regional, and local budgets; in Europe usually >70% of university income, in many countries >80%. [eua+1](#)
- Private households: tuition fees, administrative fees, accommodation and canteen services; Proportion varies greatly between countries (very low in the Nordics, very high in England and the US). [oecd+1](#)
- Companies and non-profit organizations: research contracts, endowed professorships, collaborations, donations; in Europe, usually <30% of total university income. [\[eua\]](#)
- Financial aid for students: Public expenditure goes partly directly to students (scholarships, loans) and partly to universities; a high proportion of tertiary education expenditure, particularly in the Netherlands, Scandinavia, and Ireland. [europa+1](#)

Germany

- According to the OECD, around 80–90% of expenditure on tertiary education institutions in Germany comes from public sources (federal, state, and local governments), with private sources (households and businesses) accounting for around 10–20%. [gpseducation.oecd+1](#)
- In the non-tertiary sector, the private share is around 11%, while in the tertiary sector it is slightly higher due to fees charged by non-governmental providers and third-party funding. [\[gpseducation.oecd\]](#)
- Third-party funding (e.g., DFG, EU, industry) is important within university budgets: in 2022, German universities received approximately €27 billion in basic funding and €10.4 billion in third-party funding (third-party funding ratio well over a quarter of total income). [\[foerderatlas.dfg\]](#)

Germany – rough structure of tertiary financing

Category	Share/Classification
Public funds	predominantly >70% of university income. gpseducation.oecd+1
Total private funds	approx. 10–20% (households, companies). gpseducation.oecd+1
Third-party funding in higher	approx. 28–30% of higher education funding (including research). [foerderatlas.dfg]

Category	Share/Classification
education budgets	

France

- In France, approximately 69% of tertiary education expenditure (before transfers) comes from public sources; the remainder comes from private households and other private sources. [\[oecd\]](#)
- France is thus one of the more publicly funded systems, but is slightly below the OECD average for public funding. [\[oecd+1\]](#)

France – rough structure of tertiary financing

Category	Share/Classification
Public funds	approx. 69% of tertiary expenditure. [oecd]
Total private sources	Approximately 30% (households, businesses, non-profit organizations). [oecd+1]

United Kingdom (especially England)

- The UK, especially England, is an example of high tuition fees with comprehensive loan systems: the private share of higher education funding (especially fees) is well above the OECD average and among the highest in Europe. [\[oecd+1\]](#)
- Public funds are mainly allocated to research funding and targeted teaching/student support; basic funding for teaching is largely covered by (loan-financed) fees. [\[oecd+1\]](#)

United Kingdom – rough structure of tertiary funding

Category	Classification
Public funds	significantly lower than continental European systems, but still relevant for research and grants. [oecd+1]
Private funding (mainly fees)	one of the highest shares in Europe, core funding for teaching. [oecd+1]

(Exact percentages vary depending on the data set and year; the OECD identifies England as a “high tuition–high aid” model with predominantly private cost sharing by students.) [\[oecd\]](#)

USA

- In the USA, only about 39% of tertiary education expenditure comes from public sources, with the rest coming from private sources (households, donations, companies, fees); the public share is thus well below the OECD average. [oecd+1](#)
- Tuition fees play a central role: at public universities, fees account for an average of around 20-30% of income, at private non-profit universities significantly more, and at private for-profit providers, over 90% of income comes from fees. [educationdata+1](#)

USA – rough structure of tertiary financing

Category	Classification/Order of magnitude
Total public funding	Approximately 39% of tertiary funding. oecd+1
Total private funding	approx. 60%+ (fees, donations, contracts). oecd+1
Public university fees	Approximately 20–30% of revenue. [educationdata]
Private for-profit tuition fees	>90% of revenue. [journalistsresource]

Europe as a whole – pattern

- Eurostat and Eurydice show that, in the EU, education is overwhelmingly publicly funded; private sources (including households) play a greater role in tertiary education, but remain below 30% of total expenditure in many countries. [europa+1](#)
- Models range from “low fee–high public” (Nordic countries, Germany, France) to mixed models to “high fee–high aid” (England) and systems dominated by the private sector (USA). [eua+2](#)